

# Sistemi Operativi I

Corso di Laurea in Informatica  
2025-2026



SAPIENZA  
UNIVERSITÀ DI ROMA

Gabriele Tolomei

Dipartimento di Informatica  
Sapienza Università di Roma

[tolomei@di.uniroma1.it](mailto:tolomei@di.uniroma1.it)

# Scheduling Algorithms: An Overview

- First-Come-First-Serve (FCFS)
- Round Robin (RR)
- Shortest-Job-First (SJF)
- Priority Scheduling

- Multilevel Queue (MQ)
- Multilevel Feedback-Queue (MFQ)

LAST TIME

# Scheduling Algorithms: An Overview

- First-Come-First-Serve (FCFS)
- Round Robin (RR)
- Shortest-Job-First (SJF)
- Priority Scheduling
- Multilevel Queue (MQ)
- Multilevel Feedback-Queue (MFQ)

TODAY

# Scheduling Algorithms: An Overview

- First-Come-First-Serve (FCFS)
- Round Robin (RR)
- Shortest-Job-First (SJF)
- Priority Scheduling
- **Multilevel Queue (MLQ)**
- Multilevel Feedback-Queue (MLFQ)

## MLQ: Idea

- Use multiple separate queues, one for each job **category**

## MLQ: Idea

- Use multiple separate queues, one for each job **category**
- Each queue implements whatever scheduling algorithm is most appropriate for that type of jobs

## MLQ: Idea

- Use multiple separate queues, one for each job **category**
- Each queue implements whatever scheduling algorithm is most appropriate for that type of jobs
- Scheduling must be done between queues!

# MLQ: Idea

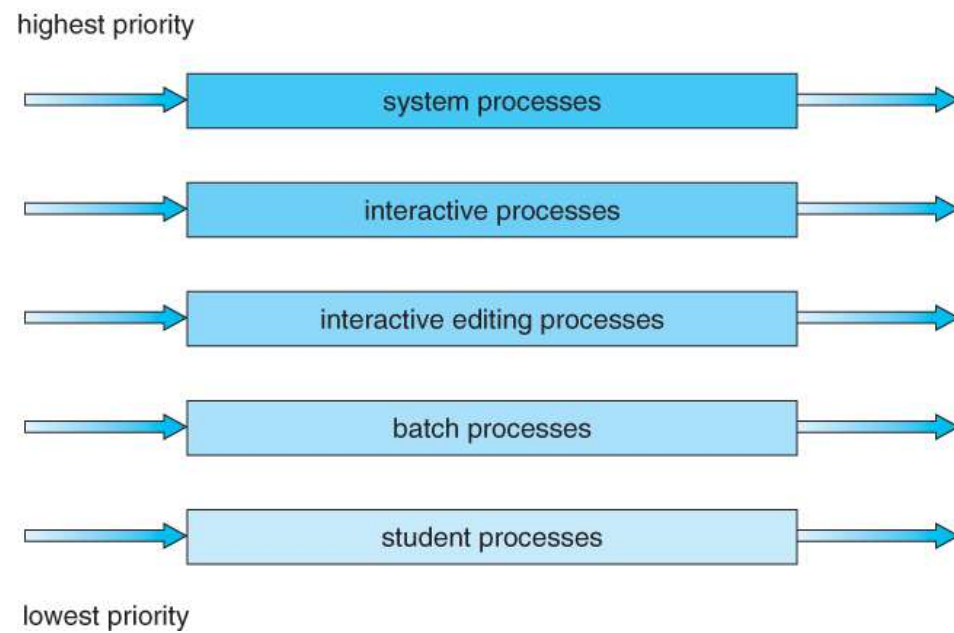
- Use multiple separate queues, one for each job **category**
- Each queue implements whatever scheduling algorithm is most appropriate for that type of jobs
- Scheduling must be done between queues!
- Two common options are:
  - **strict priority** → no job in a lower priority queue runs until all higher priority queues are empty
  - **round-robin** → each queue gets a time slice in turn, possibly of different sizes



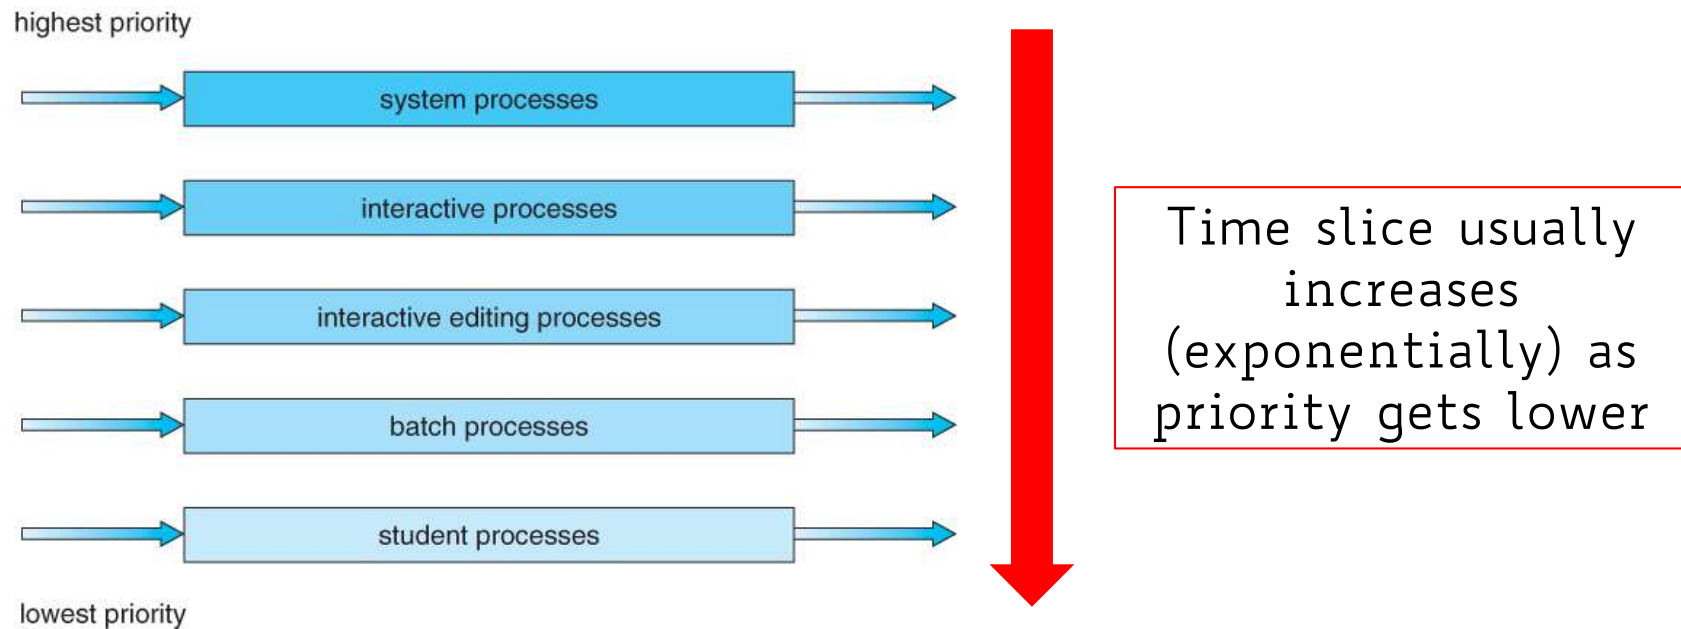
# MLQ: Idea

- Use multiple separate queues, one for each job **category**
- Each queue implements whatever scheduling algorithm is most appropriate for that type of jobs
- Scheduling must be done between queues!
- Two common options are:
  - **strict priority** → no job in a lower priority queue runs until all higher priority queues are empty
  - **round-robin** → each queue gets a time slice in turn, possibly of different sizes

# MLQ: Overview



# MLQ: Overview



# Scheduling Algorithms: An Overview


- First-Come-First-Serve (FCFS)
- Round Robin (RR)
- Shortest-Job-First (SJF)
- Priority Scheduling
- Multilevel Queue (MLQ)
- **Multilevel Feedback-Queue (MLFQ)**

## MLFQ: Idea

- Similar to the ordinary MLQ scheduling, except jobs may be moved from one queue to another

# MLFQ: Idea

- Similar to the ordinary MLQ scheduling, except jobs may be moved from one queue to another
- Moving jobs may be required when:
  - The characteristics of a job change between CPU-intensive and I/O-intensive
  - A job that has waited for a long time can get bumped up into a higher priority queue for a while (to compensate the aging problem)



# MLFQ: Adjusting Job Priority

- Job starts in the highest priority queue (by default)

# MLFQ: Adjusting Job Priority

- Job starts in the highest priority queue (by default)
- If job's time slice expires → drop its priority level by one unit



# MLFQ: Adjusting Job Priority

- Job starts in the highest priority queue (by default)
- If job's time slice expires → drop its priority level by one unit
- If job's time slice does not expire (i.e., the context switch occurs due to an I/O request, instead) → increase its priority level by one unit (up to the top) or stay the same

# MLFQ: Adjusting Job Priority

- Job starts in the highest priority queue (by default)
- If job's time slice expires → drop its priority level by one unit
- If job's time slice does not expire (i.e., the context switch occurs due to an I/O request, instead) → increase its priority level by one unit (up to the top) or stay the same
- CPU-bound jobs will quickly drop their priority

# MLFQ: Adjusting Job Priority

- Job starts in the highest priority queue (by default)
- If job's time slice expires → drop its priority level by one unit
- If job's time slice does not expire (i.e., the context switch occurs due to an I/O request, instead) → increase its priority level by one unit (up to the top) or stay the same
- CPU-bound jobs will quickly drop their priority
- I/O-bound jobs will stay at higher priority levels

# MLFQ: Idea

- MLFQ is the most flexible but it is also the most complex to implement

# MLFQ: Idea

- MLFQ is the most flexible but it is also the most complex to implement
- Some of the (many) parameters which define MLFQ systems include:
  - The number of queues
  - The scheduling algorithm for each queue
  - The methods used to upgrade or demote processes from one queue to another
  - The method used to determine which queue a process enters initially

# Multilevel Feedback Queue (MLFQ): Example I

New A B C

Order	Job	CPU burst (time units)
1	A	30
2	B	20
3	C	10

# Multilevel Feedback Queue (MLFQ): Example I

New A B C

Order	Job	CPU burst (time units)
1	A	30
2	B	20
3	C	10

No I/O burst

Initial time quantum = 1

Context switch = 0

3 queues

# Multilevel Feedback Queue (MLFQ): Example I

New A B C

Order	Job	CPU burst (time units)
1	A	30
2	B	20
3	C	10

No I/O burst

Initial time quantum = 1

Context switch = 0

3 queues

strict priority  
between queues



# Multilevel Feedback Queue (MLFQ): Example I

New A B C

Order	Job	CPU burst (time units)
1	A	30
2	B	20
3	C	10

$\text{JOB\_ID}^{\text{job\_exec\_time}}_{\text{total\_elapsed\_time}}$  = The job JOB\_ID has executed *job\_exec\_time* time units after *total\_elapsed\_time* time units

$A^2_7$  = The job A has executed 2 time units after 7 time units overall

# Multilevel Feedback Queue (MLFQ): Example I

New A B C

Order	Job	CPU burst (time units)
1	A	30
2	B	20
3	C	10

Queue	Time Slice (time units)	Jobs
1	1	
2	2	
3	4	

# Multilevel Feedback Queue (MLFQ): Example I

New A B C

Order	Job	CPU burst (time units)
1	A	30
2	B	20
3	C	10

Queue	Time Slice (time units)	Jobs
1	1	$A^1_1, B^1_2, C^1_3$
2	2	
3	4	

# Multilevel Feedback Queue (MLFQ): Example I

New A B C

Order	Job	CPU burst (time units)
1	A	30
2	B	20
3	C	10

Queue	Time Slice (time units)	Jobs
1	1	A <sup>1</sup> <sub>1</sub> , B <sup>1</sup> <sub>2</sub> , C <sup>1</sup> <sub>3</sub>
2	2	A <sup>3</sup> <sub>5</sub> , B <sup>3</sup> <sub>7</sub> , C <sup>3</sup> <sub>9</sub>
3	4	

# Multilevel Feedback Queue (MLFQ): Example I

New A B C

Order	Job	CPU burst (time units)
1	A	30
2	B	20
3	C	10

Queue	Time Slice (time units)	Jobs
1	1	$A^1_1$ , $B^1_2$ , $C^1_3$
2	2	$A^3_5$ , $B^3_7$ , $C^3_9$
3	4	$A^7_{13}$ , $B^7_{17}$ , $C^7_{21}$

# Multilevel Feedback Queue (MLFQ): Example I

New A B C

Order	Job	CPU burst (time units)
1	A	30
2	B	20
3	C	10

Queue	Time Slice (time units)	Jobs
1	1	$A^1_1$ , $B^1_2$ , $C^1_3$
2	2	$A^3_5$ , $B^3_7$ , $C^3_9$
3	4	$A^7_{13}$ , $B^7_{17}$ , $C^7_{21}$ $A^{11}_{25}$ , $B^{11}_{29}$ , $C^{10}_{32}$

# Multilevel Feedback Queue (MLFQ): Example II

New A B C

Order	Job	CPU burst (time units)
1	A	30
2	B	20
3	C	10

# Multilevel Feedback Queue (MLFQ): Example II

New A B C

Order	Job	CPU burst (time units)
1	A	30
2	B	20
3	C	10

2 queues and C now  
alternates 1 time  
unit of CPU with 1  
time unit of I/O

Queue	Time Slice (time units)	Jobs
1	1	
2	2	



# Multilevel Feedback Queue (MLFQ): Example II

New A B C

Order	Job	CPU burst (time units)
1	A	30
2	B	20
3	C	10

2 queues and C now alternates 1 time unit of CPU with 1 time unit of I/O

Queue	Time Slice (time units)	Jobs
1	1	<span style="color: purple;">A</span> <sup>1</sup> <sub>1</sub> , <span style="color: green;">B</span> <sup>1</sup> <sub>2</sub> , <span style="color: red;">C</span> <sup>1</sup> <sub>3</sub>
2	2	

# Multilevel Feedback Queue (MLFQ): Example II

New A B C

Order	Job	CPU burst (time units)
1	A	30
2	B	20
3	C	10

2 queues and C now alternates 1 time unit of CPU with 1 time unit of I/O

Queue	Time Slice (time units)	Jobs
1	1	<span style="color: purple;">A</span> <sup>1</sup> <sub>1</sub> , <span style="color: green;">B</span> <sup>1</sup> <sub>2</sub> , <span style="color: red;">C</span> <sup>1</sup> <sub>3</sub>
2	2	<span style="color: purple;">A</span> <sup>3</sup> <sub>5</sub>

# Multilevel Feedback Queue (MLFQ): Example II

New A B C

Order	Job	CPU burst (time units)
1	A	30
2	B	20
3	C	10

2 queues and C now alternates 1 time unit of CPU with 1 time unit of I/O

Queue	Time Slice (time units)	Jobs
1	1	<span style="color: purple;">A</span> <sub>1</sub> <sup>1</sup> , <span style="color: green;">B</span> <sub>2</sub> <sup>1</sup> , <span style="color: red;">C</span> <sub>3</sub> <sup>1</sup> , <span style="color: red;">C</span> <sub>6</sub> <sup>2</sup>
2	2	<span style="color: purple;">A</span> <sub>5</sub> <sup>3</sup>

# Multilevel Feedback Queue (MLFQ): Example II

New A B C

Order	Job	CPU burst (time units)
1	A	30
2	B	20
3	C	10

2 queues and C now alternates 1 time unit of CPU with 1 time unit of I/O

Queue	Time Slice (time units)	Jobs
1	1	<span style="color: purple;">A</span> <sub>1</sub> <sup>1</sup> , <span style="color: green;">B</span> <sub>2</sub> <sup>1</sup> , <span style="color: red;">C</span> <sub>3</sub> <sup>1</sup> , <span style="color: red;">C</span> <sub>6</sub> <sup>2</sup>
2	2	<span style="color: purple;">A</span> <sub>5</sub> <sup>3</sup> , <span style="color: green;">B</span> <sub>8</sub> <sup>3</sup>

# Multilevel Feedback Queue (MLFQ): Example II

New A B C

Order	Job	CPU burst (time units)
1	A	30
2	B	20
3	C	10

2 queues and C now alternates 1 time unit of CPU with 1 time unit of I/O

Queue	Time Slice (time units)	Jobs
1	1	<span style="color: purple;">A</span> <sup>1</sup> <sub>1</sub> , <span style="color: green;">B</span> <sup>1</sup> <sub>2</sub> , <span style="color: red;">C</span> <sup>1</sup> <sub>3</sub> , <span style="color: red;">C</span> <sup>2</sup> <sub>6</sub> , <span style="color: red;">C</span> <sup>3</sup> <sub>9</sub>
2	2	<span style="color: purple;">A</span> <sup>3</sup> <sub>5</sub> , <span style="color: green;">B</span> <sup>3</sup> <sub>8</sub>

# Multilevel Feedback Queue (MLFQ): Example II

New A B C

Order	Job	CPU burst (time units)
1	A	30
2	B	20
3	C	10

2 queues and C now alternates 1 time unit of CPU with 1 time unit of I/O

Queue	Time Slice (time units)	Jobs
1	1	<span style="color: purple;">A</span> <sup>1</sup> <sub>1</sub> , <span style="color: green;">B</span> <sup>1</sup> <sub>2</sub> , <span style="color: red;">C</span> <sup>1</sup> <sub>3</sub> , <span style="color: red;">C</span> <sup>2</sup> <sub>6</sub> , <span style="color: red;">C</span> <sup>3</sup> <sub>9</sub>
2	2	<span style="color: purple;">A</span> <sup>3</sup> <sub>5</sub> , <span style="color: green;">B</span> <sup>3</sup> <sub>8</sub> , <span style="color: purple;">A</span> <sup>5</sup> <sub>11</sub>

# Multilevel Feedback Queue (MLFQ): Example II

New A B C

Order	Job	CPU burst (time units)
1	A	30
2	B	20
3	C	10

2 queues and C now alternates 1 time unit of CPU with 1 time unit of I/O

Queue	Time Slice (time units)	Jobs
1	1	$A^1_1, B^1_2, C^1_3, C^2_6, C^3_9, C^4_{12}, \dots, C^{10}_{30}$
2	2	$A^3_5, B^3_8, A^5_{11}, B^5_{14}, \dots, B^{12}_{32}, A^{14}_{34}, \dots$

# MLFQ: Fairness Issue

- MLFQ tries to mimic the optimal behavior of SJF in terms of average waiting time



# MLFQ: Fairness Issue

- MLFQ tries to mimic the optimal behavior of SJF in terms of average waiting time
- It explicitly promotes short jobs (i.e., I/O-bound ones) by design

## MLFQ: Fairness Issue

- MLFQ tries to mimic the optimal behavior of SJF in terms of average waiting time
- It explicitly promotes short jobs (i.e., I/O-bound ones) by design
- **Problem:** SJF (and MLFQ) might be unfair (as opposed to RR)

Any increase in fairness by giving long jobs a fraction of the CPU when shorter jobs could be instead selected will increase waiting time

# MLFQ: Improving Fairness

- Give each queue a fraction of the CPU time
  - This is fair only if jobs are evenly distributed (i.e., uniformly) across queues

# MLFQ: Improving Fairness

- Give each queue a fraction of the CPU time
  - This is fair only if jobs are evenly distributed (i.e., uniformly) across queues
- Adjust dynamically the priority of jobs as they don't get scheduled
  - This avoids starvation but average waiting time might increase when the system is overloaded (all jobs get to the highest priority queue, eventually)

# Advanced Topics: Lottery Scheduling

- Give every job a certain number of lottery tickets

# Advanced Topics: Lottery Scheduling

- Give every job a certain number of lottery tickets
- On each time slice, pick a winning ticket **uniformly at random**

# Advanced Topics: Lottery Scheduling

- Give every job a certain number of lottery tickets
- On each time slice, pick a winning ticket **uniformly at random**
- CPU time will be assigned to jobs according to the original distribution of tickets

# Advanced Topics: Lottery Scheduling

- Give every job a certain number of lottery tickets
- On each time slice, pick a winning ticket **uniformly at random**
- CPU time will be assigned to jobs according to the original distribution of tickets

As the number of time slices  
(i.e., the number of random picks) goes to infinity

**Law of Large Numbers**



# Lottery Scheduling: Policy

- Assign tickets to jobs as follows:
  - Give more tickets to short running jobs
  - Give few tickets to long running jobs

# Lottery Scheduling: Policy

- Assign tickets to jobs as follows:
    - Give more tickets to short running jobs
    - Give few tickets to long running jobs
- } simulating SJF

# Lottery Scheduling: Policy

- Assign tickets to jobs as follows:
    - Give more tickets to short running jobs
    - Give few tickets to long running jobs
- } simulating SJF
- To avoid starvation, each job gets at least one ticket

# Lottery Scheduling: Policy

- Assign tickets to jobs as follows:
  - Give more tickets to short running jobs
  - Give few tickets to long running jobs
- To avoid starvation, each job gets at least one ticket
- Degrades gracefully as system load changes
  - Adding/deleting a job affects all the other jobs proportionally

simulating SJF

# Lottery Scheduling vs. All

## Question:

What is the main difference between lottery scheduling and any other algorithm we have seen so far?

# Lottery Scheduling vs. All

## Question:

What is the main difference between lottery scheduling and any other algorithm we have seen so far?

## Answer:

This is the only example of **randomized** scheduler  
(rather than deterministic one)

# Lottery Scheduling: Example

#short jobs / #long jobs	% of CPU for each short job	% of CPU for each long job

# Lottery Scheduling: Example

short jobs get 10 tickets each   long jobs get 1 ticket each

#short jobs / #long jobs	% of CPU for each short job	% of CPU for each long job



# Lottery Scheduling: Example

short jobs get 10 tickets each   long jobs get 1 ticket each

#short jobs / #long jobs	% of CPU for each short job	% of CPU for each long job
1/1		

# Lottery Scheduling: Example

short jobs get 10 tickets each   long jobs get 1 ticket each

#short jobs / #long jobs	% of CPU for each short job	% of CPU for each long job
1/1	~91% (10/11)	

# Lottery Scheduling: Example

short jobs get 10 tickets each   long jobs get 1 ticket each

#short jobs / #long jobs	% of CPU for each short job	% of CPU for each long job
1/1	~91% (10/11)	~9% (1/11)

# Lottery Scheduling: Example

short jobs get 10 tickets each   long jobs get 1 ticket each

#short jobs / #long jobs	% of CPU for each short job	% of CPU for each long job
1/1	~91% (10/11)	~9% (1/11)
0/2		

# Lottery Scheduling: Example

short jobs get 10 tickets each   long jobs get 1 ticket each

#short jobs / #long jobs	% of CPU for each short job	% of CPU for each long job
1/1	~91% (10/11)	~9% (1/11)
0/2	–	

# Lottery Scheduling: Example

short jobs get 10 tickets each   long jobs get 1 ticket each

#short jobs / #long jobs	% of CPU for each short job	% of CPU for each long job
1/1	~91% (10/11)	~9% (1/11)
0/2	–	50% (1/2)

# Lottery Scheduling: Example

short jobs get 10 tickets each   long jobs get 1 ticket each

#short jobs / #long jobs	% of CPU for each short job	% of CPU for each long job
1/1	~91% (10/11)	~9% (1/11)
0/2	–	50% (1/2)
2/0	50% (10/20)	–

# Lottery Scheduling: Example

short jobs get 10 tickets each   long jobs get 1 ticket each

#short jobs / #long jobs	% of CPU for each short job	% of CPU for each long job
1/1	~91% (10/11)	~9% (1/11)
0/2	–	50% (1/2)
2/0	50% (10/20)	–
10/1	~9.9% (10/101)	~0.99% (1/101)



# Lottery Scheduling: Example

short jobs get 10 tickets each   long jobs get 1 ticket each

#short jobs / #long jobs	% of CPU for each short job	% of CPU for each long job
1/1	~91% (10/11)	~9% (1/11)
0/2	–	50% (1/2)
2/0	50% (10/20)	–
10/1	~9.9% (10/101)	~0.99% (1/101)
1/10	50% (10/20)	5% (1/20)

# Lottery Scheduling: CPU Assignment

$n_{short}$  = total number of *short* jobs

$n_{long}$  = total number of *long* jobs

$N = n_{short} + n_{long}$  = total number of jobs

$m_{short}$  = number of tickets assigned to each *short* job

$m_{long}$  = number of tickets assigned to each *long* job

$M = m_{short} * n_{short} + m_{long} * n_{long}$  = total number of tickets

# Lottery Scheduling: CPU Assignment

$n_{short}$  = total number of *short* jobs

$n_{long}$  = total number of *long* jobs

$N = n_{short} + n_{long}$  = total number of jobs

$m_{short}$  = number of tickets assigned to each *short* job

$m_{long}$  = number of tickets assigned to each *long* job

$M = m_{short} * n_{short} + m_{long} * n_{long}$  = total number of tickets

$$\begin{aligned} \text{CPU}_{short} &= \frac{m_{short}}{M} \\ \text{CPU}_{long} &= \frac{m_{long}}{M} \end{aligned}$$

# Lottery Scheduling: CPU Assignment Probability

$$\begin{aligned} m_i &= \text{number of tickets assigned to job } i \\ N &= \text{total number of jobs} \\ M &= \sum_{i=1}^N m_i = \text{total number of tickets} \\ P(i) &= \frac{m_i}{M} = \text{probability of job } i \text{ being scheduled} \end{aligned}$$

# Summary of CPU Scheduling Algorithms

- **First-Come-First-Serve (FCFS)**: Very simple, non-preemptive, generally unfair, and highly variant waiting time

# Summary of CPU Scheduling Algorithms

- **First-Come-First-Serve (FCFS):** Very simple, non-preemptive, generally unfair, and highly variant waiting time
- **Round Robin (RR):** Fair but average waiting time may be long

# Summary of CPU Scheduling Algorithms

- **First-Come-First-Serve (FCFS):** Very simple, non-preemptive, generally unfair, and highly variant waiting time
- **Round Robin (RR):** Fair but average waiting time may be long
- **Shortest Job First (SJF):** Generally unfair and possible starvation, but provably optimal in terms of average waiting time (assuming we are able to correctly predict the length of the next CPU burst)

# Summary of CPU Scheduling Algorithms

- **First-Come-First-Serve (FCFS)**: Very simple, non-preemptive, generally unfair, and highly variant waiting time
- **Round Robin (RR)**: Fair but average waiting time may be long
- **Shortest Job First (SJF)**: Generally unfair and possible starvation, but provably optimal in terms of average waiting time (assuming we are able to correctly predict the length of the next CPU burst)
- **Multilevel Queuing (MLQ/MLFQ)**: An approximation of SJF



# Summary of CPU Scheduling Algorithms

- **First-Come-First-Serve (FCFS):** Very simple, non-preemptive, generally unfair, and highly variant waiting time
- **Round Robin (RR):** Fair but average waiting time may be long
- **Shortest Job First (SJF):** Generally unfair and possible starvation, but provably optimal in terms of average waiting time (assuming we are able to correctly predict the length of the next CPU burst)
- **Multilevel Queuing (MLQ/MLFQ):** An approximation of SJF
- **Lottery:** Fairer with a low average waiting time yet less predictable due to randomization

# Summary of CPU Scheduling Algorithms

- **First-Come-First-Serve (FCFS)**: Very simple, non-preemptive, generally unfair, and highly variant waiting time
- **Round Robin (RR)**: Fair but average waiting time may be long
- **Shortest Job First (SJF)**: Generally unfair and possible starvation, but provably optimal in terms of average waiting time (assuming we are able to correctly predict the length of the next CPU burst)
- **Multilevel Queuing (MLQ/MLFQ)**: An approximation of SJF
- **Lottery**: Fairer with a low average waiting time yet less predictable due to randomization