

# Sistemi Operativi I

Corso di Laurea in Informatica

2025-2026



SAPIENZA  
UNIVERSITÀ DI ROMA

Gabriele Tolomei

Dipartimento di Informatica  
Sapienza Università di Roma  
[tolomei@di.uniroma1.it](mailto:tolomei@di.uniroma1.it)

# Where Do We Stand?

- So far, we have addressed:
  - Processes and Threads
  - CPU Scheduling
  - Synchronization and Deadlock

# Where Do We Stand?

- So far, we have addressed:
  - Processes and Threads
  - CPU Scheduling
  - Synchronization and Deadlock
- Today, we will be talking about:
  - Memory Management

# Where Do We Stand?

- So far, we have addressed:
  - Processes and Threads
  - CPU Scheduling
  - Synchronization and Deadlock
- Today, we will be talking about:
  - Memory Management
- ... Later on:
  - File Systems and I/O Storage
  - Advanced Topics (?)

# Part IV: Memory Management

# Goals of Memory Management

- Allocate memory resources among multiple competing processes
  - maximizing memory utilization and system throughput

# Goals of Memory Management

- Allocate memory resources among multiple competing processes
  - maximizing memory utilization and system throughput
- Guarantee isolation between processes
  - addressability and protection

# Goals of Memory Management

- Allocate memory resources among multiple competing processes
  - maximizing memory utilization and system throughput
- Guarantee isolation between processes
  - addressability and protection
- Provide a convenient abstraction to the programmer
  - illusion of unlimited amount of memory

# From Source Code To Binary Executable

- User programs typically refer to instructions and data with **symbolic names**, such as "+", "&&", "x", "count", etc.

# From Source Code To Binary Executable

- User programs typically refer to instructions and data with **symbolic names**, such as "+", "&&", "x", "count", etc.
- For a user program to be executed it first needs to be:
  - 1) **Translated** from source code to binary executable and stored on disk

# From Source Code To Binary Executable

- User programs typically refer to instructions and data with **symbolic names**, such as "+", "&&", "x", "count", etc.
- For a user program to be executed it first needs to be:
  - 1) **Translated** from source code to binary executable and stored on disk
  - 2) **Loaded** from disk into main memory (RAM)

# From Source Code To Binary Executable

- Translation is done via Compiler/Assembler/Linker

# From Source Code To Binary Executable

- Translation is done via Compiler/Assembler/Linker
- Loading is done by the (OS) Loader

# From Source Code To Binary Executable

- Translation is done via Compiler/Assembler/Linker
- Loading is done by the (OS) Loader

## NOTE:

In case of purely-interpreted language implementations, translation from source code to executable is done "on-the-fly" by the loaded interpreter

# CPU Must Refer To Memory Addresses

- CPU repeatedly **fetches**, **decodes**, and **executes** instructions from memory while executing a program

# CPU Must Refer To Memory Addresses

- CPU repeatedly **fetches**, **decodes**, and **executes** instructions from memory while executing a program
- Most instructions involve retrieving data from memory or storing data in memory, or both

# CPU Must Refer To Memory Addresses

- CPU repeatedly **fetches**, **decodes**, and **executes** instructions from memory while executing a program
- Most instructions involve retrieving data from memory or storing data in memory, or both
- Memory chips only respond to actual physical addresses

# CPU Must Refer To Memory Addresses

- CPU repeatedly **fetches**, **decodes**, and **executes** instructions from memory while executing a program
- Most instructions involve retrieving data from memory or storing data in memory, or both
- Memory chips only respond to actual physical addresses
- It turns out that symbolic names used by user programs must be eventually bound to actual **physical memory** addresses

# CPU Must Refer To Memory Addresses

- CPU repeatedly **fetches**, **decodes**, and **executes** instructions from memory while executing a program
- Most instructions involve retrieving data from memory or storing data in memory, or both
- Memory chips only respond to actual physical addresses
- It turns out that symbolic names used by user programs must be eventually bound to actual **physical memory** addresses

How?

# Generating Memory Addresses: Example

```
void foo() {  
    int x = 42;  
    ...  
    x = x + 5;  
    ...  
}
```

# Generating Memory Addresses: Example

```
void foo() {  
    int x = 42;
```

...

```
    x = x + 5;
```

...

```
}
```

**x** is a symbolic  
name

# Generating Memory Addresses: Example

```
void foo() {  
    int x = 42;  
    ...  
    x = x + 5;  
    ...  
}
```

x is a symbolic  
name



```
128: MOV %R1, [%R2] // assuming R2 contains the  
      // address of x in memory  
136: ADD 5, %R1  
144: MOV [%R2], %R1
```

# Generating Memory Addresses: Example

```
void foo() {  
    int x = 42;  
    ...  
    x = x + 5;  
    ...  
}
```

x is a symbolic  
name



```
128: MOV %R1, [%R2] // assuming R2 contains the  
      // address of x in memory  
136: ADD 5, %R1  
144: MOV [%R2], %R1
```

references to logical addresses  
containing both instructions and data

# Generating Memory Addresses: Example

```
void foo() {  
    int x = 42;  
    ...  
    x = x + 5;  
    ...  
}
```

x is a symbolic name



```
128: MOV %R1, [%R2] // assuming R2 contains the  
      // address of x in memory  
136: ADD 5, %R1  
144: MOV [%R2], %R1
```

references to logical addresses  
containing both instructions and data

1. Fetch instruction at address 128

# Generating Memory Addresses: Example

```
void foo() {  
    int x = 42;  
    ...  
    x = x + 5;  
    ...  
}
```

x is a symbolic name



```
128: MOV %R1, [%R2] // assuming R2 contains the  
      // address of x in memory  
136: ADD 5, %R1  
144: MOV [%R2], %R1
```

references to logical addresses containing both instructions and data

1. Fetch instruction at address 128
2. Execute instruction: load from address [%R2] (e.g., 1234)

# Generating Memory Addresses: Example

```
void foo() {  
    int x = 42;  
    ...  
    x = x + 5;  
    ...  
}
```

x is a symbolic name



```
128: MOV %R1, [%R2] // assuming R2 contains the  
      // address of x in memory  
136: ADD 5, %R1  
144: MOV [%R2], %R1
```

references to logical addresses containing both instructions and data

1. Fetch instruction at address 128
2. Execute instruction: load from address [%R2] (e.g., 1234)
3. Fetch instruction at address 136
4. Execute instruction: addition (no memory reference)
5. Fetch instruction at address 144
6. Execute instruction: store to address [%R2] (1234)

# Symbolic Name vs. Logical vs. Physical Address

symbolic name: symbolic memory reference used by  
user programs

# Symbolic Name vs. Logical vs. Physical Address


**symbolic name:** symbolic memory reference used by user programs



**logical address:** memory address generated by user programs via the CPU

# Symbolic Name vs. Logical vs. Physical Address

**symbolic name:** symbolic memory reference used by user programs



**logical address:** memory address generated by user programs via the CPU



**physical address:** actual memory address which memory chip operates on

# Address Binding

Mapping from logical to physical address

# Address Binding

Mapping from logical to physical address

Compile  
time


# Address Binding

Mapping from logical to physical address



# Address Binding

Mapping from logical to physical address



# Address Binding: Compile Time

- The starting physical location  $k$  of a program in memory is known at compile time (e.g.,  $k = 0$ )

# Address Binding: Compile Time

- The starting physical location  $k$  of a program in memory is known at compile time (e.g.,  $k = 0$ )
- The compiler generates so-called **absolute code**

# Address Binding: Compile Time

- The starting physical location  $k$  of a program in memory is known at compile time (e.g.,  $k = 0$ )
- The compiler generates so-called **absolute code**
- No intervention by the OS

# Address Binding: Compile Time

- The starting physical location  $k$  of a program in memory is known at compile time (e.g.,  $k = 0$ )
- The compiler generates so-called **absolute code**
- No intervention by the OS
- **physical address == logical address**

# Address Binding: Compile Time

- The starting physical location  $k$  of a program in memory is known at compile time (e.g.,  $k = 0$ )
- The compiler generates so-called **absolute code**
- No intervention by the OS
- **physical address == logical address**

What if the starting physical memory address  $k$  where the program is loaded changes into  $k'$  at some point later?


# Address Binding: Compile Time

- The starting physical location  $k$  of a program in memory is known at compile time (e.g.,  $k = 0$ )
- The compiler generates so-called **absolute code**
- No intervention by the OS
- **physical address == logical address**

What if the starting physical memory address  $k$  where the program is loaded changes into  $k'$  at some point later?


The program must be recompiled!

# Address Binding: Compile Time




# Address Binding: Compile Time

starting physical address known by the compiler (e.g.,  $k = 0$ )




# Address Binding: Compile Time

starting physical address known by the compiler (e.g.,  $k = 0$ )




# Address Binding: Compile Time

starting physical address known by the compiler (e.g.,  $k = 0$ )



# Address Binding: Compile Time

starting physical address known by the compiler (e.g.,  $k = 0$ )



logical addresses **do not change** after they are emitted as the result of compilation/assembling  
physical addresses == logical addresses

# Address Binding: Load Time

- If the starting physical location  $k$  of a program is not known at compile time

# Address Binding: Load Time

- If the starting physical location  $k$  of a program is not known at compile time
- The compiler generates (**statically**) **relocatable code**, which references relative addresses to  $k$

# Address Binding: Load Time

- If the starting physical location  $k$  of a program is not known at compile time
- The compiler generates (**statically**) **relocatable code**, which references relative addresses to  $k$
- The OS loader determines each process' starting physical location  $k$

# Address Binding: Load Time

- If the starting physical location  $k$  of a program is not known at compile time
- The compiler generates (**statically**) **relocatable code**, which references relative addresses to  $k$
- The OS loader determines each process' starting physical location  $k$
- **physical address == logical address**

# Address Binding: Load Time

- If the starting physical location  $k$  of a program is not known at compile time
- The compiler generates (**statically**) **relocatable code**, which references relative addresses to  $k$
- The OS loader determines each process' starting physical location  $k$
- **physical address == logical address**


What if the OS decides to use a different starting physical memory address  $k'$ ?

# Address Binding: Load Time


- If the starting physical location  $k$  of a program is not known at compile time
- The compiler generates (**statically**) **relocatable code**, which references relative addresses to  $k$
- The OS loader determines each process' starting physical location  $k$
- **physical address == logical address**

What if the OS decides to use a different starting physical memory address  $k'$ ?


# Address Binding: Load Time




# Address Binding: Load Time




# Address Binding: Load Time




# Address Binding: Load Time




# Address Binding: Load Time



# Address Binding: Load Time



# Address Binding: Load Time



# Address Binding: Execution Time

- If the program can be moved around in main memory during its execution

# Address Binding: Execution Time

- If the program can be moved around in main memory during its execution
- The compiler generates (**dynamically**) **relocatable code** or **virtual addresses**

# Address Binding: Execution Time

- If the program can be moved around in main memory during its execution
- The compiler generates (**dynamically**) **relocatable code** or **virtual addresses**
- The OS maps virtual addresses to physical memory locations using special HW support (**MMU**)

# Address Binding: Execution Time


- If the program can be moved around in main memory during its execution
- The compiler generates (**dynamically**) **relocatable code** or **virtual addresses**
- The OS maps virtual addresses to physical memory locations using special HW support (**MMU**)
- **physical address != logical/virtual address**

# Address Binding: Execution Time


- If the program can be moved around in main memory during its execution
- The compiler generates (**dynamically**) **relocatable code** or **virtual addresses**
- The OS maps virtual addresses to physical memory locations using special HW support (**MMU**)
- **physical address != logical/virtual address**

Most flexible solution implemented by the majority of modern OSs


# Address Binding: Execution Time




# Address Binding: Execution Time




# Address Binding: Execution Time




# Address Binding: Execution Time




# Address Binding: Execution Time



# Address Binding: Execution Time



# Address Binding: Execution Time



# Managing Memory

- Key Assumptions:

- The virtual address space (VAS) of each process has the same size
- VAS must be entirely placed contiguously in main memory
- Main memory is (way) larger than VAS

# Managing Memory


- Key Assumptions:

- The virtual address space (VAS) of each process has the same size
- VAS must be entirely placed contiguously in main memory
- Main memory is (way) larger than VAS

Don't worry! We will soon relax those assumptions

# Managing Memory: VAS

virtual address space = 16 KiB



# Manage Memory: Goals (1)

- **Sharing**

- Several processes coexist in main memory at the same time
- Cooperating processes can share portions of address space

# Manage Memory: Goals (1)

- **Sharing**

- Several processes coexist in main memory at the same time
- Cooperating processes can share portions of address space

- **Transparency**

- Processes should not be aware that memory is shared
- Processes should not be aware of which portions of physical memory they are assigned to

# Manage Memory: Goals (2)

- **Protection/Security**
  - Processes must not be able to corrupt each other or the OS
  - Processes must not be able to read data of other processes

# Manage Memory: Goals (2)

- **Protection/Security**

- Processes must not be able to corrupt each other or the OS
- Processes must not be able to read data of other processes

- **Efficiency**

- CPU and memory performance should not degrade badly due to sharing
- Keep memory fragmentation low

# Relocation: Initial Idea

- The OS is allocated to the lowest/highest memory addresses

# Relocation: Initial Idea

- The OS is allocated to the lowest/highest memory addresses
- Assume logical addresses generated by each user process starts at 0 up to  $(\text{memory\_size} - \text{os\_size} - 1)$

# Relocation: Initial Idea

- The OS is allocated to the lowest/highest memory addresses
- Assume logical addresses generated by each user process starts at 0 up to  $(\text{memory\_size} - \text{os\_size} - 1)$
- Load a process by allocating it in one of the free **contiguous segments** of memory


# Relocation: Initial Idea

- The OS is allocated to the lowest/highest memory addresses
- Assume logical addresses generated by each user process starts at 0 up to  $(\text{memory\_size} - \text{os\_size} - 1)$
- Load a process by allocating it in one of the free **contiguous segments** of memory
- Allow transparent sharing of memory: each process' address space may be placed anywhere in memory

# Managing Memory: VAS Relocation

RAM size = 64 KiB

OS size = 16 KiB



# Static Relocation

- The OS loader rewrites the addresses generated by a process to main memory addresses (**load time binding**)

# Static Relocation

- The OS loader rewrites the addresses generated by a process to main memory addresses (**load time binding**)
- PRO:
  - No HW support is needed

# Static Relocation

- The OS loader rewrites the addresses generated by a process to main memory addresses (**load time binding**)
- PRO:
  - No HW support is needed
- CONs:
  - No protection/privacy → processes can corrupt the OS or other processes
  - Address space must be entirely allocated contiguously → free space between stack and heap can be huge and wasted!
  - The OS cannot move a process (address space) once allocated in memory

# Dynamic Relocation

- Protect OS and processes from one another

# Dynamic Relocation

- Protect OS and processes from one another
- Requires hardware support (**M**emory **M**anagement **U**nit)

# Dynamic Relocation


- Protect OS and processes from one another
- Requires hardware support (**M**emory **M**anagement **U**nit)
- Address binding at execution/run time

# Dynamic Relocation

- Protect OS and processes from one another
- Requires hardware support (**M**emory **M**anagement **U**nit)
- Address binding at execution/run time
- **MMU** dynamically translates each logical/virtual address generated by a process into a corresponding physical address

# Dynamic Relocation

- Protect OS and processes from one another
- Requires hardware support (**M**emory **M**anagement **U**nit)
- Address binding at execution/run time
- **MMU** dynamically translates each logical/virtual address generated by a process into a corresponding physical address



# Dynamic Relocation

- Protect OS and processes from one another
- Requires hardware support (**M**emory **M**anagement **U**nit)
- Address binding at execution/run time
- **MMU** dynamically translates each logical/virtual address generated by a process into a corresponding physical address



# HW Support for Dynamic Relocation

- MMU contains at least 2 registers:

# HW Support for Dynamic Relocation

- MMU contains at least 2 registers:
  - **base** → start physical memory location of address space

# HW Support for Dynamic Relocation

- MMU contains at least 2 registers:
  - **base** → start physical memory location of address space
  - **limit** → size limit of address space

# HW Support for Dynamic Relocation

- MMU contains at least 2 registers:
  - **base** → start physical memory location of address space
  - **limit** → size limit of address space
- CPU supports at least 2 operating modes:


# HW Support for Dynamic Relocation

- MMU contains at least **2 registers**:
  - **base** → start physical memory location of address space
  - **limit** → size limit of address space
- CPU supports at least **2 operating modes**:
  - **privileged (kernel) mode** when the OS is running
    - after issuing any trap (system call, interruption, or exception)
    - when manipulating sensitive resources (e.g., the content of MMU registers)

# HW Support for Dynamic Relocation


- MMU contains at least **2 registers**:
  - **base** → start physical memory location of address space
  - **limit** → size limit of address space
- CPU supports at least **2 operating modes**:
  - **privileged (kernel) mode** when the OS is running
    - after issuing any trap (system call, interruption, or exception)
    - when manipulating sensitive resources (e.g., the content of MMU registers)
  - **user mode** when user process is running
    - while executing process instructions on the CPU

# Base and Limit Registers: Idea



Each process is given a **contiguous segment** of main memory when loaded


# Base and Limit Registers: Idea



Each process is given a **contiguous segment** of main memory when loaded

As such, each process can only access to memory locations belonging to the segment it is assigned to

# Base and Limit Registers: Idea



Each process is given a **contiguous segment** of main memory when loaded

As such, each process can only access to memory locations belonging to the segment it is assigned to

Protection implemented using two MMU registers: **base** and **limit**

# Implementing Dynamic Relocation

CPU must check every memory access generated in user mode (i.e., by a user process) is within the correct [base, base + limit) range for that process


# Implementing Dynamic Relocation

CPU must check every memory access generated in user mode (i.e., by a user process) is within the correct [base, base + limit) range for that process




# Implementing Dynamic Relocation

CPU must check every memory access generated in user mode (i.e., by a user process) is within the correct [base, base + limit) range for that process




# Implementing Dynamic Relocation

CPU must check every memory access generated in user mode (i.e., by a user process) is within the correct [base, base + limit) range for that process




# Implementing Dynamic Relocation

CPU must check every memory access generated in user mode (i.e., by a user process) is within the correct [base, base + limit) range for that process




# Implementing Dynamic Relocation

CPU must check every memory access generated in user mode (i.e., by a user process) is within the correct [base, base + limit) range for that process




# Implementing Dynamic Relocation

CPU must check every memory access generated in user mode (i.e., by a user process) is within the correct [base, base + limit) range for that process




# Implementing Dynamic Relocation

CPU must check every memory access generated in user mode (i.e., by a user process) is within the correct [base, base + limit) range for that process




# Implementing Dynamic Relocation

CPU must check every memory access generated in user mode (i.e., by a user process) is within the correct [base, base + limit) range for that process



# Implementing Dynamic Relocation

CPU must check every memory access generated in user mode (i.e., by a user process) is within the correct [base, base + limit) range for that process



# Dynamic Relocation

- PROs:
  - Provides protection (both read and write) across address spaces
  - OS can easily move a process during execution
  - OS can allow process to dynamically grow over time
  - Simple, fast hardware implementation (MMU):
    - 2 special registers, one add and one compare operation (can be done in parallel)

# Dynamic Relocation

- CONs:
  - Little hardware overhead to pay at each memory reference
  - Each VAS of a process must still be allocated contiguously in physical memory (possible memory waste, e.g., stack/heap)  
→ **segmentation**
  - Degree of multiprogramming is bound since all VAS of all active processes must fit entirely in memory → **paging**
  - Process is still limited to physical memory size → **virtual memory**

# To Wrap Up

- Modern OSs manage memory ensuring:
  - Transparency → logical/virtual vs. physical address space
  - Protection/Flexibility → dynamic relocation
  - Efficiency → hardware support (e.g., MMU)
- We are still assuming the whole virtual address space of a process is fully and contiguously loaded in main memory → **serious limitation!**

# Contiguous Memory Allocation

- So far, we have assumed each process is allocated into a contiguous space of physical memory

# Contiguous Memory Allocation

- So far, we have assumed each process is allocated into a contiguous space of physical memory
- One simple method is to divide upfront all available memory dedicated to user processes into **equally-sized** segments/partitions

# Contiguous Memory Allocation

- So far, we have assumed each process is allocated into a contiguous space of physical memory
- One simple method is to divide upfront all available memory dedicated to user processes into **equally-sized** segments/partitions
  - Assign each process to a segment

# Contiguous Memory Allocation

- So far, we have assumed each process is allocated into a contiguous space of physical memory
- One simple method is to divide upfront all available memory dedicated to user processes into **equally-sized** segments/partitions
  - Assign each process to a segment
  - Implicitly restricts the grade of multiprogramming (i.e., the number of simultaneous processes) and their size

# Contiguous Memory Allocation


- So far, we have assumed each process is allocated into a contiguous space of physical memory
- One simple method is to divide upfront all available memory dedicated to user processes into **equally-sized** segments/partitions
  - Assign each process to a segment
  - Implicitly restricts the grade of multiprogramming (i.e., the number of simultaneous processes) and their size
  - No longer used!

# Contiguous Memory Allocation

An alternative approach is for the OS to keep track of **free** (unused) memory segments, as processes enter the system, grow, and terminate


# Contiguous Memory Allocation

An alternative approach is for the OS to keep track of **free** (unused) memory segments, as processes enter the system, grow, and terminate




# Contiguous Memory Allocation

An alternative approach is for the OS to keep track of **free** (unused) memory segments, as processes enter the system, grow, and terminate




# Contiguous Memory Allocation

An alternative approach is for the OS to keep track of **free** (unused) memory segments, as processes enter the system, grow, and terminate




# Contiguous Memory Allocation


An alternative approach is for the OS to keep track of **free** (unused) memory segments, as processes enter the system, grow, and terminate




# Contiguous Memory Allocation



# Contiguous Memory Allocation



# Contiguous Memory Allocation




Unused segments of memory are called **holes**

The OS keeps a list of holes

Whenever a process has to be loaded, the OS must select a hole of suitable size

# Contiguous Memory Allocation



Unused segments of memory are called **holes**

The OS keeps a list of holes

Whenever a process has to be loaded, the OS must select a hole of suitable size

How?

# Memory Allocation Policies: First-Fit

- Linearly scan the list of holes until one is found that is big enough to satisfy the request

# Memory Allocation Policies: First-Fit

- Linearly scan the list of holes until one is found that is big enough to satisfy the request
- Subsequent requests may either start from the beginning of the list or from the end of previous search

# Memory Allocation Policies: First-Fit

- Linearly scan the list of holes until one is found that is big enough to satisfy the request
- Subsequent requests may either start from the beginning of the list or from the end of previous search
- Complexity:  $O(n)$ , where  $n$  is the number of holes


# Memory Allocation Policies: First-Fit

Suppose process X needs 100B of memory to be loaded, and the list of holes is as follows:



# Memory Allocation Policies: First-Fit


Suppose process X needs 100B of memory to be loaded, and the list of holes is as follows:



Which segment will be picked by the OS using first-fit?

# Memory Allocation Policies: First-Fit


Suppose process X needs 100B of memory to be loaded, and the list of holes is as follows:




Which segment will be picked by the OS using first-fit?

# Memory Allocation Policies: First-Fit

Suppose process X needs 100B of memory to be loaded, and the list of holes is as follows:




Which segment will be picked by the OS using first-fit?




# Memory Allocation Policies: First-Fit

Suppose process X needs 100B of memory to be loaded, and the list of holes is as follows:



Which segment will be picked by the OS using first-fit?




What if afterwards process Y requires 350B?

# Memory Allocation Policies: First-Fit

Suppose process X needs 100B of memory to be loaded, and the list of holes is as follows:



Which segment will be picked by the OS using first-fit?



What if afterwards process Y requires 350B?

We will not be able to satisfy this request even if theoretically we could

# Memory Allocation Policies: Best-Fit

- Allocate the smallest hole that is big enough to satisfy the request

# Memory Allocation Policies: Best-Fit

- Allocate the smallest hole that is big enough to satisfy the request
- This saves large holes for other process requests that may need them

# Memory Allocation Policies: Best-Fit

- Allocate the smallest hole that is big enough to satisfy the request
- This saves large holes for other process requests that may need them
- However, the resulting unused portions of holes may be too small to be of any use, and will therefore be wasted

# Memory Allocation Policies: Best-Fit

- Allocate the smallest hole that is big enough to satisfy the request
- This saves large holes for other process requests that may need them
- However, the resulting unused portions of holes may be too small to be of any use, and will therefore be wasted
- **Complexity:** still  $O(n)$  but can be  $O(\log n)$  if the list of holes is kept sorted

# Memory Allocation Policies: Best-Fit

- Allocate the smallest hole that is big enough to satisfy the request
- This saves large holes for other process requests that may need them
- However, the resulting unused portions of holes may be too small to be of any use, and will therefore be wasted
- **Complexity:** still  $O(n)$  but can be  $O(\log n)$  if the list of holes is kept sorted

Do you know how which data structure can be used to achieve this?

# Memory Allocation Policies: Best-Fit

- Allocate the smallest hole that is big enough to satisfy the request
- This saves large holes for other process requests that may need them
- However, the resulting unused portions of holes may be too small to be of any use, and will therefore be wasted
- **Complexity:** still  $O(n)$  but can be  $O(\log n)$  if the list of holes is kept sorted

Do you know how which data structure can be used to achieve this?

Binary Search Tree (BST)

# Memory Allocation Policies: Best-Fit

Suppose process X needs 100B of memory to be loaded, and the list of holes is as follows:



# Memory Allocation Policies: Best-Fit

Suppose process X needs 100B of memory to be loaded, and the list of holes is as follows:



Which segment will be picked by the OS using best-fit?

# Memory Allocation Policies: Best-Fit


Suppose process X needs 100B of memory to be loaded, and the list of holes is as follows:




Which segment will be picked by the OS using best-fit?

# Memory Allocation Policies: Best-Fit

Suppose process X needs 100B of memory to be loaded, and the list of holes is as follows:



Which segment will be picked by the OS using best-fit?




# Memory Allocation Policies: Best-Fit

Suppose process X needs 100B of memory to be loaded, and the list of holes is as follows:



Which segment will be picked by the OS using best-fit?




What if afterwards process Y requires 350B?

# Memory Allocation Policies: Best-Fit

Suppose process X needs 100B of memory to be loaded, and the list of holes is as follows:



Which segment will be picked by the OS using best-fit?



What if afterwards process Y requires 350B?

We can now assign it the second available hole segment (400B)

# Memory Allocation Policies: Worst-Fit

- Allocate the largest hole available

# Memory Allocation Policies: Worst-Fit

- Allocate the largest hole available
- Might sound counterintuitive but this increases the likelihood that the remaining portion will be usable for satisfying future requests

# Memory Allocation Policies: Worst-Fit

- Allocate the largest hole available
- Might sound counterintuitive but this increases the likelihood that the remaining portion will be usable for satisfying future requests
- Simulations show that First-Fit and Best-Fit usually work best

# Memory Allocation Policies: Worst-Fit

- Allocate the largest hole available
- Might sound counterintuitive but this increases the likelihood that the remaining portion will be usable for satisfying future requests
- Simulations show that First-Fit and Best-Fit usually work best
- First-Fit is also generally faster than Best-Fit

# Fragmentation

## Problem


Individual holes may be too small to serve a process request but they can be large enough if combined together

# Fragmentation

## Problem

Individual holes may be too small to serve a process request but they can be large enough if combined together

External  
Fragmentation



# Fragmentation

## Problem

Individual holes may be too small to serve a process request but they can be large enough if combined together

External  
Fragmentation

Internal  
Fragmentation

# External Fragmentation


- Frequent loading and unloading processes causes holes to be broken into small (i.e., unusable) chunks

# External Fragmentation

- Frequent loading and unloading processes causes holes to be broken into small (i.e., unusable) chunks
- It happens when there is enough memory to load a process in memory but space is not contiguous

# External Fragmentation

- Frequent loading and unloading processes causes holes to be broken into small (i.e., unusable) chunks
- It happens when there is enough memory to load a process in memory but space is not contiguous




# External Fragmentation

- Frequent loading and unloading processes causes holes to be broken into small (i.e., unusable) chunks
- It happens when there is enough memory to load a process in memory but space is not contiguous

Simulations show that for every  $2N$  allocated blocks,  $N$  are lost due to external fragmentation

1/3 of memory space is wasted on average




# External Fragmentation

- Frequent loading and unloading processes causes holes to be broken into small (i.e., unusable) chunks
- It happens when there is enough memory to load a process in memory but space is not contiguous

Simulations show that for every  $2N$  allocated blocks,  $N$  are lost due to external fragmentation

1/3 of memory space is wasted on average

Goal:  
Allocation policy that minimizes wasted space!



# Internal Fragmentation

- It happens when memory internal to a segment is wasted


# Internal Fragmentation

- It happens when memory internal to a segment is wasted
- For example, consider a process whose size is 8,846B and a hole of size 8,848B


# Internal Fragmentation

- It happens when memory internal to a segment is wasted
- For example, consider a process whose size is 8,846B and a hole of size 8,848B
- It may be much more efficient to allocate the process the whole block (and waste 2B) rather than keep track of a tiny 2B hole


# Solution to Fragmentation: Full Compaction




# Solution to Fragmentation: Full Compaction



# Solution to Fragmentation: Full Compaction




# Solution to Fragmentation: Full Compaction




Only one hole is left but two processes need to be moved (C and D)


# Solution to Fragmentation: Partial Compaction




# Solution to Fragmentation: Partial Compaction



# Solution to Fragmentation: Partial Compaction



# Solution to Fragmentation: Partial Compaction



Still some holes left but only one process is moved (D) rather than two

# Swapping

- So far, we have assumed all processes are entirely loaded in memory (of course, when loaded!)

# Swapping

- So far, we have assumed all processes are entirely loaded in memory (of course, when loaded!)
- **Remember:** A process needs to sit physically in main memory only if the CPU executes its instructions and accesses its data

# Swapping

- So far, we have assumed all processes are entirely loaded in memory (of course, when loaded!)
- **Remember:** A process needs to sit physically in main memory only if the CPU executes its instructions and accesses its data
- If a process blocks (e.g., due to an I/O call) it doesn't need to be in memory while I/O is running

# Swapping

- So far, we have assumed all processes are entirely loaded in memory (of course, when loaded!)
- **Remember:** A process needs to sit physically in main memory only if the CPU executes its instructions and accesses its data
- If a process blocks (e.g., due to an I/O call) it doesn't need to be in memory while I/O is running
- That process can be "swapped out" from memory to disk to make room for other processes

# Swapping

- Once process becomes ready again, the OS must reload it in memory

# Swapping

- Once process becomes ready again, the OS must reload it in memory
- Swap in depends on the address binding used:
  - **compile-** or **load-time**: must be swapped back into the same memory location from which they were swapped out

# Swapping

- Once process becomes ready again, the OS must reload it in memory
- Swap in depends on the address binding used:
  - **compile-** or **load-time**: must be swapped back into the same memory location from which they were swapped out
  - **execution-time**: can be swapped back into any available location (updating base and limit registers)

# Swapping

- Once process becomes ready again, the OS must reload it in memory
- Swap in depends on the address binding used:
  - **compile-** or **load-time**: must be swapped back into the same memory location from which they were swapped out
  - **execution-time**: can be swapped back into any available location (updating base and limit registers)
- Using swapping, fragmentation can be tackled easily
  - Just run compaction before swapping-in a process

# Swapping: Example

- Swapping is a very slow process compared to other operations due to the interaction with hard disk

# Swapping: Example

- Swapping is a very slow process compared to other operations due to the interaction with hard disk
- Example:
  - 10 MB user process
  - disk transfer rate = 40 MB/sec (250 msec just to do the data transfer)


# Swapping: Example

- Swapping is a very slow process compared to other operations due to the interaction with hard disk
- Example:
  - 10 MB user process
  - disk transfer rate = 40 MB/sec (250 msec just to do the data transfer)
- Since swap-in may involve swapping-out another process, the overall time required will be ~500 msec

# Swapping: Example

- Swapping is a very slow process compared to other operations due to the interaction with hard disk
- Example:
  - 10 MB user process
  - disk transfer rate = 40 MB/sec (250 msec just to do the data transfer)
- Since swap-in may involve swapping-out another process, the overall time required will be ~500 msec
- Time slice is usually way smaller than that!

# Swapping



Most modern OSs no longer use swapping, because it is too slow and there are faster alternatives available (e.g., [paging](#))

# Summary

- Contiguous memory allocation may cause fragmentation
  - External vs. Internal Fragmentation
- Existing countermeasures (compaction) exist but they are costly
- We may want to relax the constraint on having an entire process loaded in main memory
- Paging solves all these issues!