

# Sistemi Operativi

Corso di Laurea in Informatica

a.a. 2019-2020



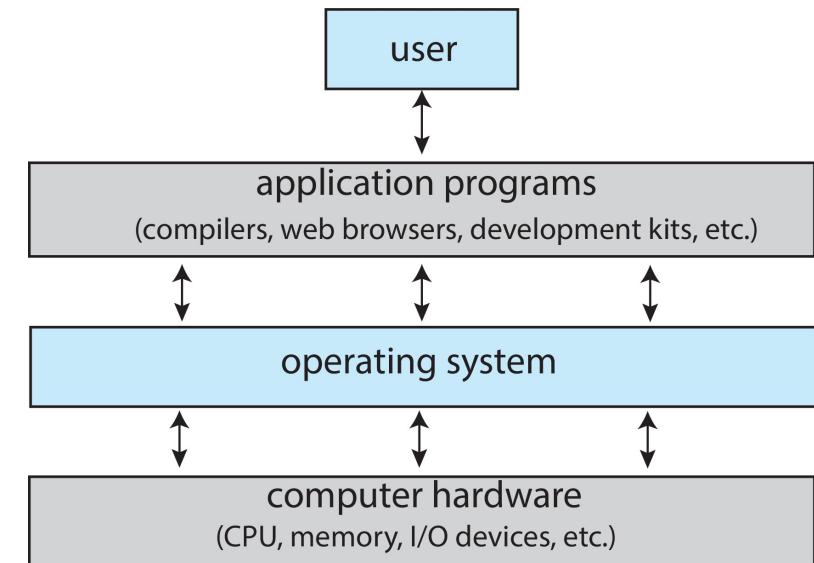
**SAPIENZA**  
UNIVERSITÀ DI ROMA

Gabriele Tolomei

Dipartimento di Informatica  
Sapienza Università di Roma  
[tolomei@di.uniroma1.it](mailto:tolomei@di.uniroma1.it)

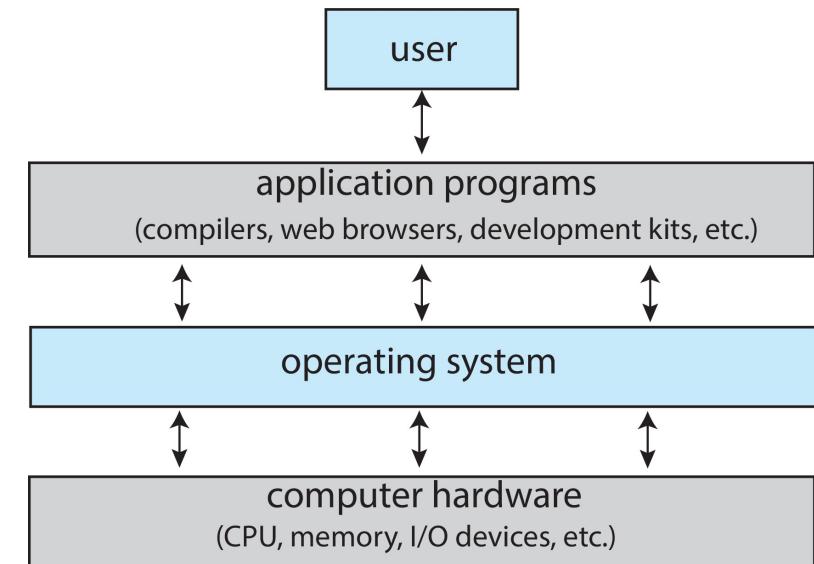
# Recap from Last Lecture

- Operating System is a complex system which plays several roles:
  - resource manager
  - virtual machine
  - HW/SW interface



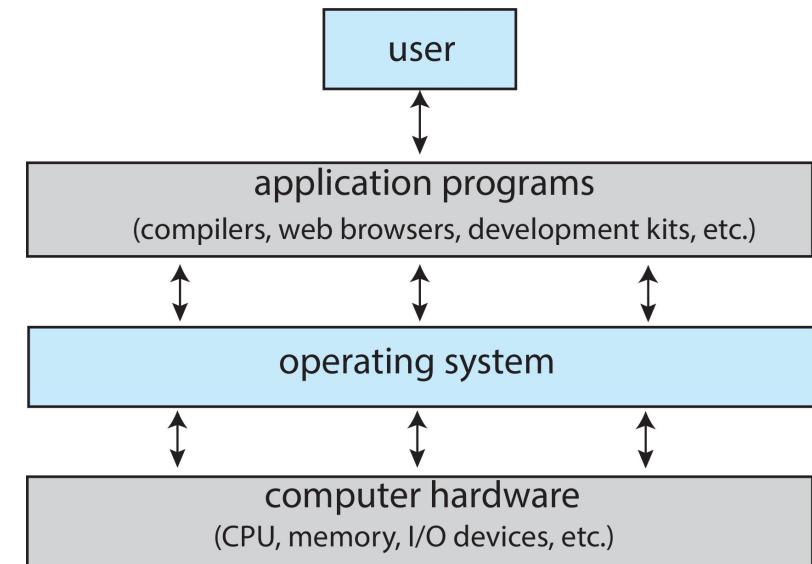
# Recap from Last Lecture

- Operating System is a complex system which plays several roles:
  - resource manager
  - virtual machine
  - HW/SW interface
- Exposes services to users/applications (SW) leveraging the physical machine (HW)



# Recap from Last Lecture

- Operating System is a complex system which plays several roles:
  - resource manager
  - virtual machine
  - HW/SW interface
- Exposes services to users/applications (SW) leveraging the physical machine (HW)
- Changes in HW may affect OS design



# OS and Computer Architecture

- Computer architecture review

# OS and Computer Architecture

- Computer architecture review
- Basic OS functionalities (enabled by architectural features)

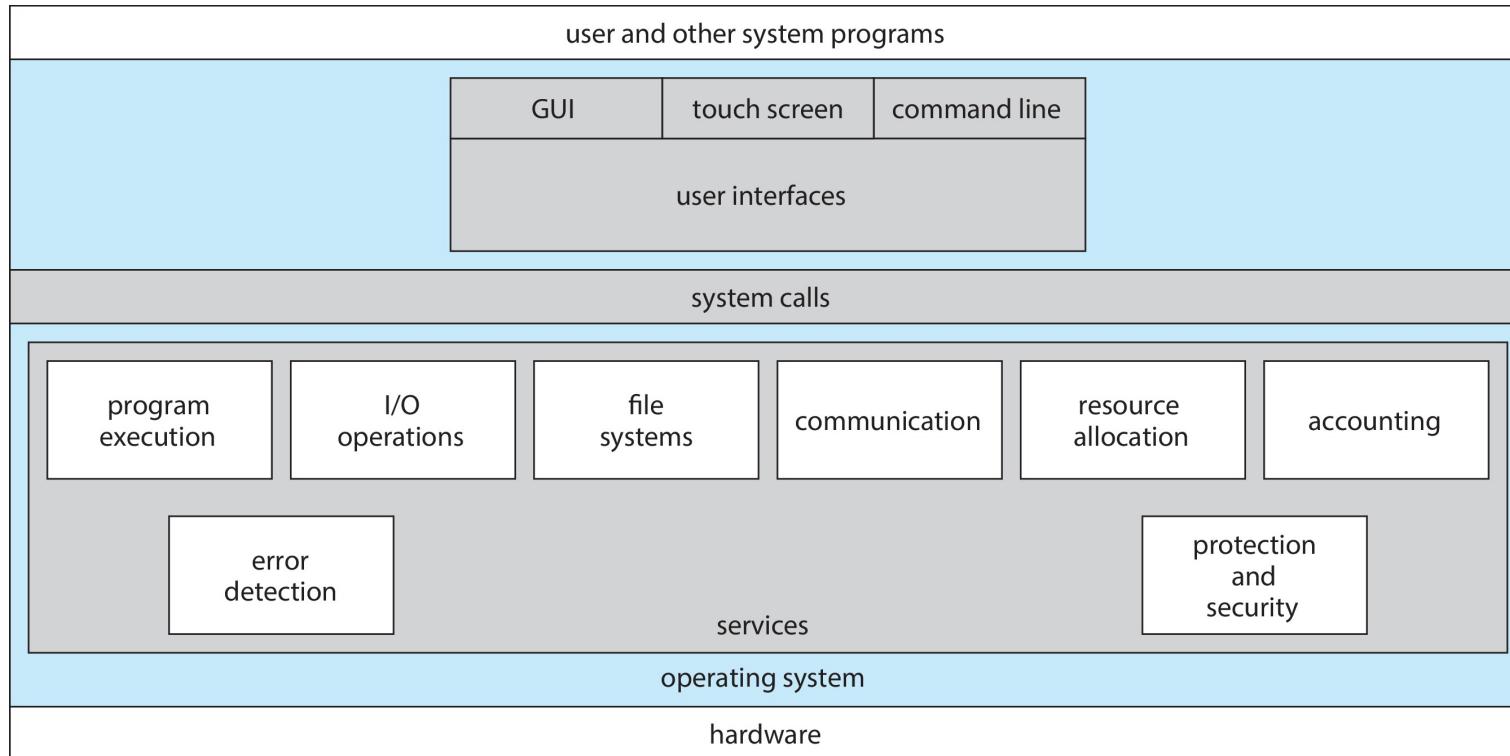
# OS and Computer Architecture

- Computer architecture review
- Basic OS functionalities (enabled by architectural features)
- What the OS can do is partially dictated by the underlying architecture

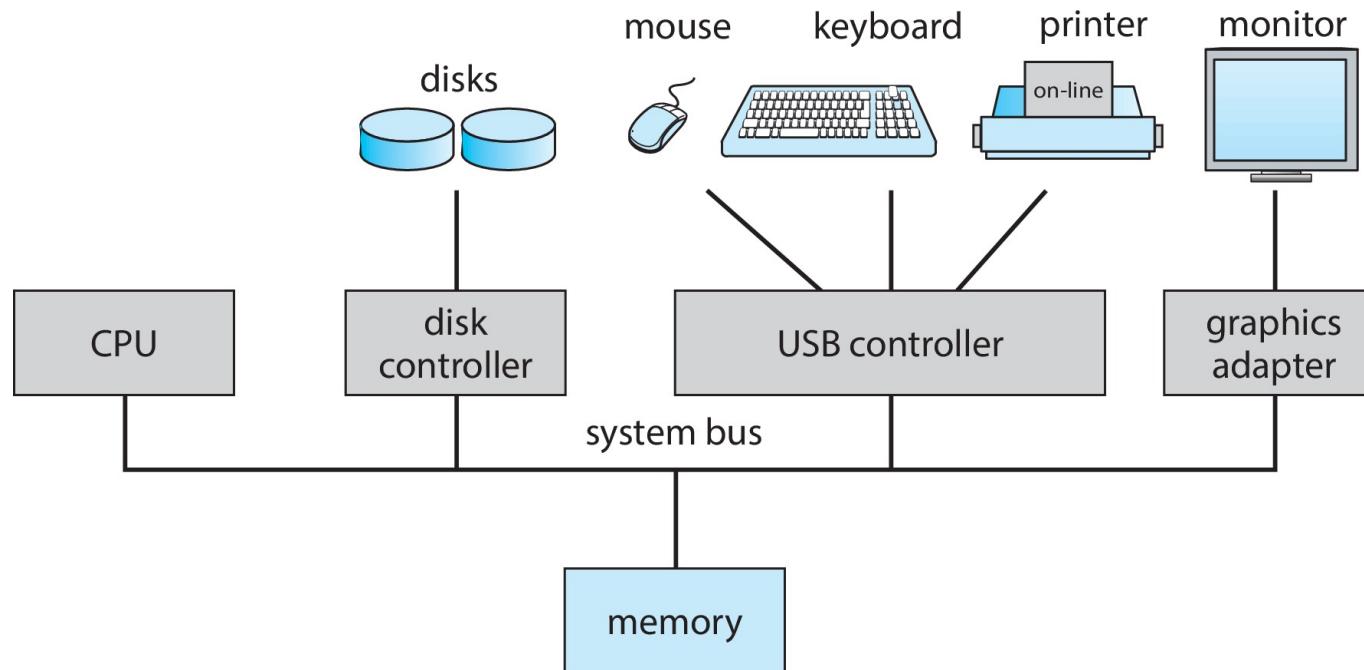
# OS and Computer Architecture

- Computer architecture review
- Basic OS functionalities (enabled by architectural features)
- What the OS can do is partially dictated by the underlying architecture
- Architectural support may significantly simplify or complicate the OS design

# Modern OS Functionalities



# Generic High-Level Computer Architecture



# Generic High-Level Computer Architecture

- **CPU** → the processor that performs the actual computation
  - Multiple cores are now common in modern architectures

# Generic High-Level Computer Architecture

- **CPU** → the processor that performs the actual computation
  - Multiple cores are now common in modern architectures
- **Main Memory** → stores data and instructions used by the CPU

# Generic High-Level Computer Architecture

- **CPU** → the processor that performs the actual computation
  - Multiple cores are now common in modern architectures
- **Main Memory** → stores data and instructions used by the CPU
- **I/O devices** → terminal, keyboard, disks, etc.
  - associated with specific device controllers

# Generic High-Level Computer Architecture

- **CPU** → the processor that performs the actual computation
  - Multiple cores are now common in modern architectures
- **Main Memory** → stores data and instructions used by the CPU
- **I/O devices** → terminal, keyboard, disks, etc.
  - associated with specific device controllers
- **System Bus** → communication medium between CPU, memory, and peripherals

# Computer Architecture Model

- Conceptually, the same architectural model for many computing devices:
  - PCs/laptops
  - High-end servers
  - Smartphones/Tablets
  - etc.

# Computer Architecture Model

- Conceptually, the same architectural model for many computing devices:
  - PCs/laptops
  - High-end servers
  - Smartphones/Tablets
  - etc.
- Based on **stored-program** concept (as opposed to fixed-program)

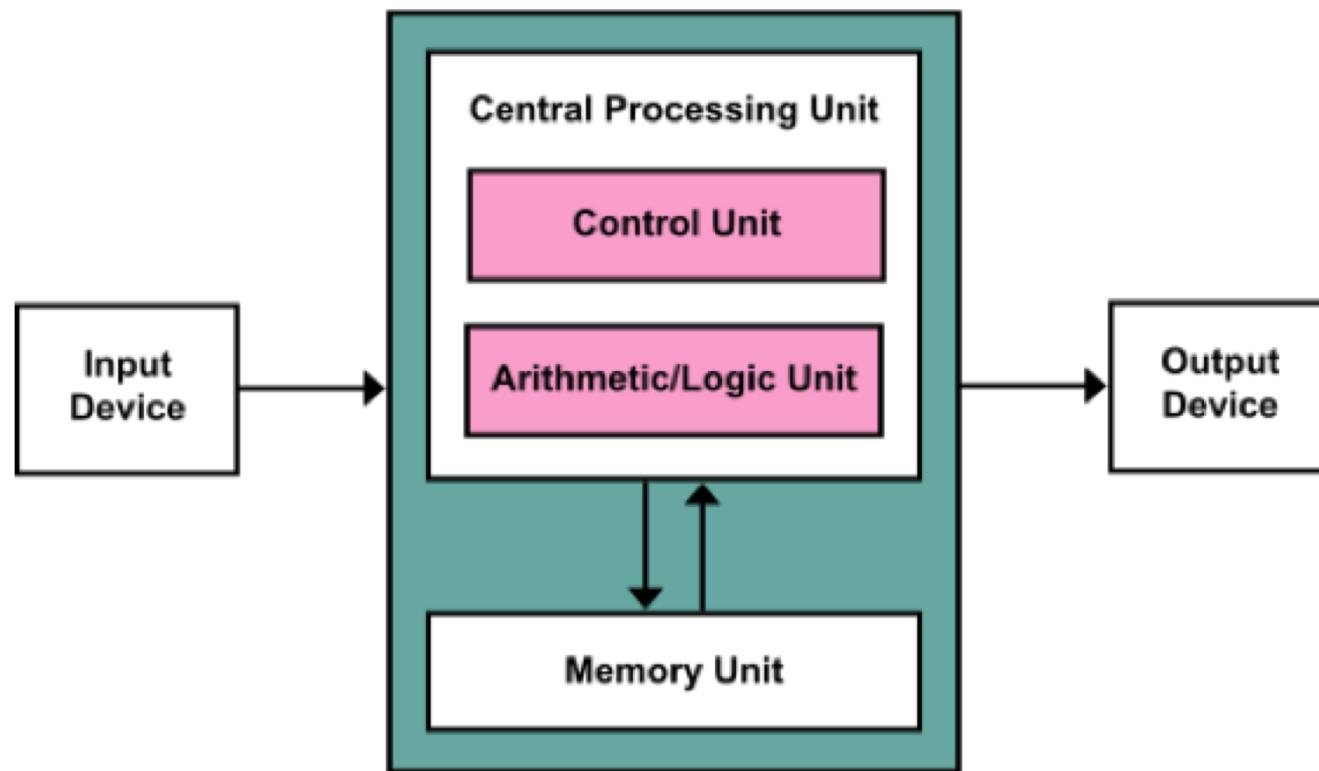
# Computer Architecture Model

- Conceptually, the same architectural model for many computing devices:
  - PCs/laptops
  - High-end servers
  - Smartphones/Tablets
  - etc.
- Based on **stored-program** concept (as opposed to fixed-program)

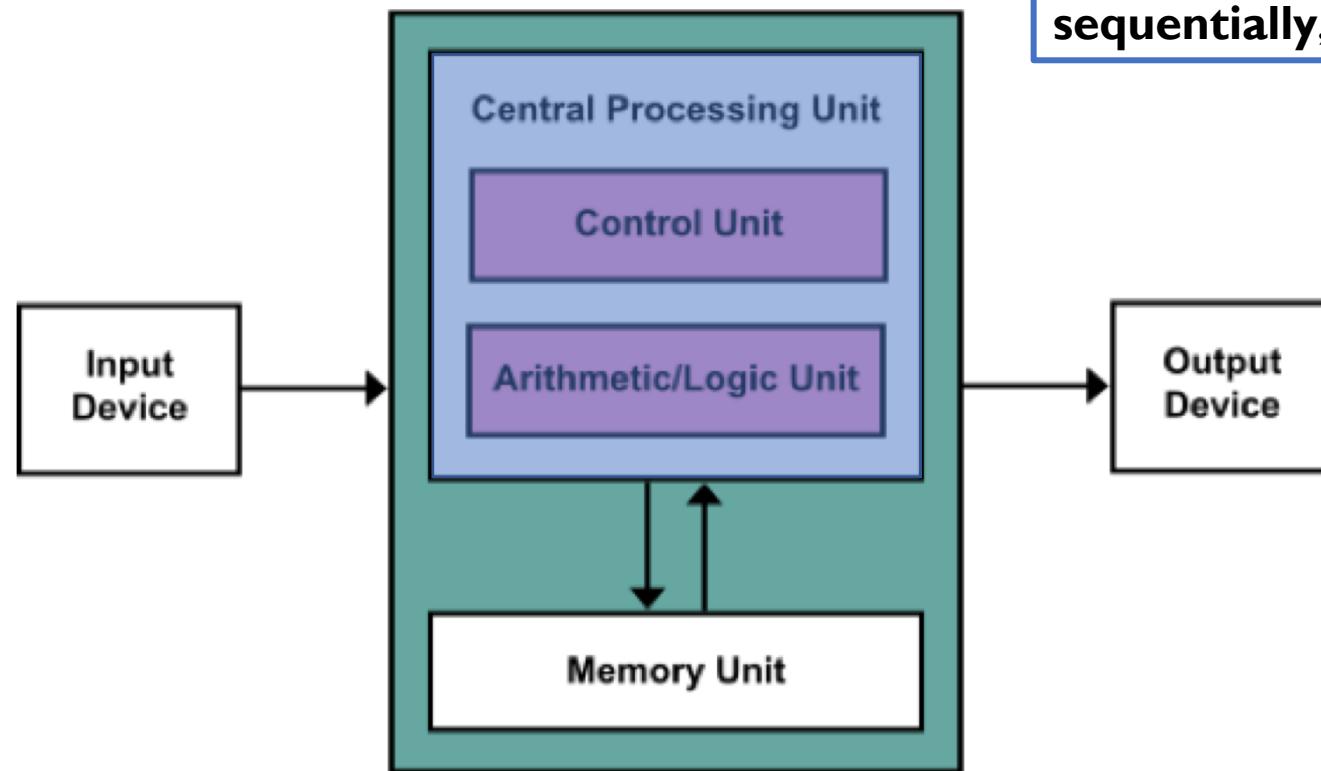


John von Neumann

# von Neumann Architecture

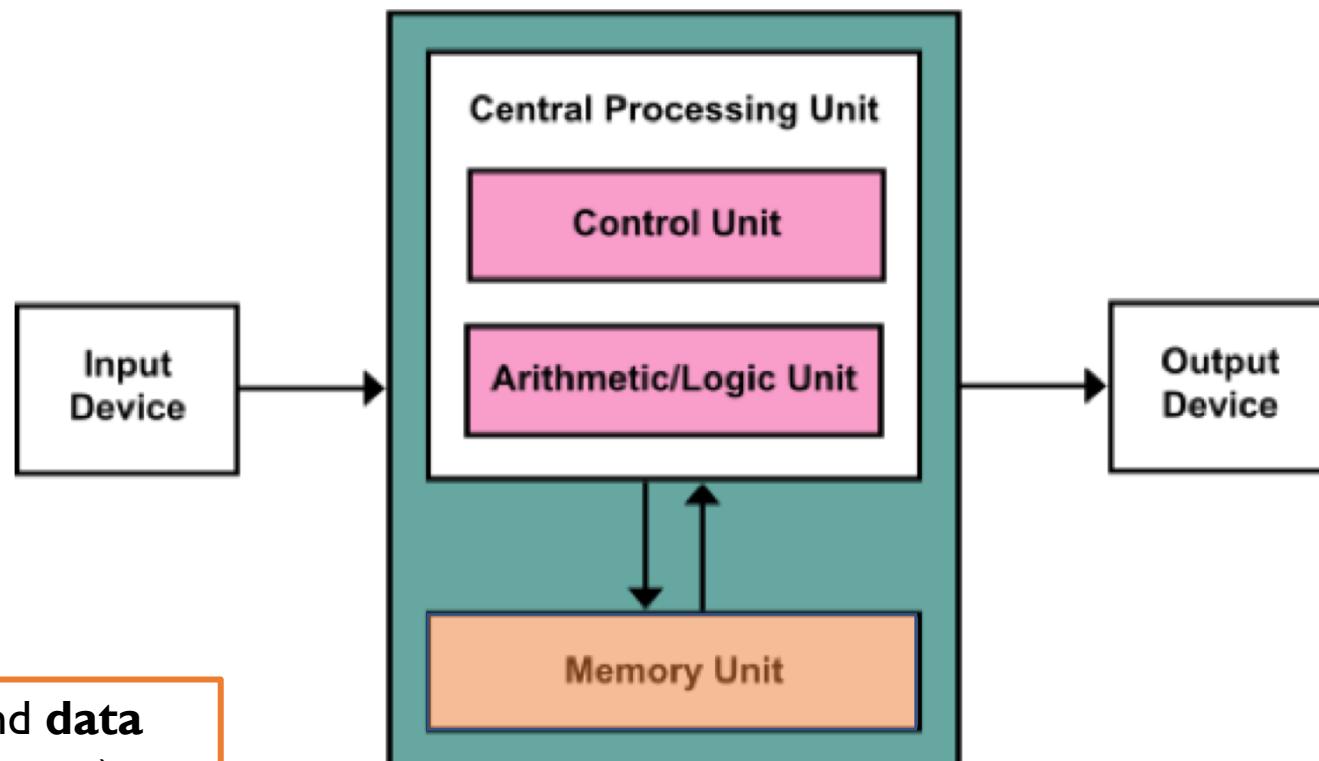


# von Neumann Architecture



Executes **instructions** (i.e., basic arithmetical/logical computations) **sequentially**, using internal **registers**

# von Neumann Architecture



# Central Processing Unit (CPU)

# Instruction Cycle: Fetch-Decode-Execute

- The CPU performs the following **3 steps** cyclically:

# Instruction Cycle: Fetch-Decode-Execute

- The CPU performs the following **3 steps** cyclically:
  - **Fetch:** retrieves an instruction from a specific memory address whose value is contained in a special CPU register, called Program Counter

# Instruction Cycle: Fetch-Decode-Execute

- The CPU performs the following **3 steps** cyclically:
  - **Fetch:** retrieves an instruction from a specific memory address whose value is contained in a special CPU register, called Program Counter
  - **Decode:** interprets the fetched instruction

# Instruction Cycle: Fetch-Decode-Execute

- The CPU performs the following **3 steps** cyclically:
  - **Fetch:** retrieves an instruction from a specific memory address whose value is contained in a special CPU register, called **Program Counter**
  - **Decode:** interprets the fetched instruction
  - **Execute:** runs the actual decoded instruction

# Machine Language

- Defines the set of (elementary) instructions the CPU is able to execute directly (i.e., on the HW)

# Machine Language

- Defines the set of (elementary) instructions the CPU is able to execute directly (i.e., on the HW)
- The machine language is represented by binary numeral system

# Machine Language

- Defines the set of (elementary) instructions the CPU is able to execute directly (i.e., on the HW)
- The machine language is represented by binary numeral system
- Each instruction is encoded as a sequence of bits
  - A single bit is the smallest unit of (digital) information
  - It takes on two possible values: 0 or 1

# Machine Language

- Defines the set of (elementary) instructions the CPU is able to execute directly (i.e., on the HW)
- The machine language is represented by binary numeral system
- Each instruction is encoded as a sequence of bits
  - A single bit is the smallest unit of (digital) information
  - It takes on two possible values: 0 or 1
- A **word** is the unit of data the CPU can directly operate on
  - today ranging from 32 to 64 bits

# Binary vs. Decimal Numeral System

- In natural language we usually implicitly refer to the decimal numeral system (base-10)

# Binary vs. Decimal Numeral System

- In natural language we usually implicitly refer to the decimal numeral system (base-10)
- In base-10 system, each digit can only take one out of 10 possible values: 0, 1, ..., 9

1	0	1
$10^2$	$10^1$	$10^0$

$$1*10^0 + 0*10^1 + 1*10^2 = 101$$

# Binary vs. Decimal Numeral System

- In natural language we usually implicitly refer to the decimal numeral system (base-10)
- In base-10 system, each digit can only take one out of 10 possible values: 0, 1, ..., 9

1	0	1
$10^2$	$10^1$	$10^0$

$$1 * 10^0 + 0 * 10^1 + 1 * 10^2 = 101$$

- In binary system (base-2), each digit is a bit

1	0	1
$2^2$	$2^1$	$2^0$

$$1 * 2^0 + 0 * 2^1 + 1 * 2^2 = 5$$

# A Side Note on Units

Prefixes for multiples of bits (bit) or bytes (B)					
Decimal		Binary			
Value	SI	Value	IEC	JEDEC	
1000	$10^3$ k kilo	1024	$2^{10}$ Ki kibi	K kilo	
$1000^2$	$10^6$ M mega	$1024^2$	$2^{20}$ Mi mebi	M mega	
$1000^3$	$10^9$ G giga	$1024^3$	$2^{30}$ Gi gibi	G giga	
$1000^4$	$10^{12}$ T tera	$1024^4$	$2^{40}$ Ti tebi	–	
$1000^5$	$10^{15}$ P peta	$1024^5$	$2^{50}$ Pi pebi	–	
$1000^6$	$10^{18}$ E exa	$1024^6$	$2^{60}$ Ei exbi	–	
$1000^7$	$10^{21}$ Z zetta	$1024^7$	$2^{70}$ Zi zebi	–	
$1000^8$	$10^{24}$ Y yotta	$1024^8$	$2^{80}$ Yi yobi	–	

# Instruction Set (Architecture)

- The collection of instructions defined by the machine language

# Instruction Set (Architecture)

- The collection of instructions defined by the machine language
- Machine language instructions are made of 2 parts:
  - An **operator** (op code)
  - Zero or more **operands** representing either CPU internal registers or memory addresses

# Instruction Set (Architecture)

- The collection of instructions defined by the machine language
- Machine language instructions are made of 2 parts:
  - An **operator** (op code)
  - Zero or more **operands** representing either CPU internal registers or memory addresses
- An abstraction of the underlying physical (hardware) architecture
  - x86, ARM, SPARC, MIPS, etc.

# Instruction Set (Architecture)

- The collection of instructions defined by the machine language
- Machine language instructions are made of 2 parts:
  - An **operator** (op code)
  - Zero or more **operands** representing either CPU internal registers or memory addresses
- An abstraction of the underlying physical (hardware) architecture
  - x86, ARM, SPARC, MIPS, etc.
- Each realization of the same instruction set is an implementation of a physical architecture
  - x86 → Intel, AMD, Cyrix, etc.

# CPU Registers

- On-chip storage whose size typically coincides with the CPU word size

# CPU Registers

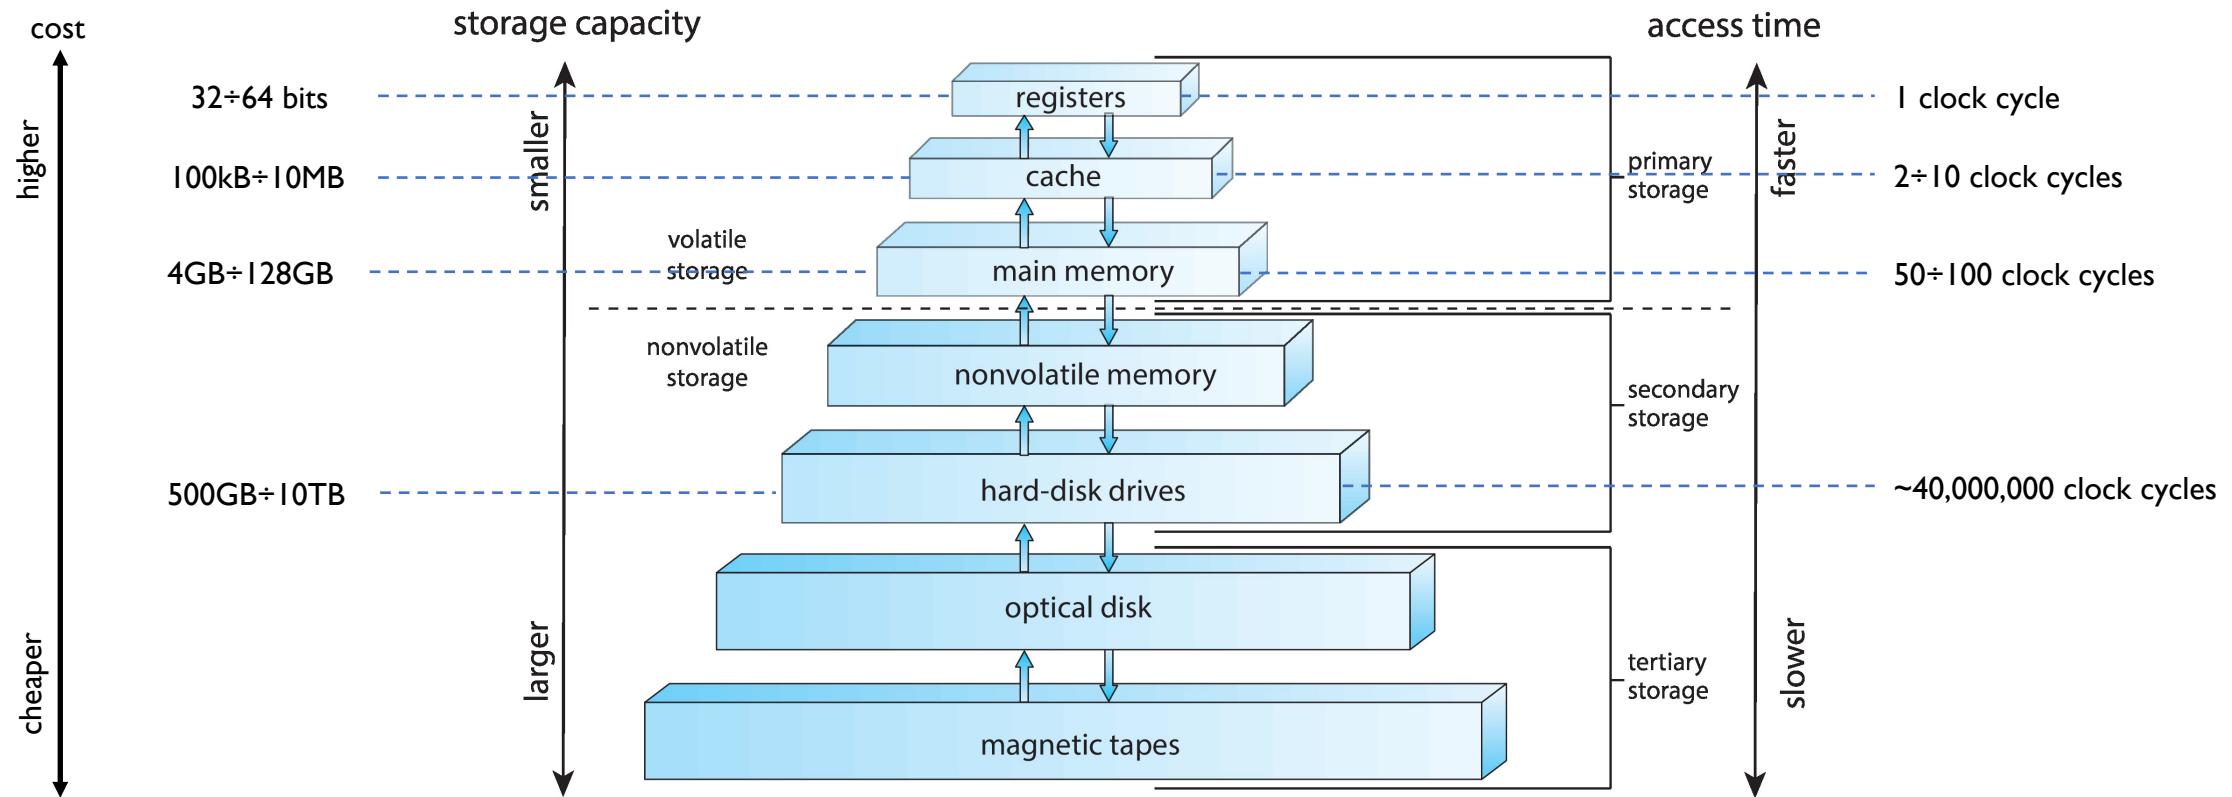
- On-chip storage whose size typically coincides with the CPU word size
- General-purpose (x86):
  - eax, ebx, ecx, etc.

# CPU Registers

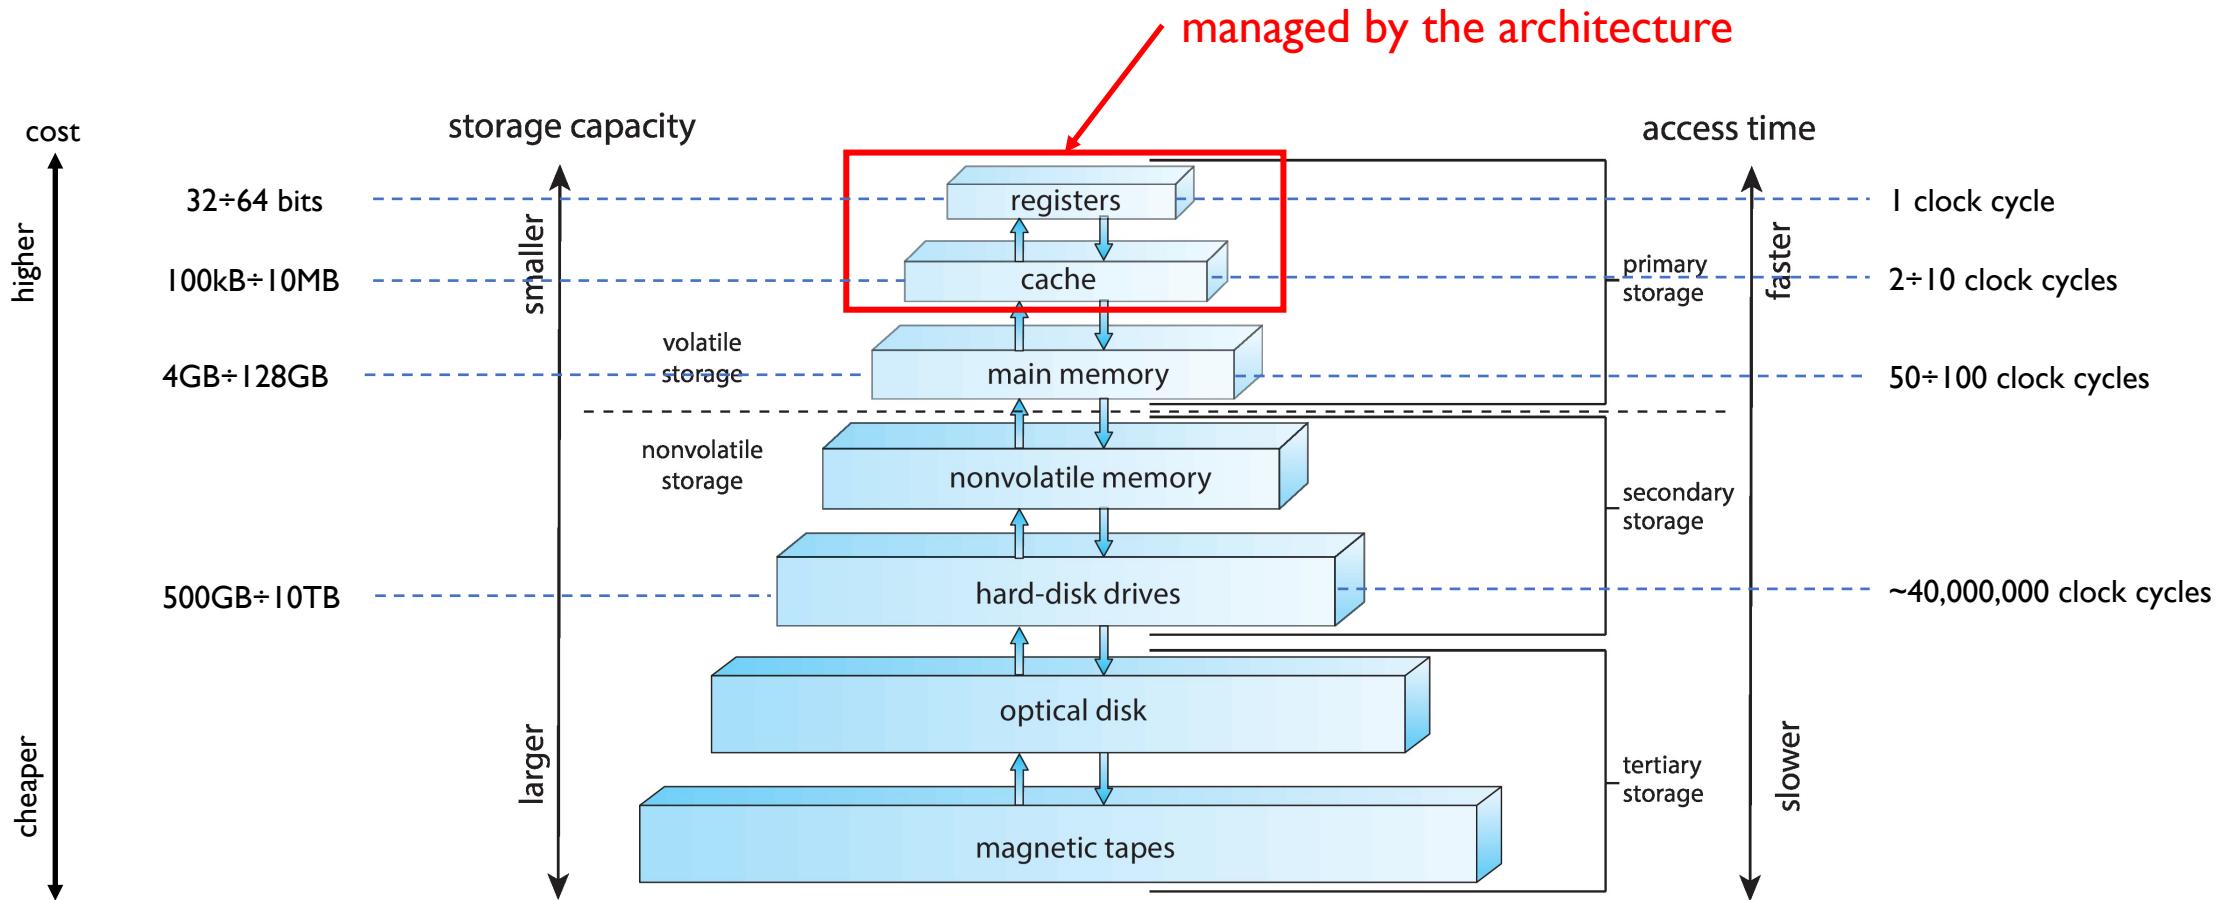
- On-chip storage whose size typically coincides with the CPU word size
- General-purpose (x86):
  - eax, ebx, ecx, etc.
- Special-purpose (x86):
  - esp → Stack pointer for top address of the stack
  - ebp → Stack base pointer for the address of the current stack frame
  - eip → Instruction pointer, holds the program counter (i.e., the address of next instruction)

# Memory

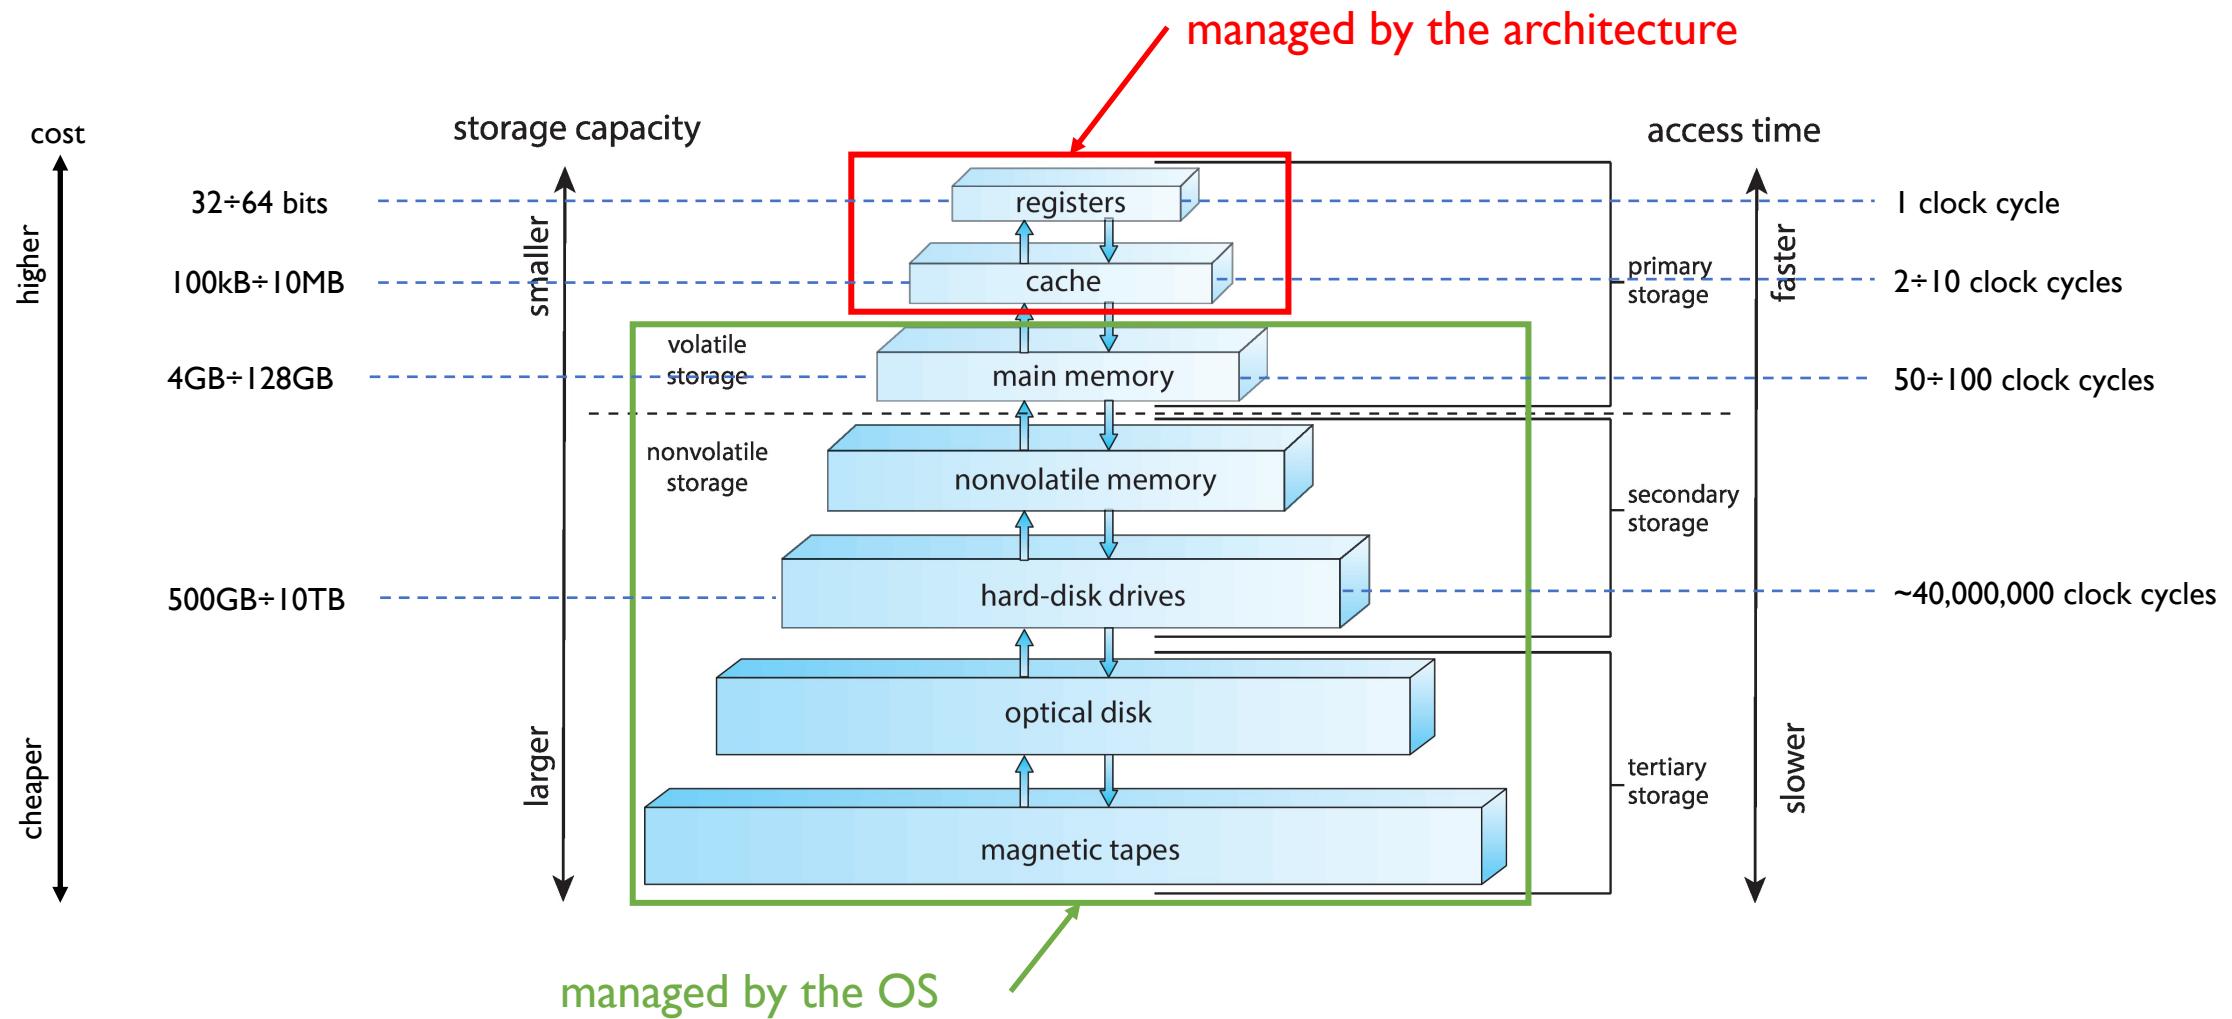
# Memory Hierarchy



# Memory Hierarchy



# Memory Hierarchy



# Main Memory Representation

- Essentially, a sequence of cells

# Main Memory Representation

- Essentially, a sequence of cells
- Each cell is logically organized in groups of 8 bits (1 Byte) or multiple of it (e.g., 32 bits = 4 Bytes)

# Main Memory Representation

- Essentially, a sequence of cells
- Each cell is logically organized in groups of 8 bits (1 Byte) or multiple of it (e.g., 32 bits = 4 Bytes)
- Each cell has its own address

# Main Memory Representation

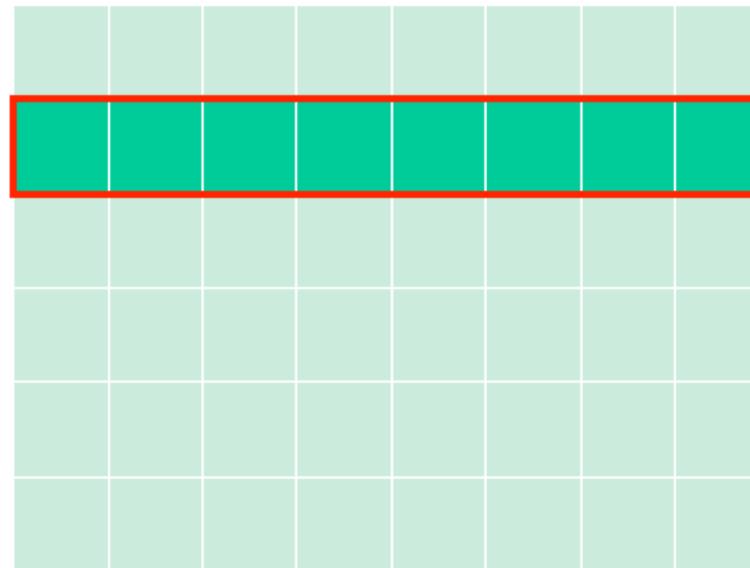
- Essentially, a sequence of cells
- Each cell is logically organized in groups of 8 bits (1 Byte) or multiple of it (e.g., 32 bits = 4 Bytes)
- Each cell has its own address
- CPU (and I/O devices) read from/write to main memory referencing memory location addresses

# Main Memory Representation

- Essentially, a sequence of cells
- Each cell is logically organized in groups of 8 bits (1 Byte) or multiple of it (e.g., 32 bits = 4 Bytes)
- Each cell has its own address
- CPU (and I/O devices) read from/write to main memory referencing memory location addresses
- The smallest addressable unit is usually 1 Byte

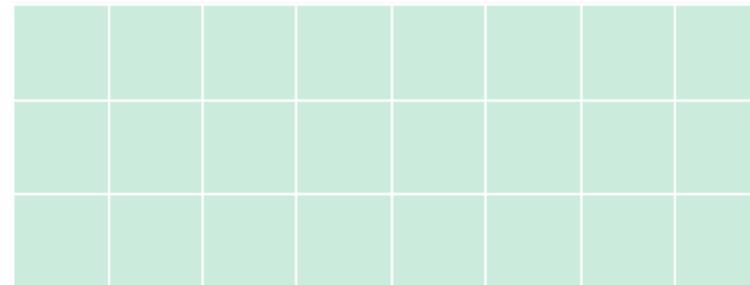
# Memory Cell (I)

Cell/Location



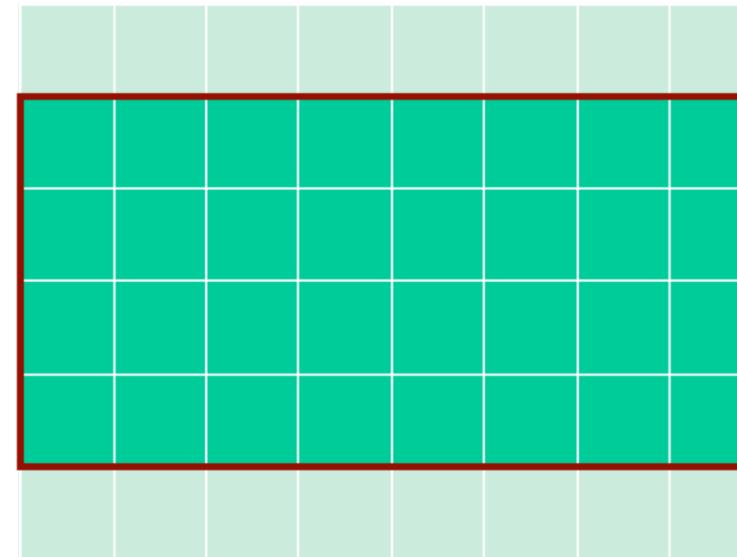
8 bits = 1 Byte

...



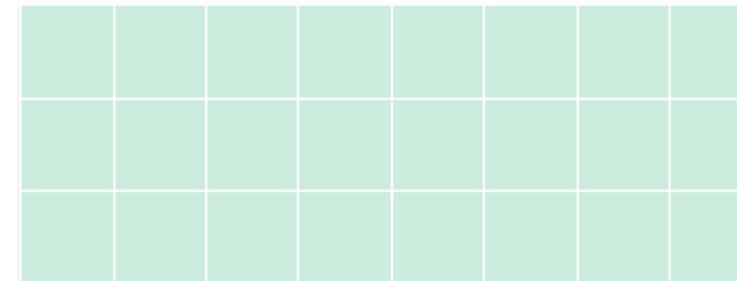
# Memory Cell (2)

Cell/Location

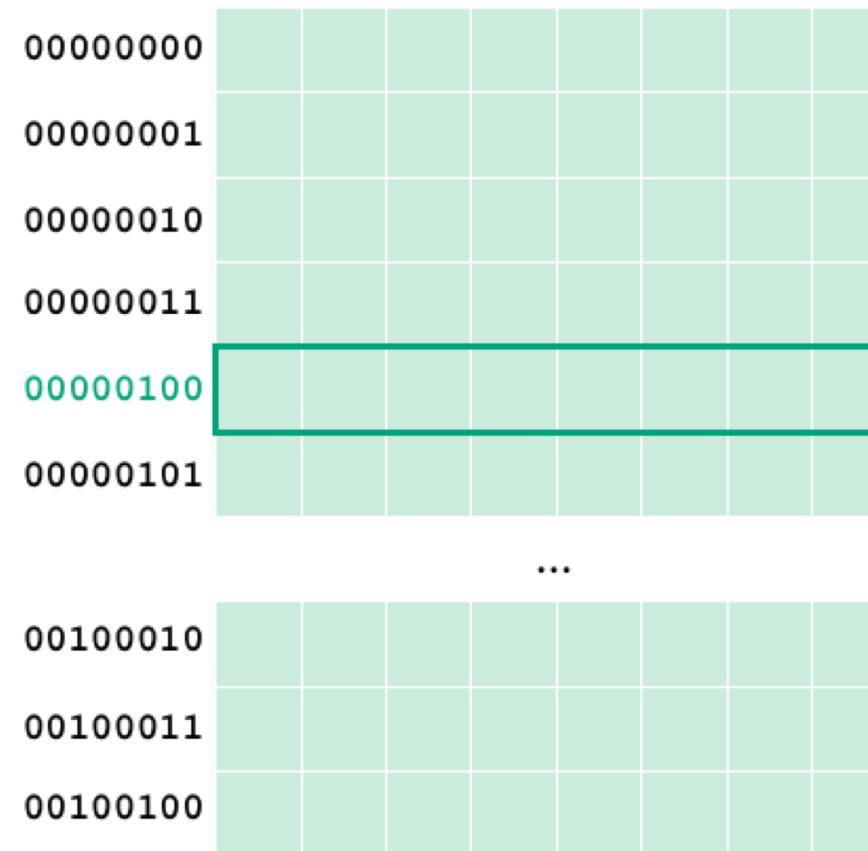


32 bits = 4 Bytes

...



# Memory Address (Single Byte)



# Computer Buses

# System Bus

- Initially, a single bus to handle all the traffic

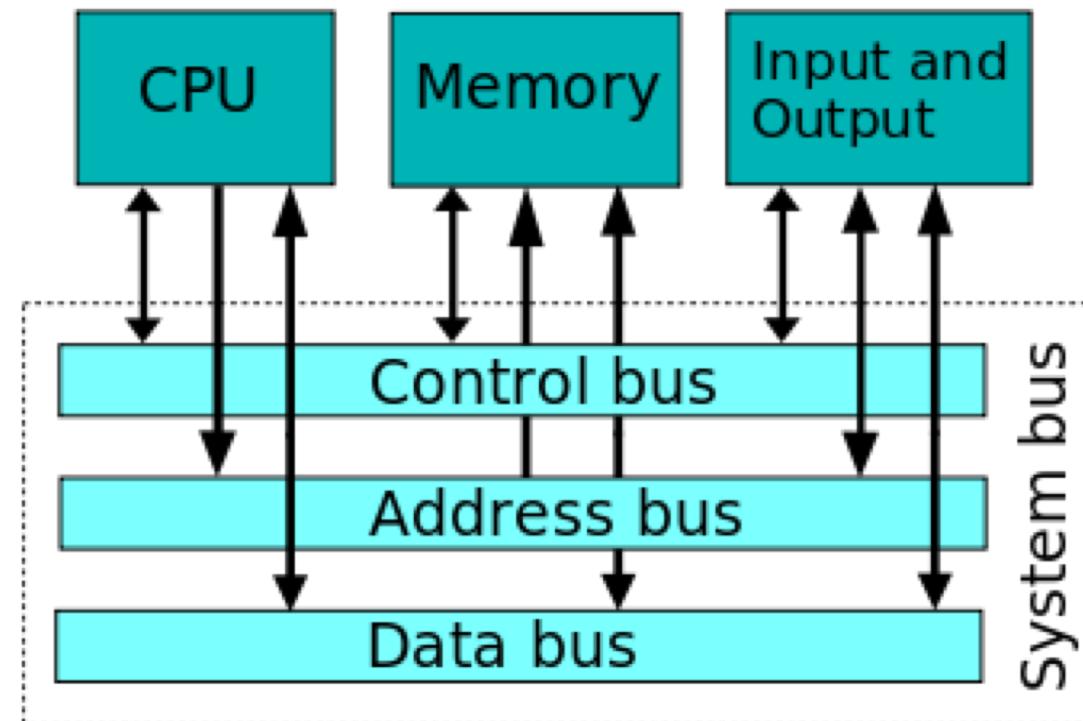
# System Bus

- Initially, a single bus to handle all the traffic
- Combines the functions of:
  - **Data bus** → to actually carry information
  - **Address bus** → to determine where such information should be sent
  - **Control bus** → to indicate which operation should be performed

# System Bus

- Initially, a single bus to handle all the traffic
- Combines the functions of:
  - **Data bus** → to actually carry information
  - **Address bus** → to determine where such information should be sent
  - **Control bus** → to indicate which operation should be performed
- More dedicated buses have been added to manage CPU-to-memory and I/O traffic
  - PCI, SATA, USB, etc.

# System Bus



# I/O Devices

# Components

- Each I/O device is made of **2 parts**:

# Components

- Each I/O device is made of **2 parts**:
  - the **physical device** itself

# Components

- Each I/O device is made of **2 parts**:
  - the **physical device** itself
  - the **device controller** (chip or set of chips controlling a family of physical devices)

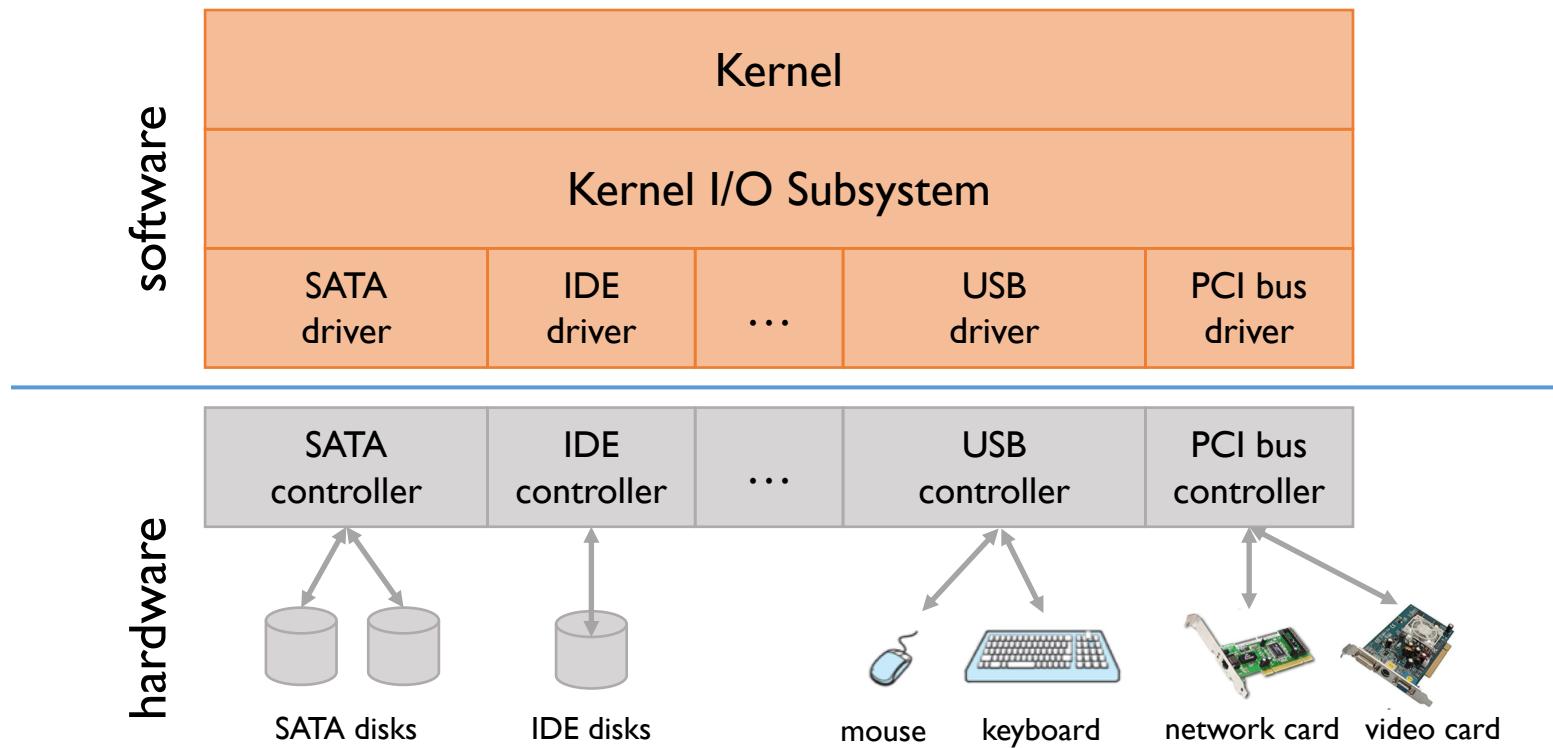
# Components

- Each I/O device is made of **2 parts**:
  - the **physical device** itself
  - the **device controller** (chip or set of chips controlling a family of physical devices)
- Can be categorized as:
  - storage, communications, user-interface, etc.

# Components

- Each I/O device is made of **2 parts**:
  - the **physical device** itself
  - the **device controller** (chip or set of chips controlling a family of physical devices)
- Can be categorized as:
  - storage, communications, user-interface, etc.
- OS talks to a device controller using a specific **device driver**

# Device Drivers: OS Software Abstraction



# Communicating with Device Controllers

- Every device controller has a number of dedicated registers to communicate with it:

# Communicating with Device Controllers

- Every device controller has a number of dedicated registers to communicate with it:
  - **Status registers** → provide status information to the CPU about the I/O device (e.g., idle, ready for input, busy, error, transaction complete)

# Communicating with Device Controllers

- Every device controller has a number of dedicated registers to communicate with it:
  - **Status registers** → provide status information to the CPU about the I/O device (e.g., idle, ready for input, busy, error, transaction complete)
  - **Configuration/Control registers** → used by the CPU to configure and control the device

# Communicating with Device Controllers

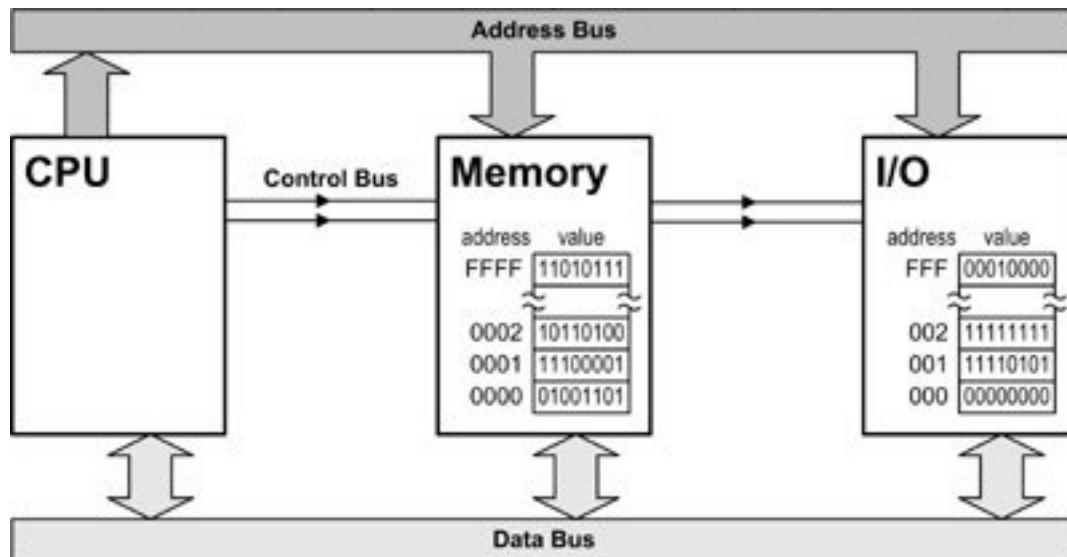
- Every device controller has a number of dedicated registers to communicate with it:
  - **Status registers** → provide status information to the CPU about the I/O device (e.g., idle, ready for input, busy, error, transaction complete)
  - **Configuration/Control registers** → used by the CPU to configure and control the device
  - **Data registers** → used to read data from or send data to the I/O device

# Communicating with Device Controllers

- Every device controller has a number of dedicated registers to communicate with it:
  - **Status registers** → provide status information to the CPU about the I/O device (e.g., idle, ready for input, busy, error, transaction complete)
  - **Configuration/Control registers** → used by the CPU to configure and control the device
  - **Data registers** → used to read data from or send data to the I/O device

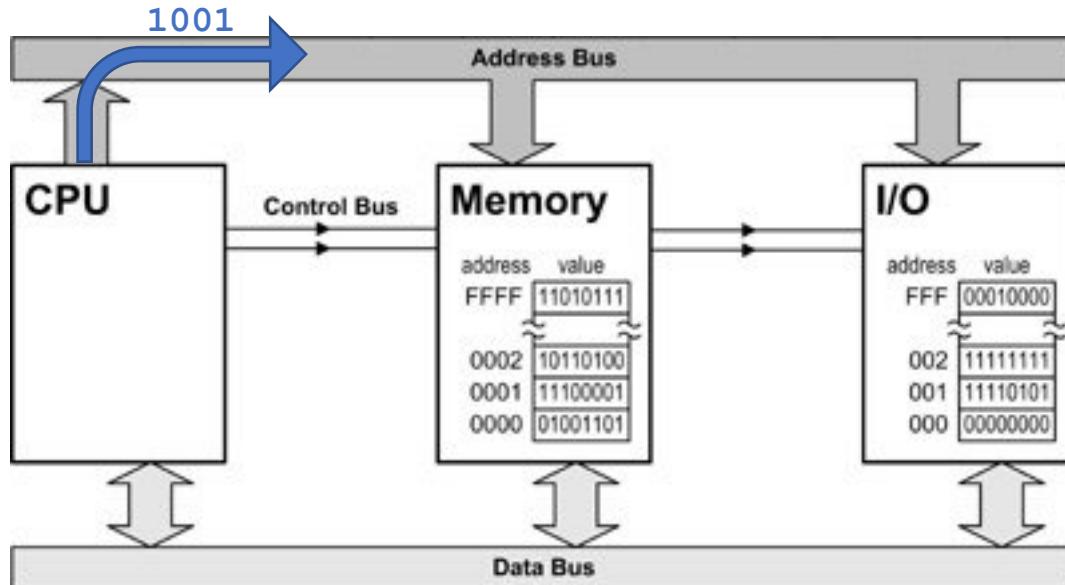
How does the CPU know how to address (registers of) I/O devices?

# Addressing Using the System Bus



How does CPU reference Memory addresses?

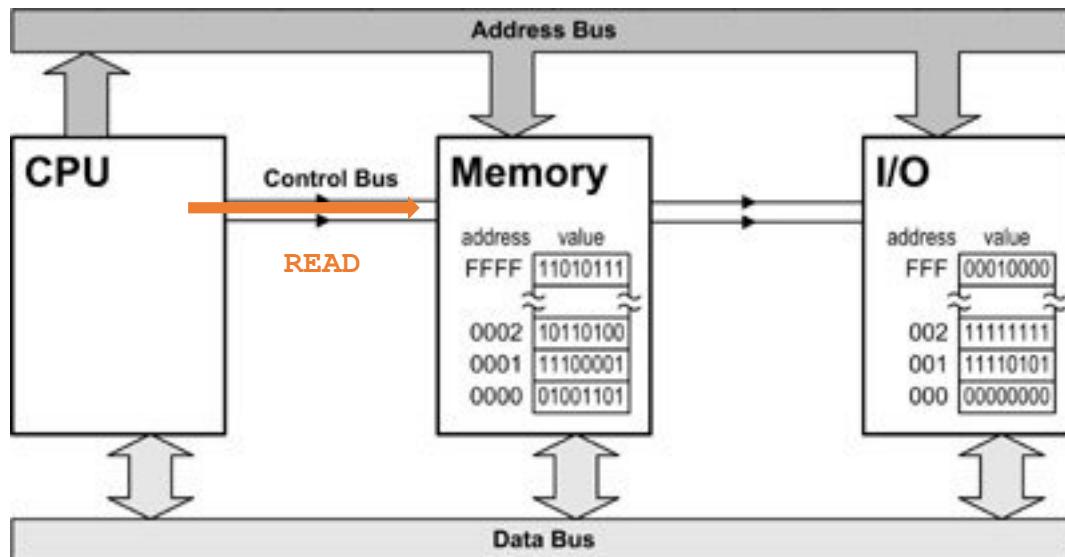
# Addressing Using the System Bus



How does CPU reference Memory addresses?

It puts the address of a byte of memory on the address bus

# Addressing Using the System Bus

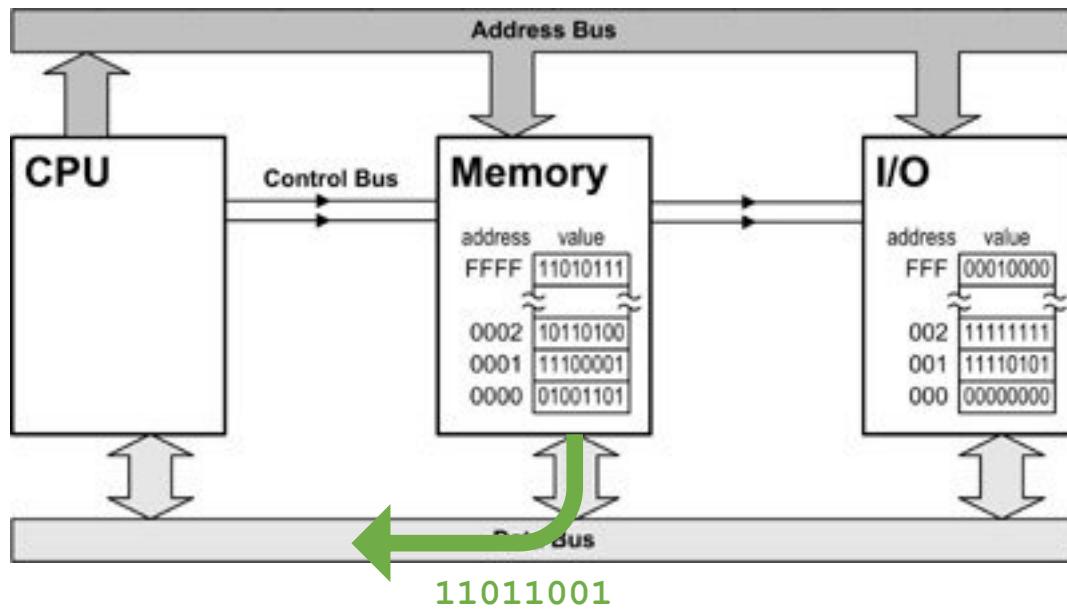


How does CPU reference Memory addresses?

It puts the address of a byte of memory on the address bus

It raises the **READ** signal on the control bus

# Addressing Using the System Bus



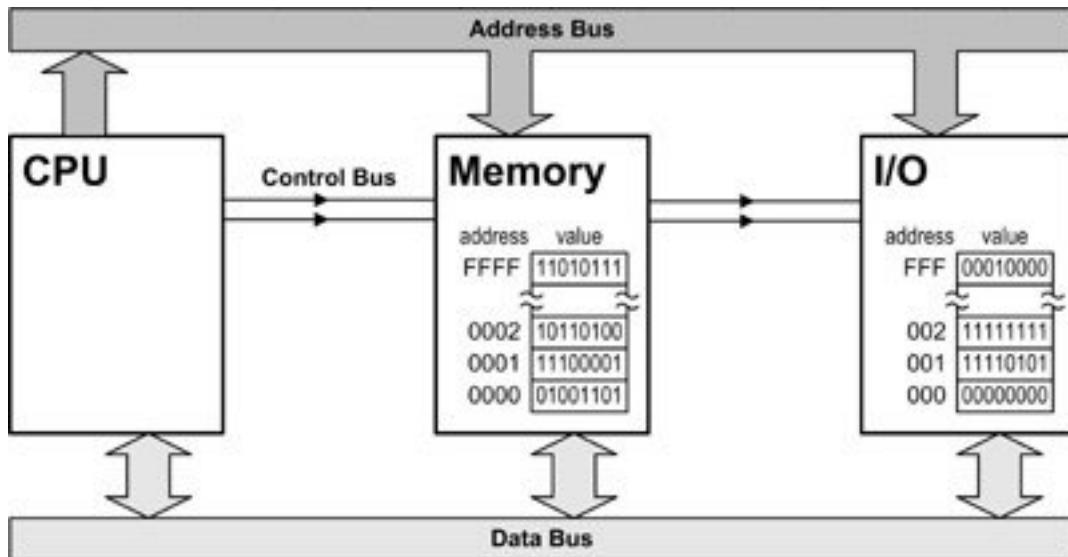
How does CPU reference Memory addresses?

It puts the address of a byte of memory on the address bus

It raises the **READ** signal on the control bus

Eventually, the RAM replies with the memory content on the data bus

# Addressing Using the System Bus



How does CPU reference Memory addresses?

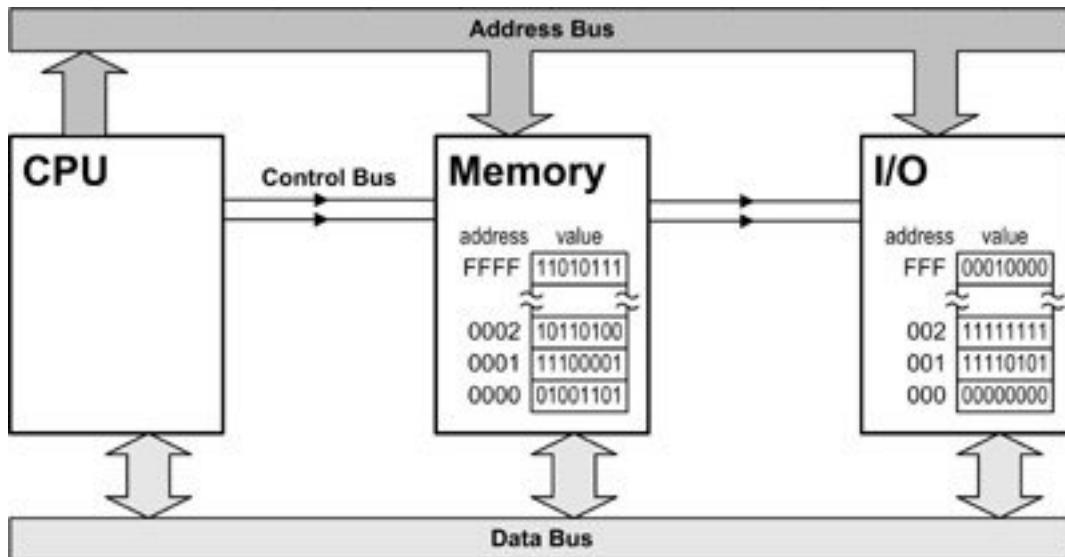
It puts the address of a byte of memory on the address bus

It raises the **READ** signal on the control bus

Eventually, the RAM replies with the memory content on the data bus

How about I/O devices? How to distinguish between Memory and I/O devices?

# Addressing Using the System Bus



How does CPU reference Memory addresses?

It puts the address of a byte of memory on the address bus

It raises the **READ** signal on the control bus

Eventually, the RAM replies with the memory content on the data bus

How about I/O devices? How to distinguish between Memory and I/O devices?

The control bus has a special line called "**M/#IO**" which asserts whether the CPU wants to talk to memory or an I/O device

# Port- vs. Memory-Mapped I/O

- CPU can talk to a device controller in **2 ways**:

# Port- vs. Memory-Mapped I/O

- CPU can talk to a device controller in **2 ways**:
  - **Port-mapped I/O** → referencing controller's registers using a separate I/O address space

# Port- vs. Memory-Mapped I/O

- CPU can talk to a device controller in **2 ways**:
  - **Port-mapped I/O** → referencing controller's registers using a separate I/O address space
  - **Memory-mapped I/O** → mapping controller's registers to the same address space used for main memory

# Port-Mapped I/O

- Each I/O device controller's register is mapped to a specific port (address)
- Requires special class of CPU instructions (e.g., IN/OUT)
  - The IN instruction reads from an I/O device, OUT writes
  - When you use the IN or OUT instructions, the M/#IO is not asserted, so memory does not respond and the I/O chip does

# Memory-Mapped I/O

- Memory-mapped I/O "wastes" some address space but doesn't need any special instruction
- To the CPU I/O device ports are just like normal memory addresses
- The CPU use MOV-like instructions to access I/O device registers
- In this way, the **M/#IO** is asserted indicating the address requested by the CPU refers to main memory

# Port- vs. Memory-Mapped I/O

```
MOV DX,1234h  
MOV AL,[DX]      ; reads memory address 1234h (memory address space)  
IN AL,DX        ; reads I/O port 1234h (I/O address space)
```

Both put the value **1234h** on the CPU address bus, and both assert a **READ** operation on control bus

# Port- vs. Memory-Mapped I/O

```
MOV DX,1234h  
MOV AL,[DX]      ; reads memory address 1234h (memory address space)  
IN AL,DX        ; reads I/O port 1234h (I/O address space)
```

The first one will assert **M/#IO** to indicate that the address belongs to memory address space

# Port- vs. Memory-Mapped I/O

```
MOV DX,1234h  
MOV AL,[DX]      ; reads memory address 1234h (memory address space)  
IN AL,DX        ; reads I/O port 1234h (I/O address space)
```

The second one will not assert M/#IO to indicate that the address belongs to I/O address space

# Performing I/O Tasks

- **Polling**
  - CPU periodically checks for the I/O task status

# Performing I/O Tasks

- **Polling**
  - CPU periodically checks for the I/O task status
- **Interrupt-driven**
  - CPU receives an interrupt from the controller (device or DMA) once the I/O task is done (either successfully or abnormally)

# Performing I/O Tasks

- **Polling**
  - CPU periodically checks for the I/O task status
- **Interrupt-driven**
  - CPU receives an interrupt from the controller (device or DMA) once the I/O task is done (either successfully or abnormally)
- **Programmed I/O**
  - CPU does the actual work of moving data

# Performing I/O Tasks

- **Polling**
  - CPU periodically checks for the I/O task status
- **Interrupt-driven**
  - CPU receives an interrupt from the controller (device or DMA) once the I/O task is done (either successfully or abnormally)
- **Programmed I/O**
  - CPU does the actual work of moving data
- **Direct Memory Access (DMA)**
  - CPU delegates off the work to a dedicated DMA controller

# Performing I/O Tasks

- **Polling**
  - CPU periodically checks for the I/O task status
- **Interrupt-driven**
  - CPU receives an interrupt from the controller (device or DMA) once the I/O task is done (either successfully or abnormally)
- **Programmed I/O**
  - CPU does the actual work of moving data
- **Direct Memory Access (DMA)**
  - CPU delegates off the work to a dedicated DMA controller

**HOW?**

# Performing I/O Tasks

- **Polling**
  - CPU periodically checks for the I/O task status
- **Interrupt-driven**
  - CPU receives an interrupt from the controller (device or DMA) once the I/O task is done (either successfully or abnormally)

- **Programmed I/O**

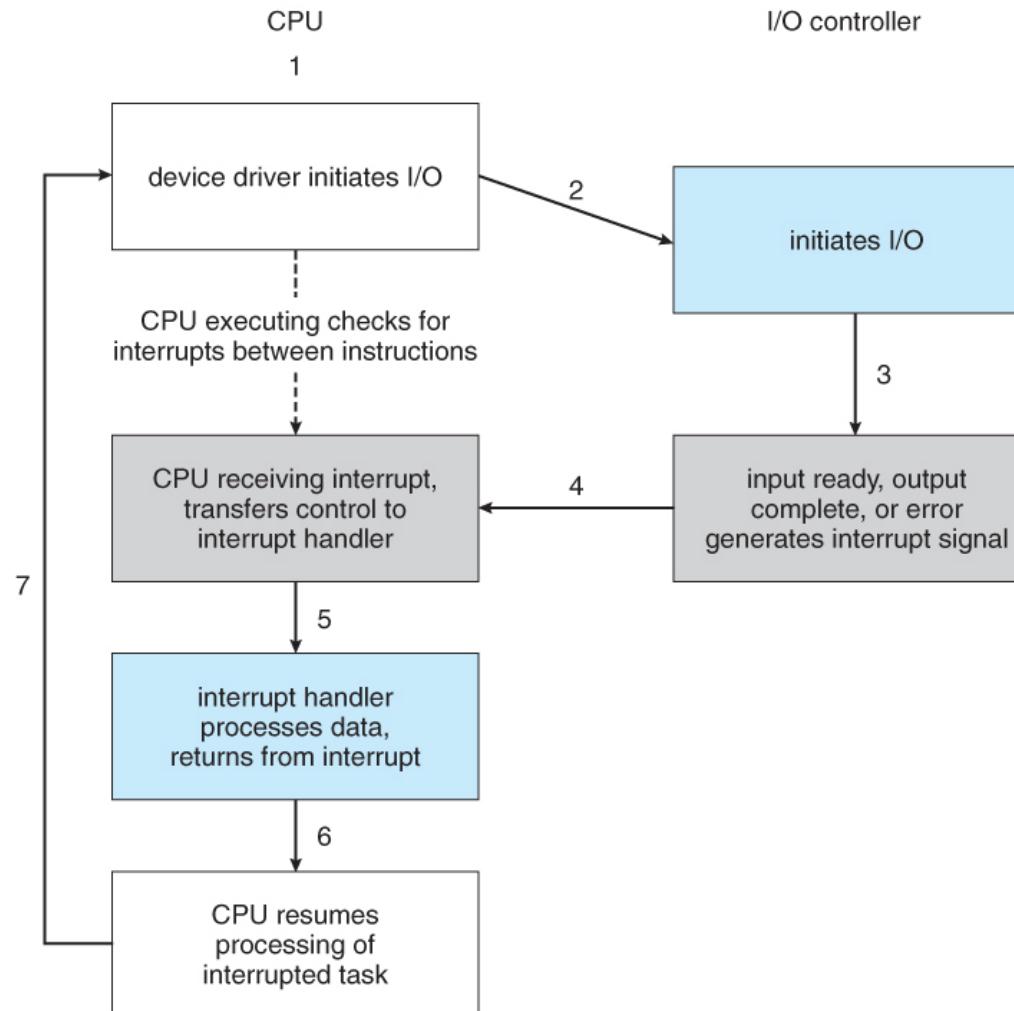
- CPU does the actual work of moving data

**WHO?**

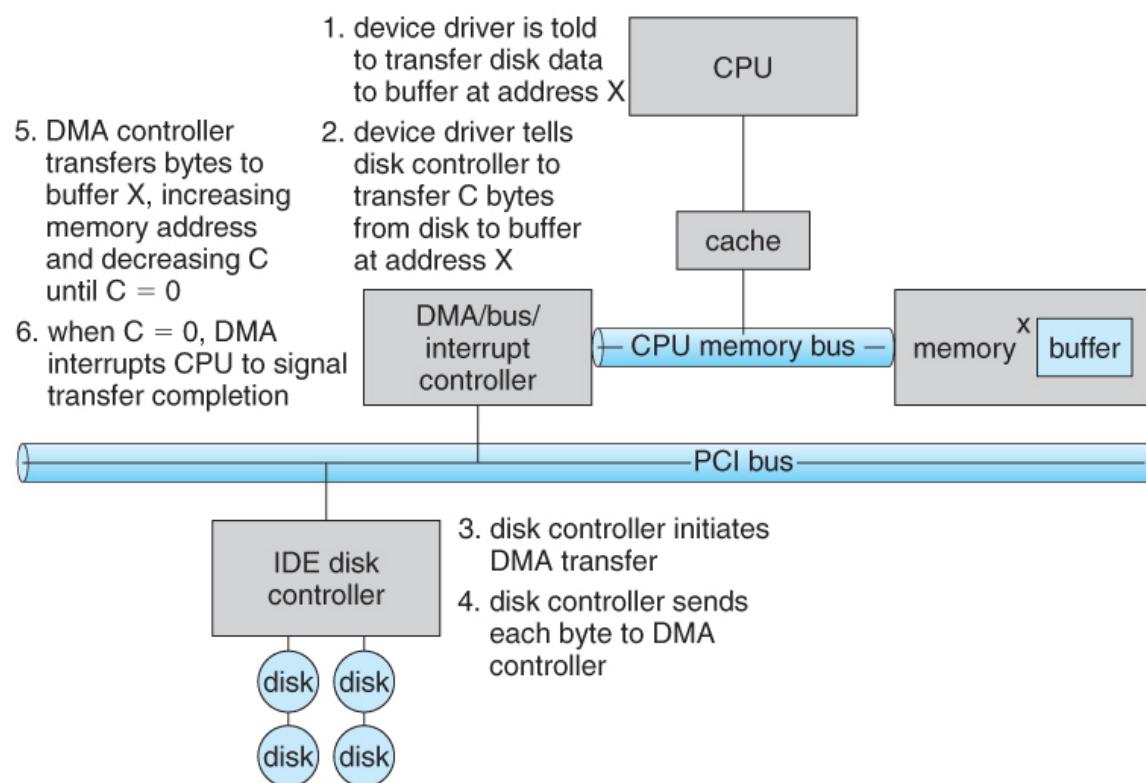
- **Direct Memory Access (DMA)**

- CPU delegates off the work to a dedicated DMA controller

# Interrupt-driven I/O



# Direct Memory Access (DMA)



Overcome the limitation of Programmed I/O

Maybe wasteful to tie up the CPU transferring data in and out of registers **one byte at a time**

Useful for devices that transfer large quantities of data (such as disk controllers)

Typically, used in combination with interrupt-driven I/O

# Architectural Features Enabling OS Services

OS Service	HW Support
Protection and Security	Kernel/user mode, protected instructions, base/limit registers
System calls	Trap instructions and interrupt vectors
Exception handling	Trap instructions and interrupt vectors
I/O operations	Trap instructions, interrupt vectors, and memory mapping
Scheduling	Timer
Synchronization	Atomic instructions
Virtual memory	Translation Look-aside Buffer (TLB)

# Protection and Security

# Privileged Instructions

- Some CPU instructions are more sensitive than others
  - `MOV %eax, %ebx` → move the content of the register ebx **into** eax
  - `MOV %eax, [%ebx]` → move the content of memory indexed by register ebx **to** eax
  - `HLT` → halt the system
  - `INT X` → generate interrupt X

# Privileged Instructions

- Some CPU instructions are more sensitive than others
  - `MOV %eax, %ebx` → move the content of the register ebx **into** eax
  - `MOV %eax, [%ebx]` → move the content of memory indexed by register ebx **to** eax
  - `HLT` → halt the system
  - `INT X` → generate interrupt X

# Privileged Instructions

- Some CPU instructions are more sensitive than others
  - `MOV %eax, %ebx` → move the content of the register ebx **into** eax
  - `MOV %eax, [%ebx]` → move the content of memory indexed by register ebx **to** eax
  - `HLT` → halt the system
  - `INT X` → generate interrupt X

**Idea:** sensitive (privileged) instructions can be executed only by the OS

# Kernel vs. User Mode

- 2 different "states" while executing CPU instructions

# Kernel vs. User Mode

- 2 different "states" while executing CPU instructions
- **Kernel mode** is unrestricted:
  - The OS can perform any instruction (including privileged ones)

# Kernel vs. User Mode

- 2 different "states" while executing CPU instructions
- **Kernel mode** is unrestricted:
  - The OS can perform any instruction (including privileged ones)
- **User mode** is restricted so that the user is not able to:
  - Address I/O (directly)
  - Manipulate the content of main memory
  - Halt the machine
  - Switch to kernel mode
  - etc.

# Kernel vs. User Mode

- 2 different "states" while executing CPU instructions
- **Kernel mode** is unrestricted:
  - The OS can perform any instruction (including privileged ones)
- **User mode** is restricted so that the user is not able to:
  - Address I/O (directly)
  - Manipulate the content of main memory
  - Halt the machine
  - Switch to kernel mode
  - etc.

Implemented in HW!

A status bit stored in a protected CPU register  
(0=kernel, 1=user)

# Beyond Kernel vs. User Mode

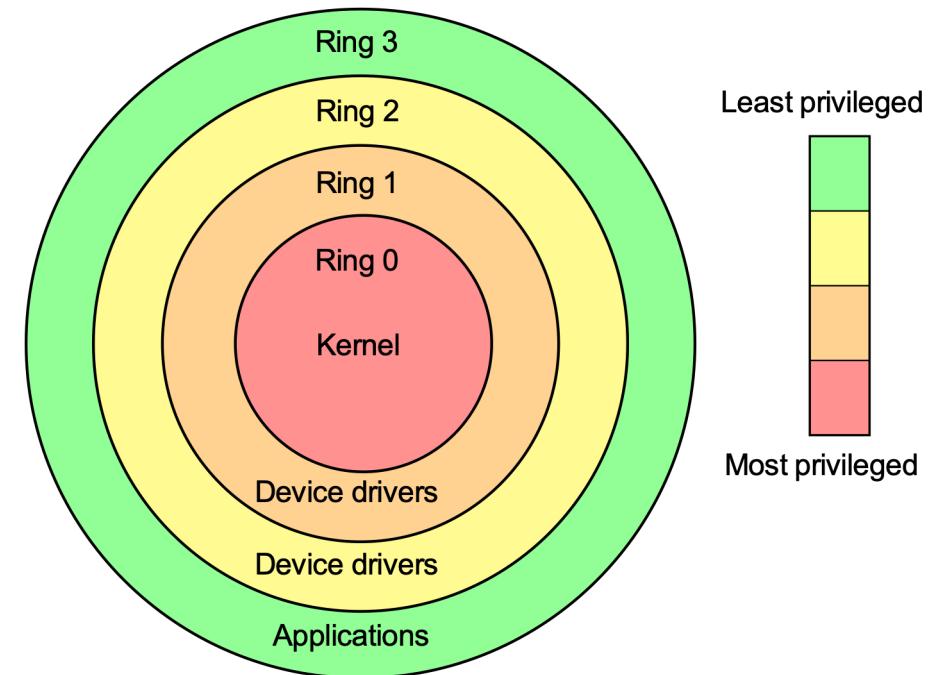
- The underlying HW must support at least kernel and user mode

# Beyond Kernel vs. User Mode

- The underlying HW must support at least kernel and user mode
- More fine-grained solutions are also possible

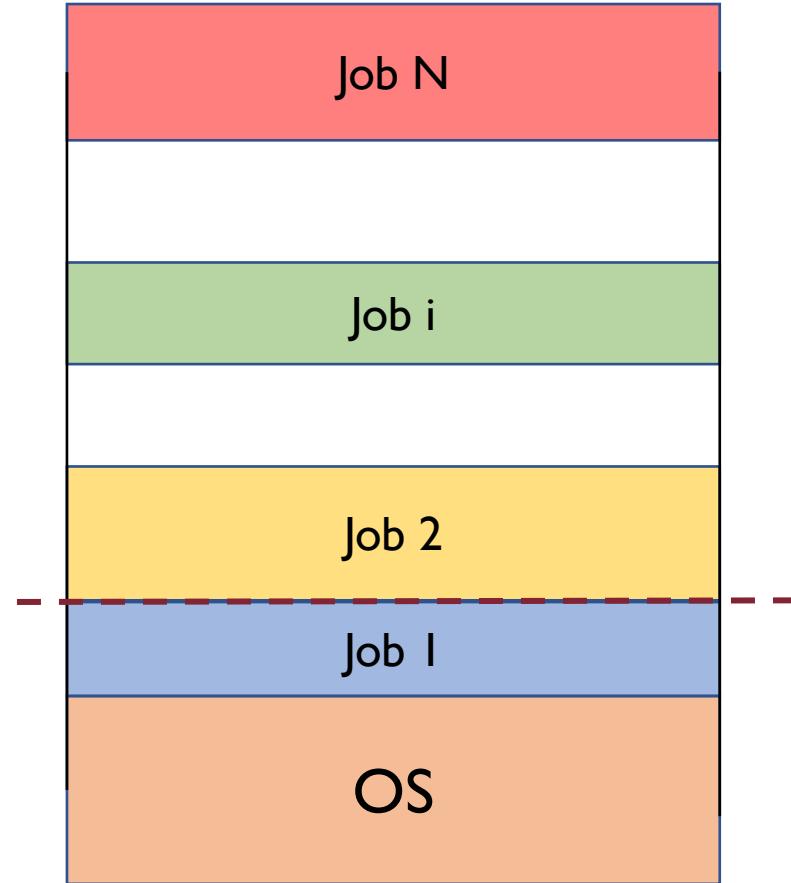
# Beyond Kernel vs. User Mode

- The underlying HW must support at least kernel and user mode
- More fine-grained solutions are also possible
- **Protection Rings**
  - 4 different privilege levels  $\{0, \dots, 3\}$
  - Still implementable in HW (2 bits)



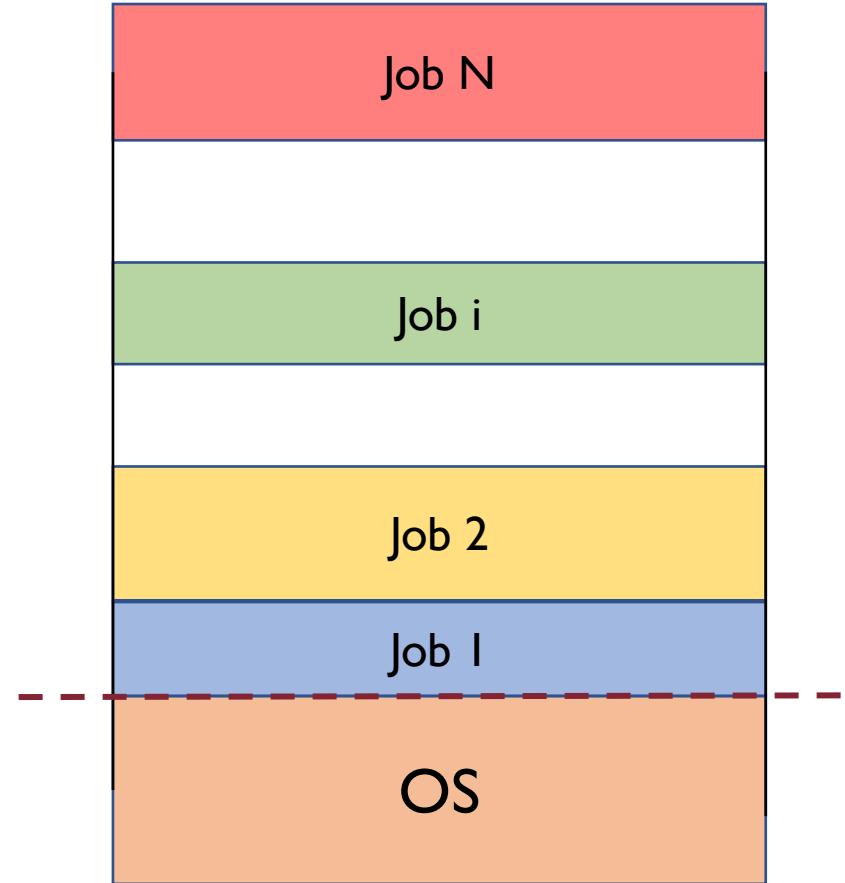
# Memory Protection

- Architecture must provide support for the OS to:
  - Protect user programs from each other



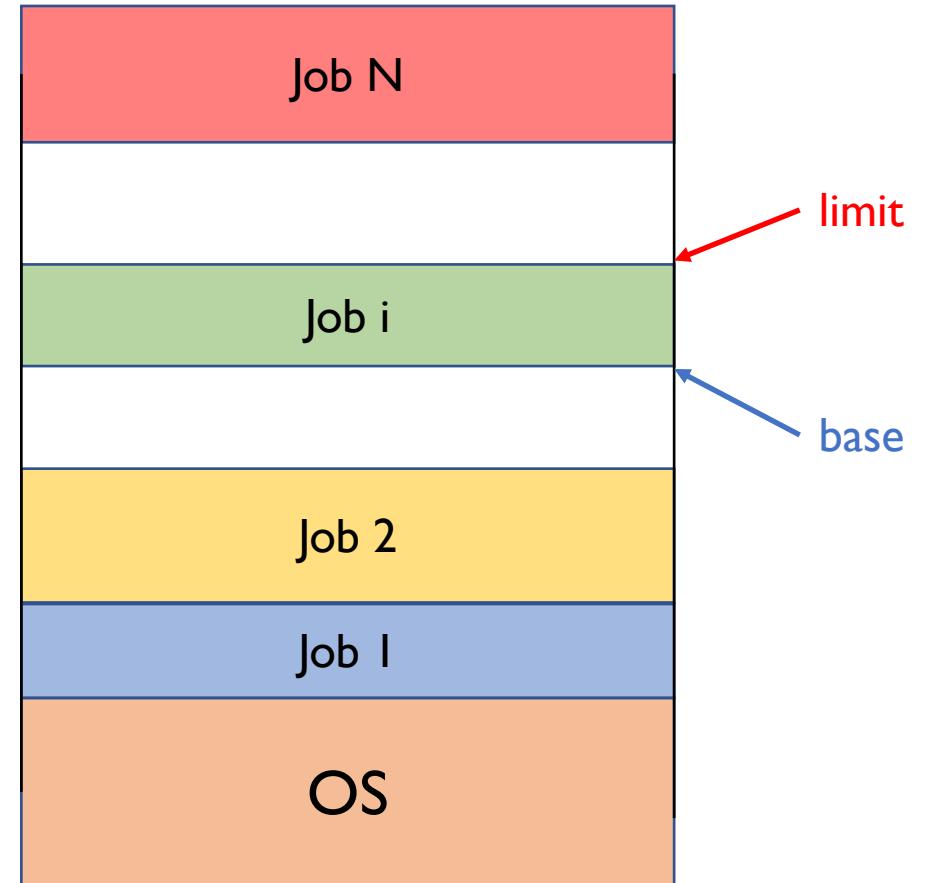
# Memory Protection

- Architecture must provide support for the OS to:
  - Protect the OS from user programs



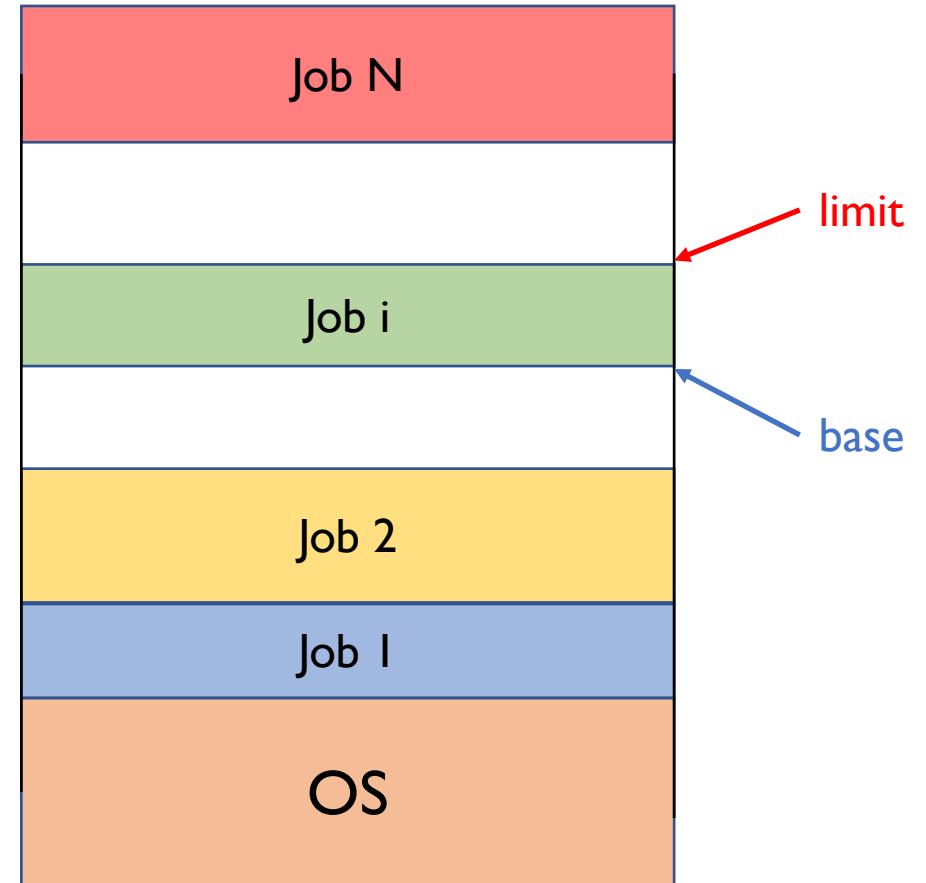
# Memory Protection

- The simplest technique is to have **2 dedicated registers**
  - **base** → contains the starting valid memory address
  - **limit** → contains the last valid memory address



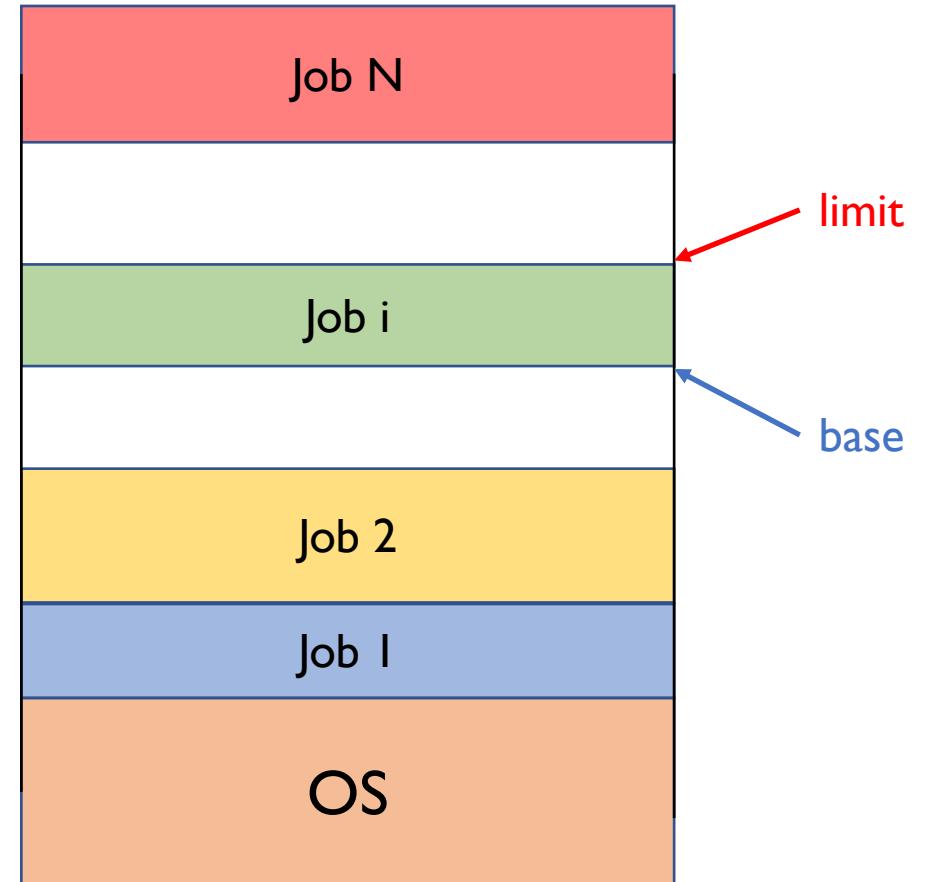
# Memory Protection

- The simplest technique is to have **2 dedicated registers**
  - **base** → contains the starting valid memory address
  - **limit** → contains the last valid memory address
- The OS loads the **base** and **limit** registers upon program startup



# Memory Protection

- The simplest technique is to have **2 dedicated registers**
  - **base** → contains the starting valid memory address
  - **limit** → contains the last valid memory address
- The OS loads the **base** and **limit** registers upon program startup
- The CPU checks each memory address referenced by user program falls between **base** and **limit** values



# Program vs. Process

- A **program** is an executable file which resides on the persistent memory (e.g., disk),
  - contains only the set of instructions needed to accomplish a specific job
  - e.g., the `ls` program is an executable file stored at `/bin/ls` on the disk of a UNIX-like OS

# Program vs. Process

- A **program** is an executable file which resides on the persistent memory (e.g., disk),
  - contains only the set of instructions needed to accomplish a specific job
  - e.g., the `ls` program is an executable file stored at `/bin/ls` on the disk of a UNIX-like OS
- A **process** is a particular instance of a program when loaded to main memory
  - e.g., multiple instances of the `ls` program above, thus multiple processes for the same program

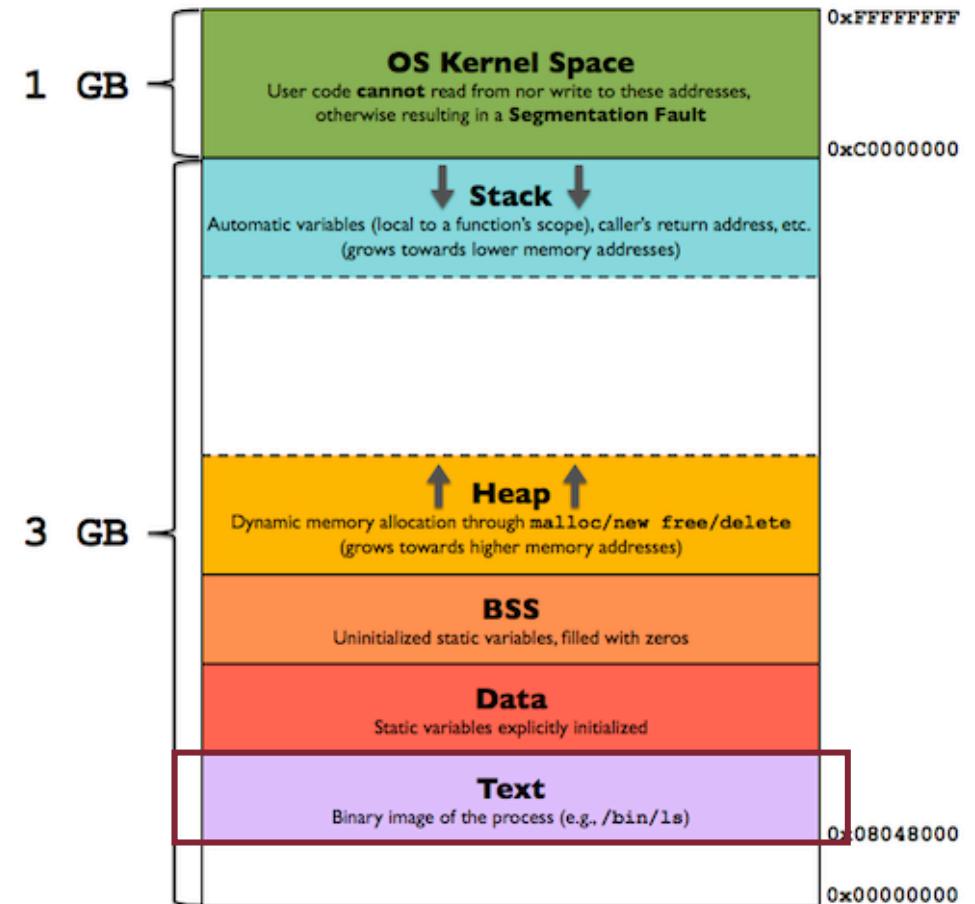
# Program vs. Process

- A **program** is an executable file which resides on the persistent memory (e.g., disk),
  - contains only the set of instructions needed to accomplish a specific job
  - e.g., the `ls` program is an executable file stored at `/bin/ls` on the disk of a UNIX-like OS
- A **process** is a particular instance of a program when loaded to main memory
  - e.g., multiple instances of the `ls` program above, thus multiple processes for the same program

program → "static/pассивный" vs. process → "dynamic/активный"

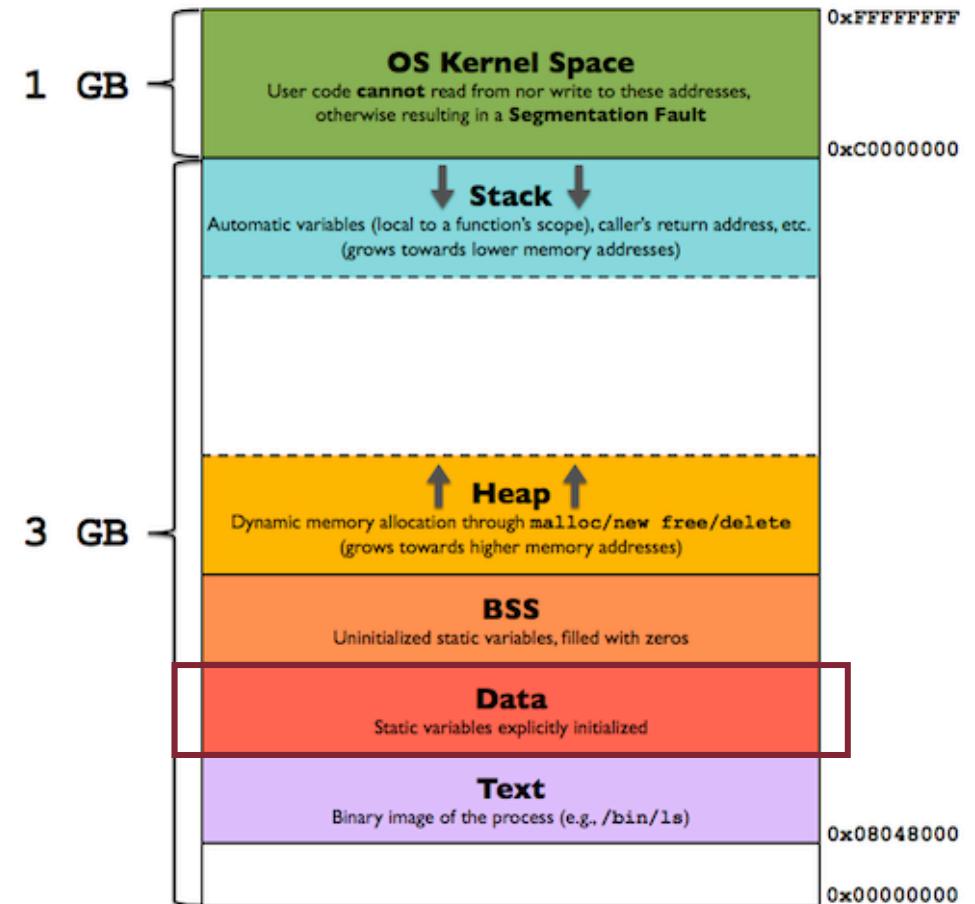
# Process: In-Memory Layout

- **Text** → contains executable instructions



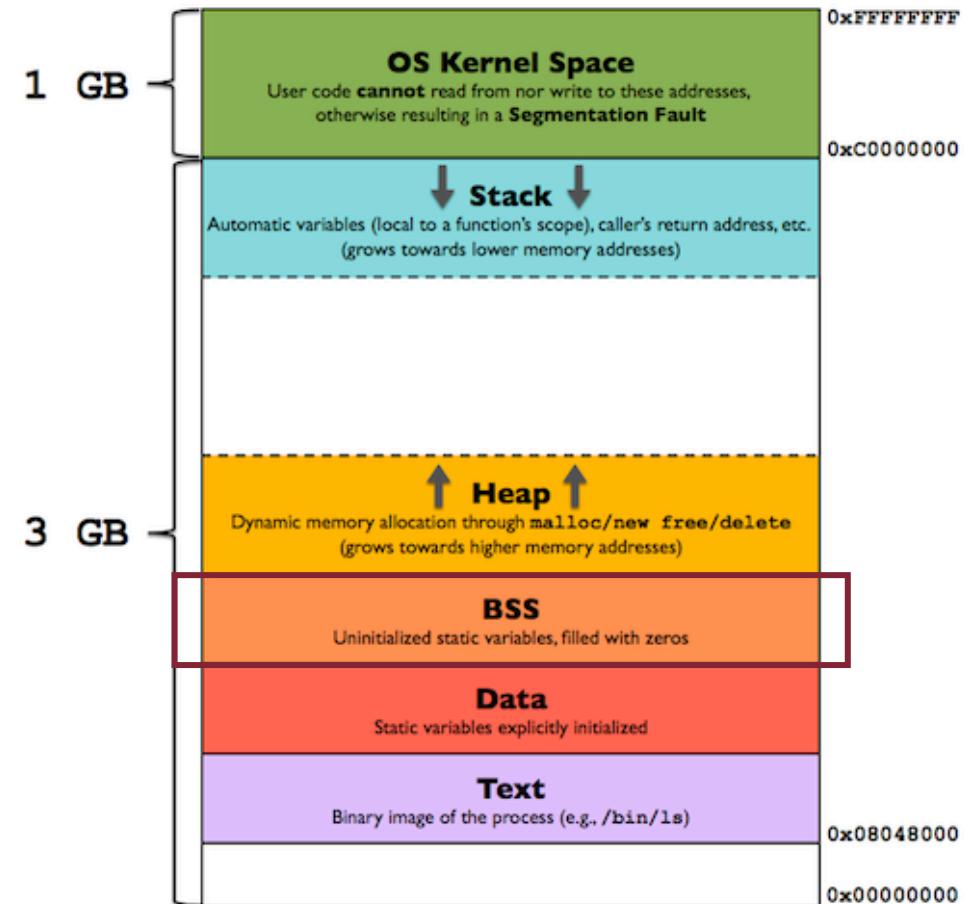
# Process: In-Memory Layout

- **Text** → contains executable instructions
- **Data** → global and static variable (initialized)



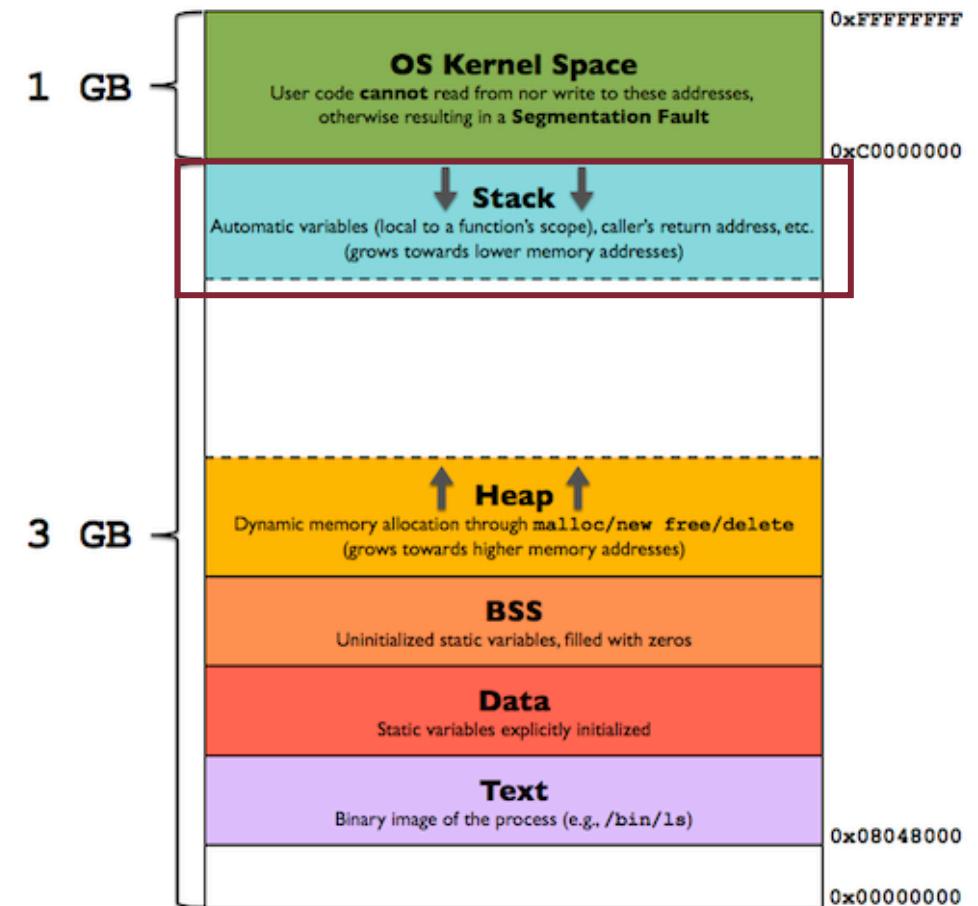
# Process: In-Memory Layout

- **Text** → contains executable instructions
- **Data** → global and static variable (initialized)
- **BSS (Block Started by Symbol)** → global and static variable (uninitialized or initialized to 0)



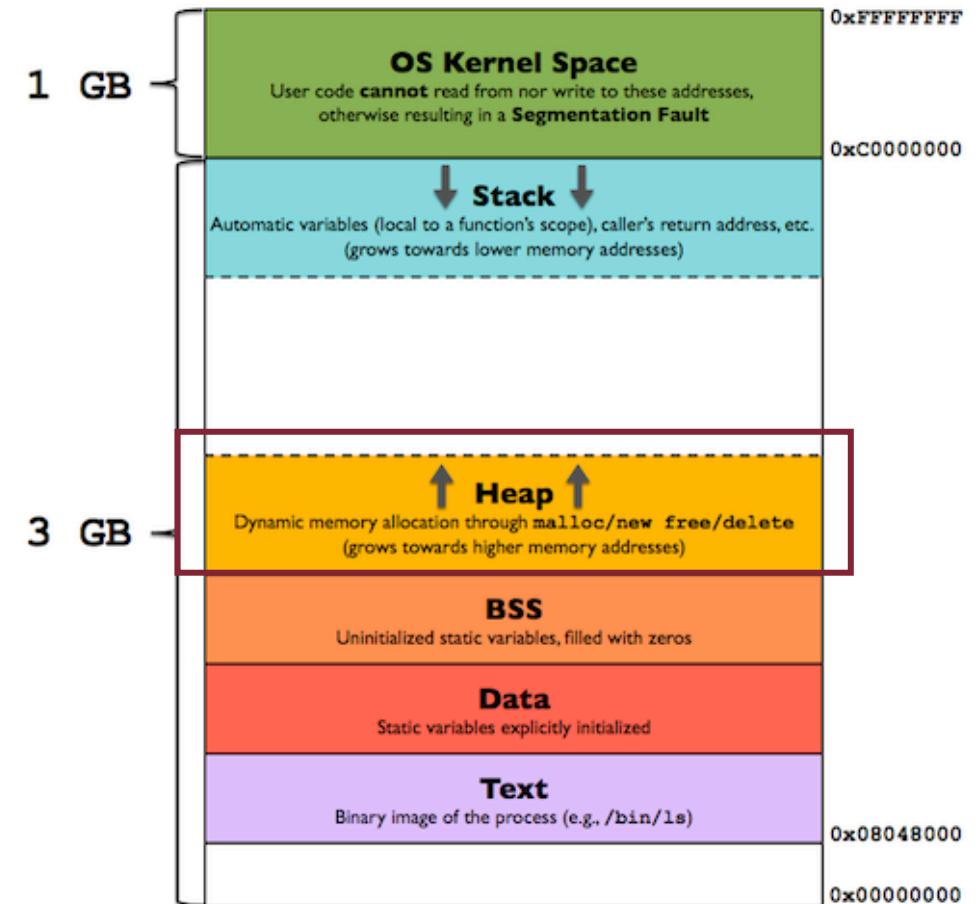
# Process: In-Memory Layout

- **Text** → contains executable instructions
- **Data** → global and static variable (initialized)
- **BSS (Block Started by Symbol)** → global and static variable (uninitialized or initialized to 0)
- **Stack** → LIFO structure used to store all the data needed by a function call (**stack frame**)



# Process: In-Memory Layout

- **Text** → contains executable instructions
- **Data** → global and static variable (initialized)
- **BSS (Block Started by Symbol)** → global and static variable (uninitialized or initialized to 0)
- **Stack** → LIFO structure used to store all the data needed by a function call (**stack frame**)
- **Heap** → used for dynamic allocation



# Function Call: Stack Frame

- Each function uses a portion of the stack, and we call it a **stack frame**
- The stack frame for each function is divided into **3 parts**:
  - function parameters
  - back-pointer to the previous stack frame
  - local variables

# Stack Frame: Function Parameters

- This part of a function's stack frame is set up by the **caller**
- **2 operations** are defined on a stack:
  - **push** → used to place items onto the stack
  - **pop** → used to remove items from the stack
- Different languages may push the parameters on in different orders

# Stack Frame: Function Parameters

```
foo (a, b, c);
```

# Stack Frame: Function Parameters



# Stack Frame: Function Parameters



- Each item is pushed onto the stack, the stack grows down
- The stack-pointer register (**esp**) is decremented by 4 bytes (in 32-bit mode), and the item is copied to the memory location pointed to by the stack-pointer register
- The **call** instruction will implicitly push the return address on the stack

[esp + 0]	- return address
[esp + 4]	- parameter 'a'
[esp + 8]	- parameter 'b'
[esp + 12]	- parameter 'c'

# Stack Frame: Base Pointer

```
[esp + 0] - return address  
[esp + 4] - parameter 'a'  
[esp + 8] - parameter 'b'  
[esp + 12] - parameter 'c'
```

This way of accessing parameters from the callee might get clumsy, especially when several local variables need to be stored

# Stack Frame: Base Pointer

```
[esp + 0] - return address  
[esp + 4] - parameter 'a'  
[esp + 8] - parameter 'b'  
[esp + 12] - parameter 'c'
```

This way of accessing parameters from the callee might get clumsy, especially when several local variables need to be stored

Save the stack frame base pointer (**ebp**)!

```
push ebp      ; save previous stackbase-pointer register  
mov ebp, esp ; ebp = esp
```

```
[ebp + 16] - parameter 'c'  
[ebp + 12] - parameter 'b'  
[ebp + 8] - parameter 'a'  
[ebp + 4] - return address  
[ebp + 0] - saved stackbase-pointer register
```

# Stack Frame: Putting All Together

```
[ebp + 16] - parameter 'c'  
[ebp + 12] - parameter 'b'  
[ebp + 8] - parameter 'a'  
[ebp + 4] - return address  
[ebp + 0] - saved stackbase-pointer register
```

Parameters (of the caller) are accessed through the stack frame base pointer (**ebp**)

# Stack Frame: Putting All Together

```
[ebp + 16] - parameter 'c'  
[ebp + 12] - parameter 'b'  
[ebp + 8] - parameter 'a'  
[ebp + 4] - return address  
[ebp + 0] - saved stackbase-pointer register
```

Parameters (of the caller) are accessed through the stack frame base pointer (**ebp**)

```
[esp + (# - 4)] - top of local variables section  
[esp + 0] - bottom of local variables section
```

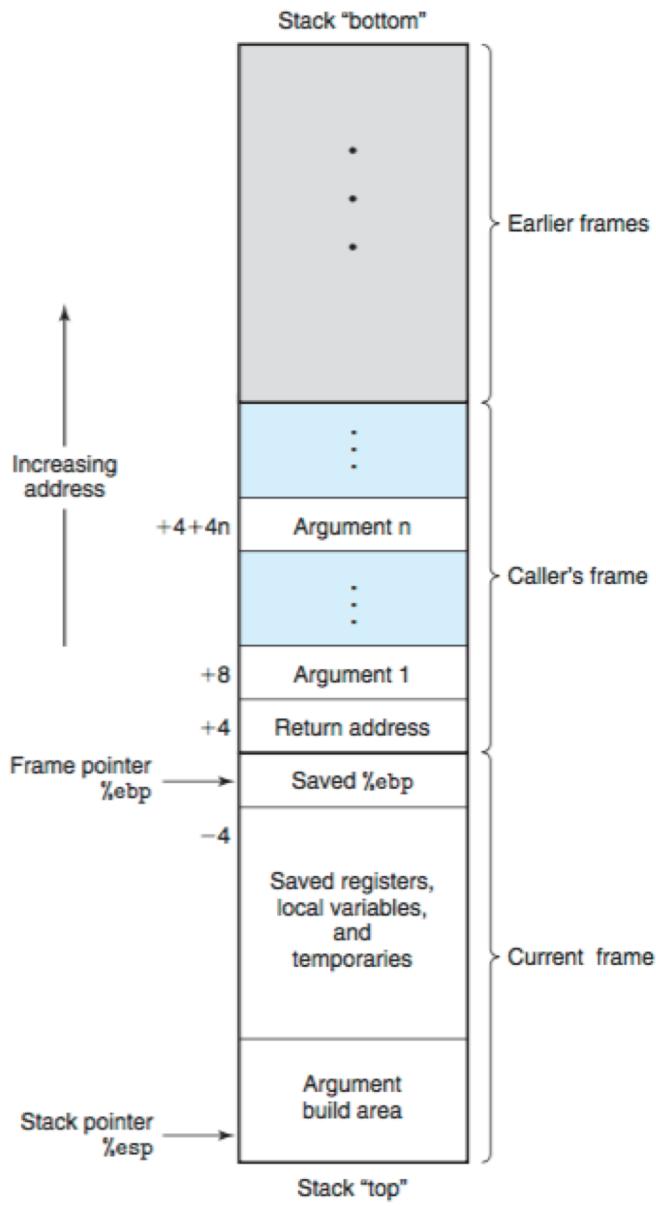
Local variables (of the callee) are accessed through the stack frame top pointer (**esp**)

# Stack Frame: Cleanup

```
mov esp, ebp ; undo the carving of space for the local variables  
pop ebp      ; restore the previous stackbase-pointer register
```

The old stack frame base pointer is restored

# Stack: Outline



# System Calls, Exceptions, I/O Interrupts

# Preserving Usability: System Calls

- Privileged instructions cannot be executed in user mode **directly**

# Preserving Usability: System Calls

- Privileged instructions cannot be executed in user mode **directly**
- But programs running in user mode can ask the OS to perform some restricted operations on their behalf (in kernel mode)

# Preserving Usability: System Calls

- Privileged instructions cannot be executed in user mode **directly**
- But programs running in user mode can ask the OS to perform some restricted operations on their behalf (in kernel mode)
- For example, a user program may require to
  - Write data to a file stored on disk
  - Send data over the network interface
  - etc.

# Preserving Usability: System Calls

- Privileged instructions cannot be executed in user mode **directly**
- But programs running in user mode can ask the OS to perform some restricted operations on their behalf (in kernel mode)
- For example, a user program may require to
  - Write data to a file stored on disk
  - Send data over the network interface
  - etc.

Crossing protection boundaries using **system calls**

# Exceptions and Interrupts

- **Exceptions**
  - software-generated
  - e.g., program error like division by 0

# Exceptions and Interrupts

- **Exceptions**

- software-generated
  - e.g., program error like division by 0

- **Interrupts**

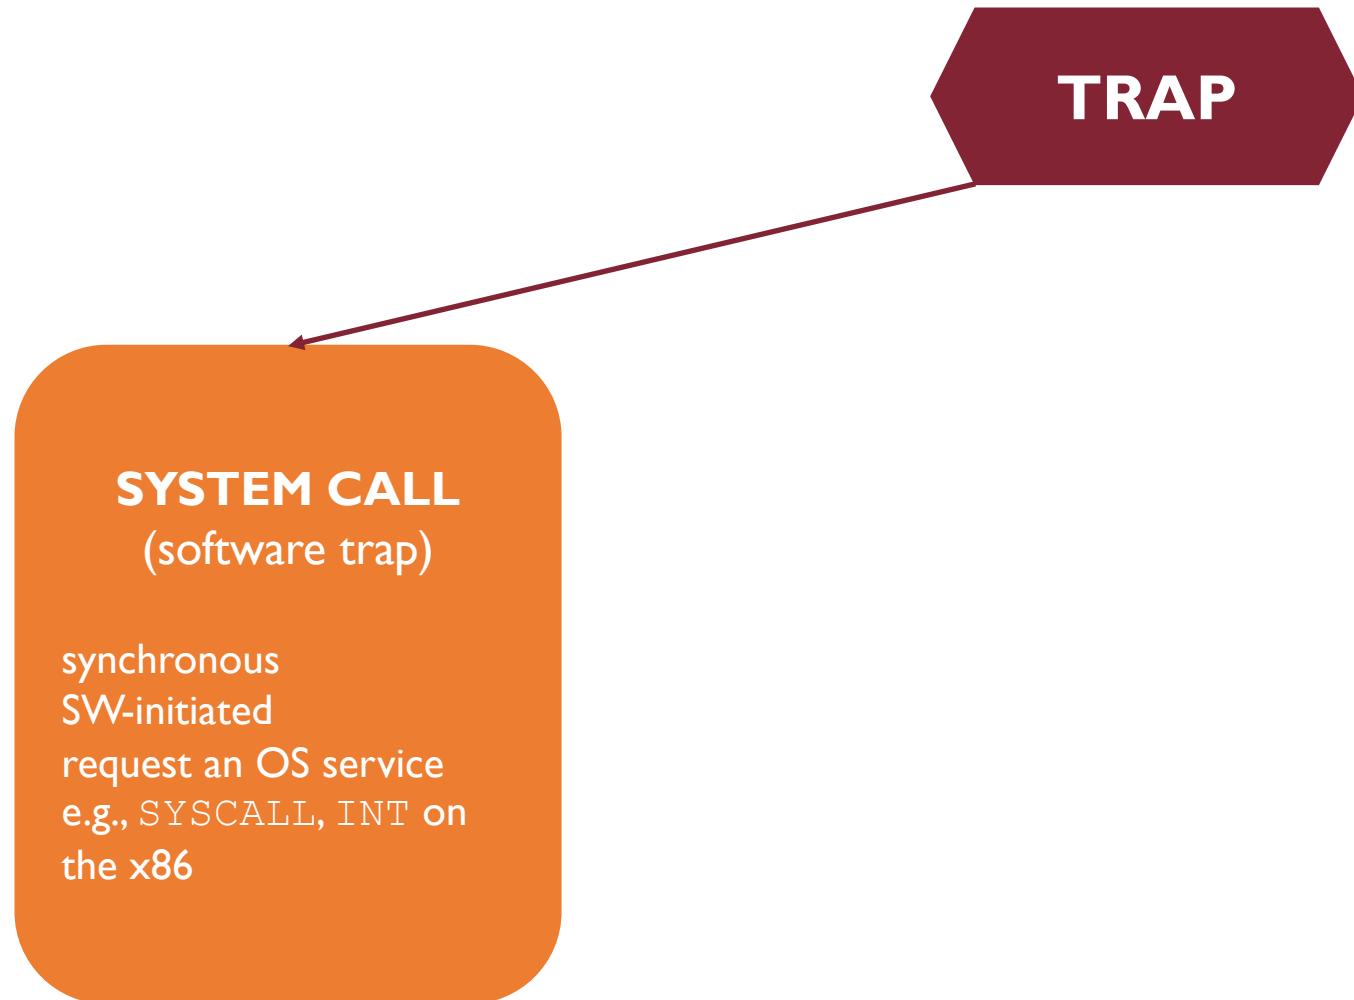
- hardware-generated (by external devices)
  - e.g., I/O completion or timer interrupt on a multi-tasking system

# A Quick Note on Terminology

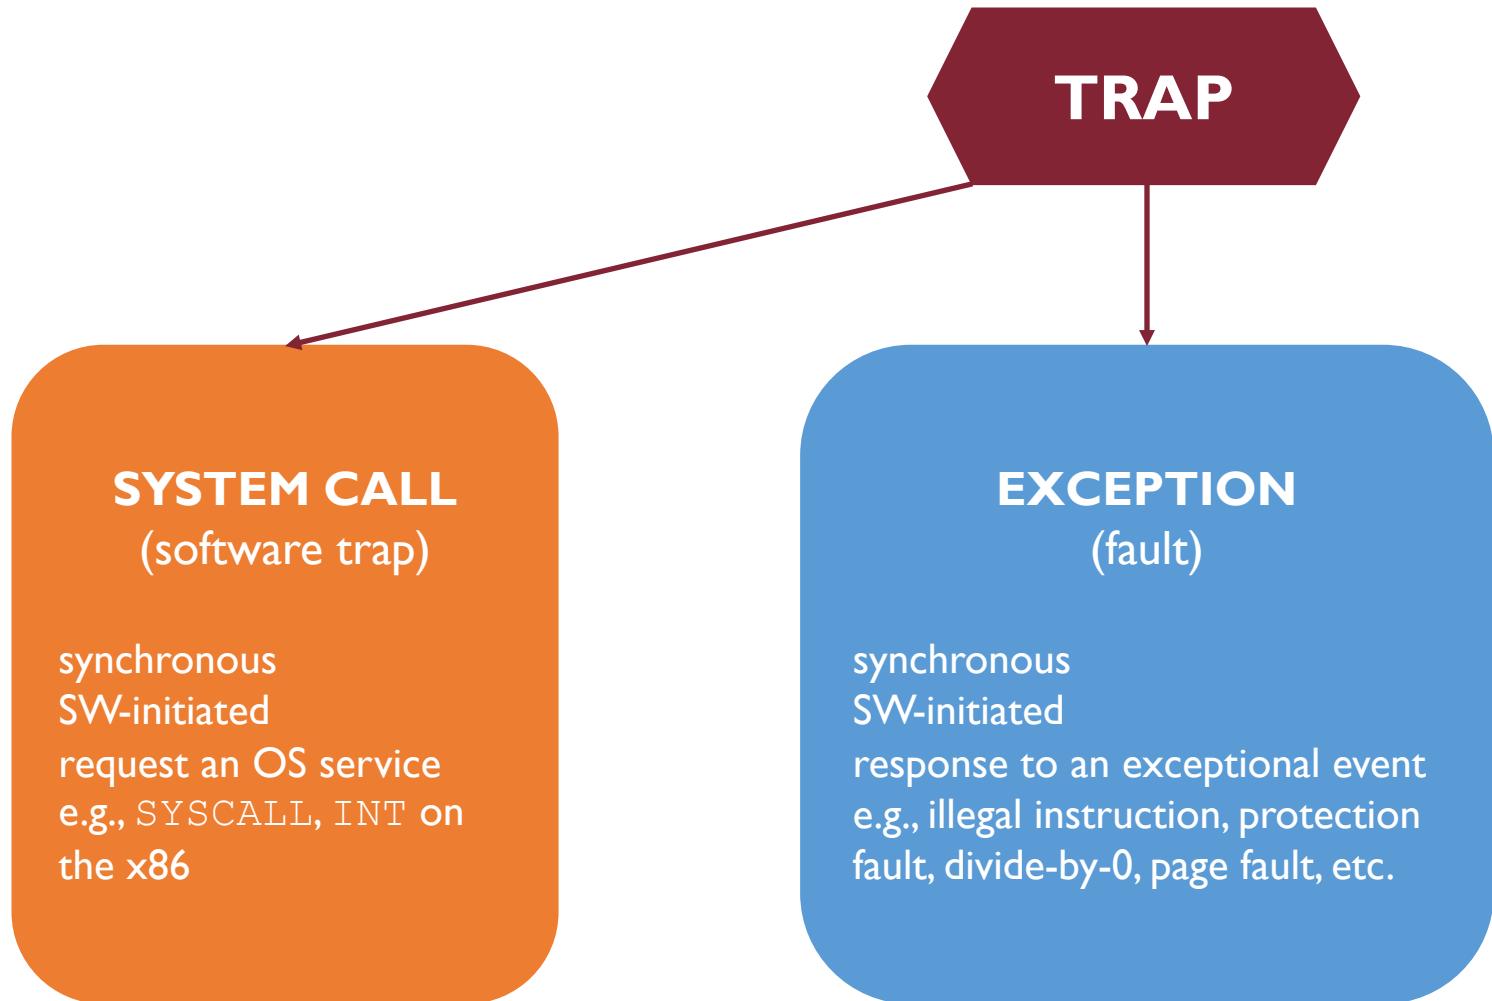
**TRAP**

We will refer to **trap** as any event that causes switch to OS kernel mode

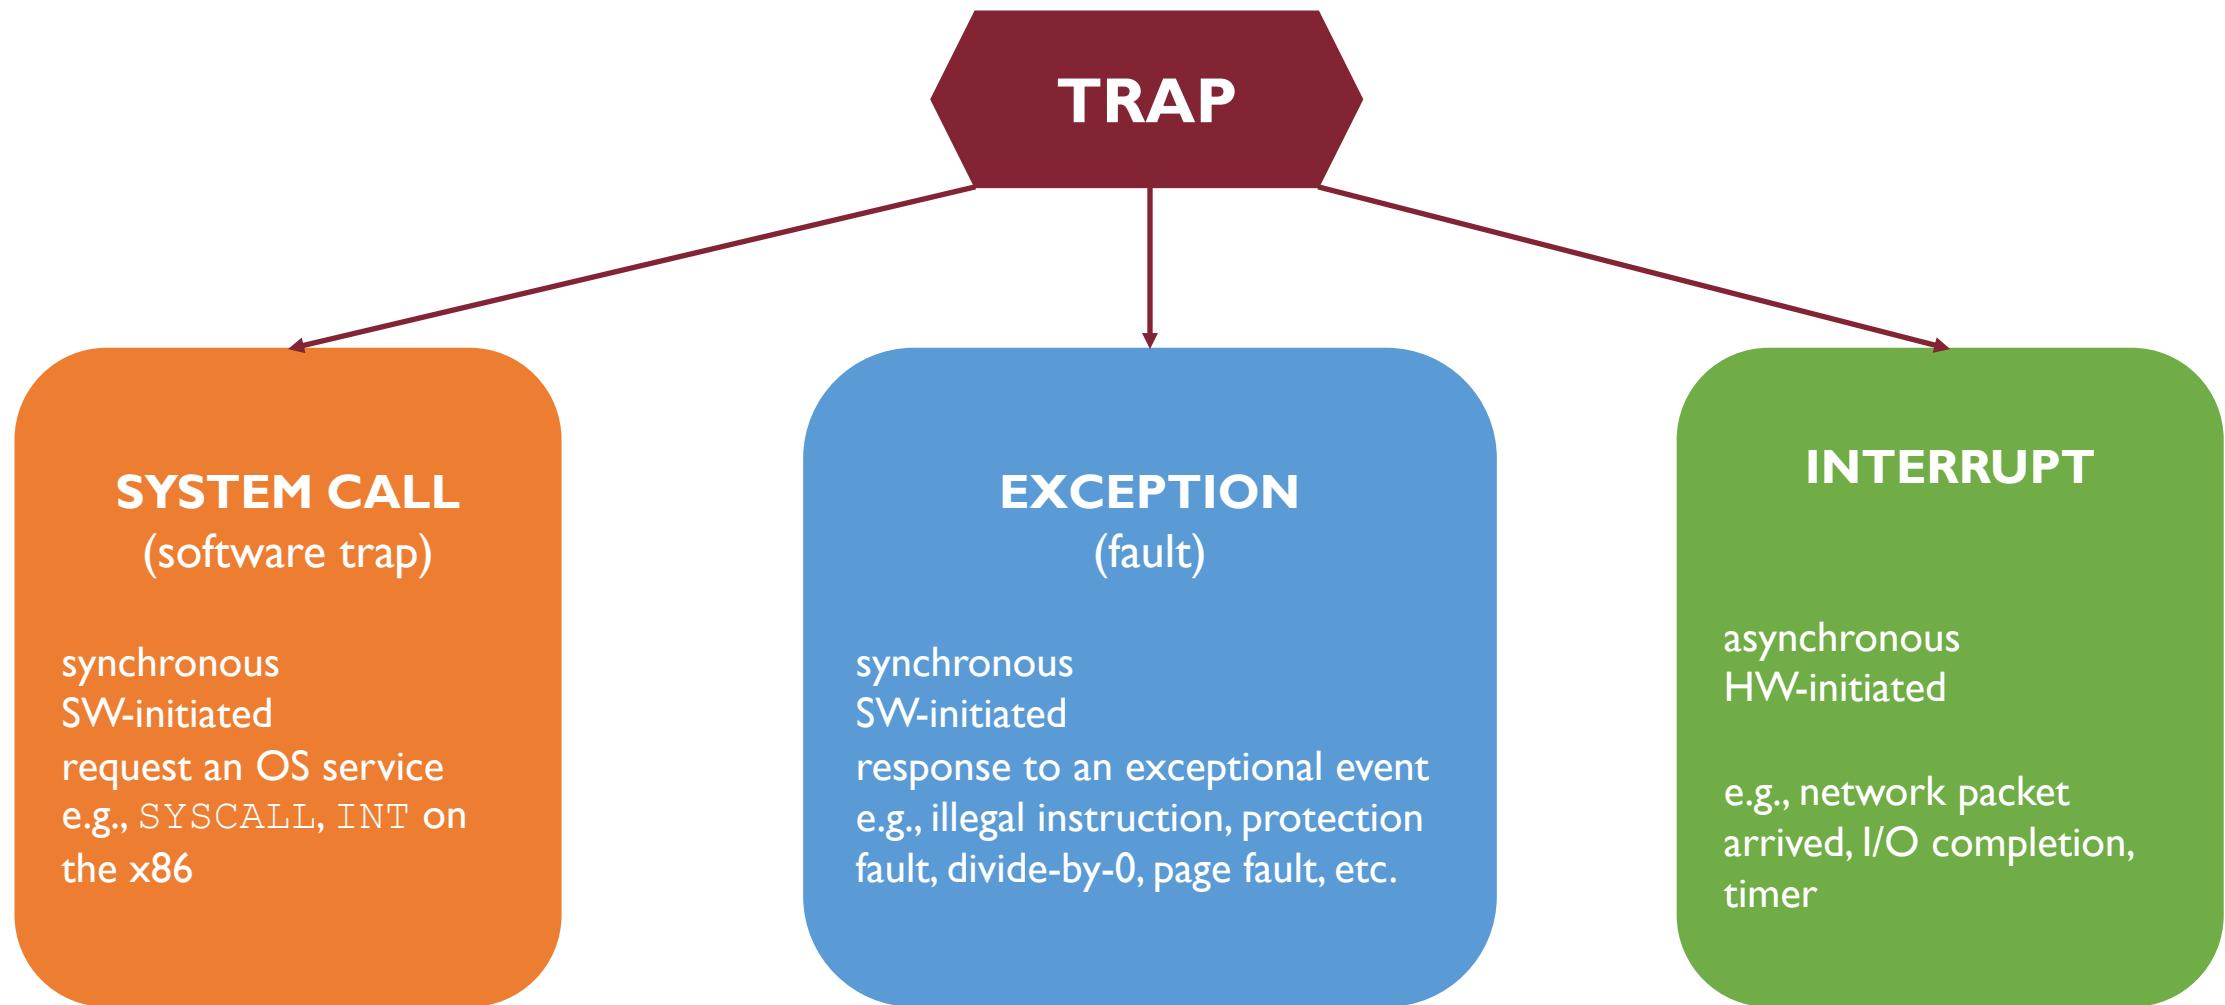
# A Quick Note on Terminology



# A Quick Note on Terminology



# A Quick Note on Terminology



# Scheduling and Synchronization

# Timer

- Hardware facility to enable CPU scheduling

# Timer

- Hardware facility to enable CPU scheduling
- It is just a clock which marks the time of the day

# Timer

- Hardware facility to enable CPU scheduling
- It is just a clock which marks the time of the day
- In multi-tasking systems, allows the CPU not to be monopolized by "selfish" processes

# Timer

- Hardware facility to enable CPU scheduling
- It is just a clock which marks the time of the day
- In multi-tasking systems, allows the CPU not to be monopolized by "selfish" processes
- The timer generates an interrupt every, say, 100 microseconds

# Timer

- Hardware facility to enable CPU scheduling
- It is just a clock which marks the time of the day
- In multi-tasking systems, allows the CPU not to be monopolized by "selfish" processes
- The timer generates an interrupt every, say, 100 microseconds
- At each timer interrupt, the CPU scheduler takes over and decides which process to execute next

# Atomic Instructions

- Interrupts may occur at any time and interfere with running processes

# Atomic Instructions

- Interrupts may occur at any time and interfere with running processes
- OS must be able to synchronize the activities of cooperating, concurrent processes

# Atomic Instructions

- Interrupts may occur at any time and interfere with running processes
- OS must be able to synchronize the activities of cooperating, concurrent processes
- Hardware must ensure that short sequences of instructions (e.g., read-modify-write) are executed **atomically** by either:
  - Disabling interrupts before the sequence and re-enable them afterwards  
or
  - Special instructions that are natively executed atomically

# Virtual Memory

# What is Virtual Memory?

- It is an **abstraction** (of the actual, physical main memory)

# What is Virtual Memory?

- It is an **abstraction** (of the actual, physical main memory)
- It gives each process the illusion that physical memory is just a contiguous address space (virtual address space)

# What is Virtual Memory?

- It is an **abstraction** (of the actual, physical main memory)
- It gives each process the illusion that physical memory is just a contiguous address space (virtual address space)
- It allows to run programs without them being entirely loaded in main memory
  - They are entirely loaded in virtual memory, though!

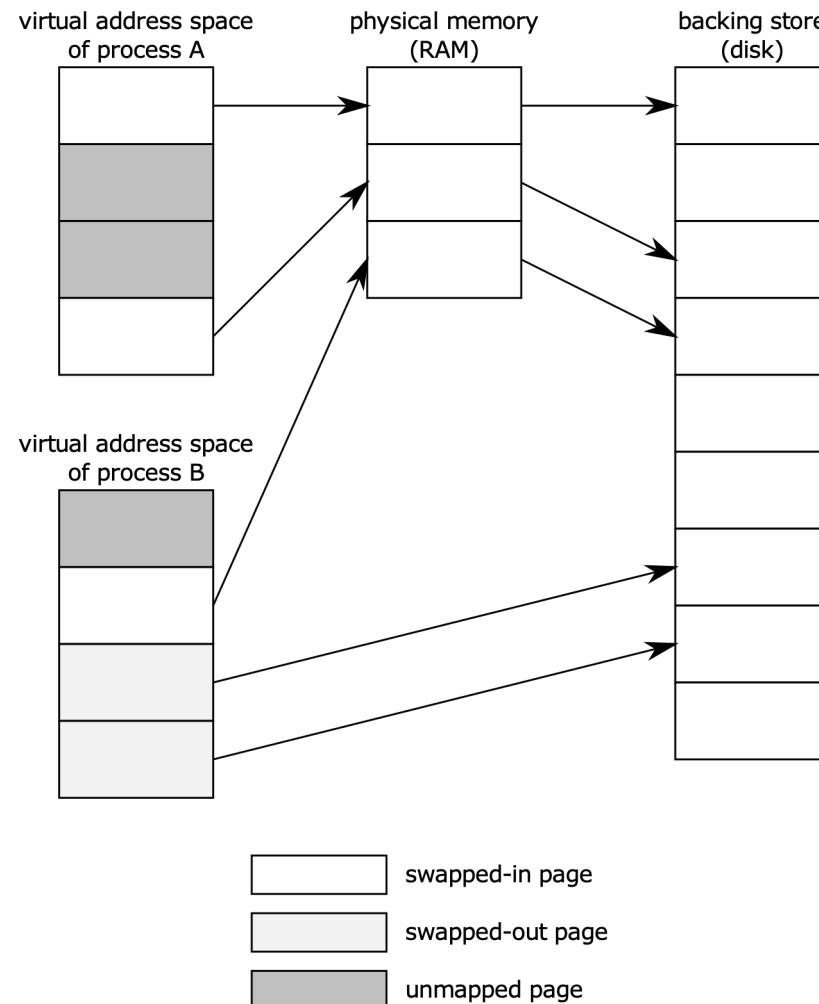
# What is Virtual Memory?

- It is an **abstraction** (of the actual, physical main memory)
- It gives each process the illusion that physical memory is just a contiguous address space (virtual address space)
- It allows to run programs without them being entirely loaded in main memory
  - They are entirely loaded in virtual memory, though!
- Implemented both in HW (**MMU**) and SW (**OS**)
  - **MMU** is responsible for translating virtual addresses into physical ones
  - **OS** is responsible for managing virtual address spaces

# Virtual vs. Physical Address Space

- On a 64 bit system the CPU is able to address  $2^{64}$  bytes = 16 exabytes (EiB)
- Virtual address space ranges from 0 to  $2^{64} - 1$
- This is about a billion times more than main memory capacity currently available!
- Virtual address space is typically divided into contiguous blocks of the same size (e.g., 4 KiB), called **pages**
- Pages which are not loaded in main memory are stored on disk

# Virtual vs. Physical Address Space



# Memory Management Unit (MMU)

- Maps virtual addresses to physical ones through a **Page Table** managed by the OS

# Memory Management Unit (MMU)

- Maps virtual addresses to physical ones through a **Page Table** managed by the OS
- Uses a cache called **Translation Look-aside Buffer (TLB)** with "recent mappings" for quicker lookups

# Memory Management Unit (MMU)

- Maps virtual addresses to physical ones through a **Page Table** managed by the OS
- Uses a cache called **Translation Look-aside Buffer (TLB)** with "recent mappings" for quicker lookups
- The OS must be aware of which pages are loaded in main memory and which ones are on disk

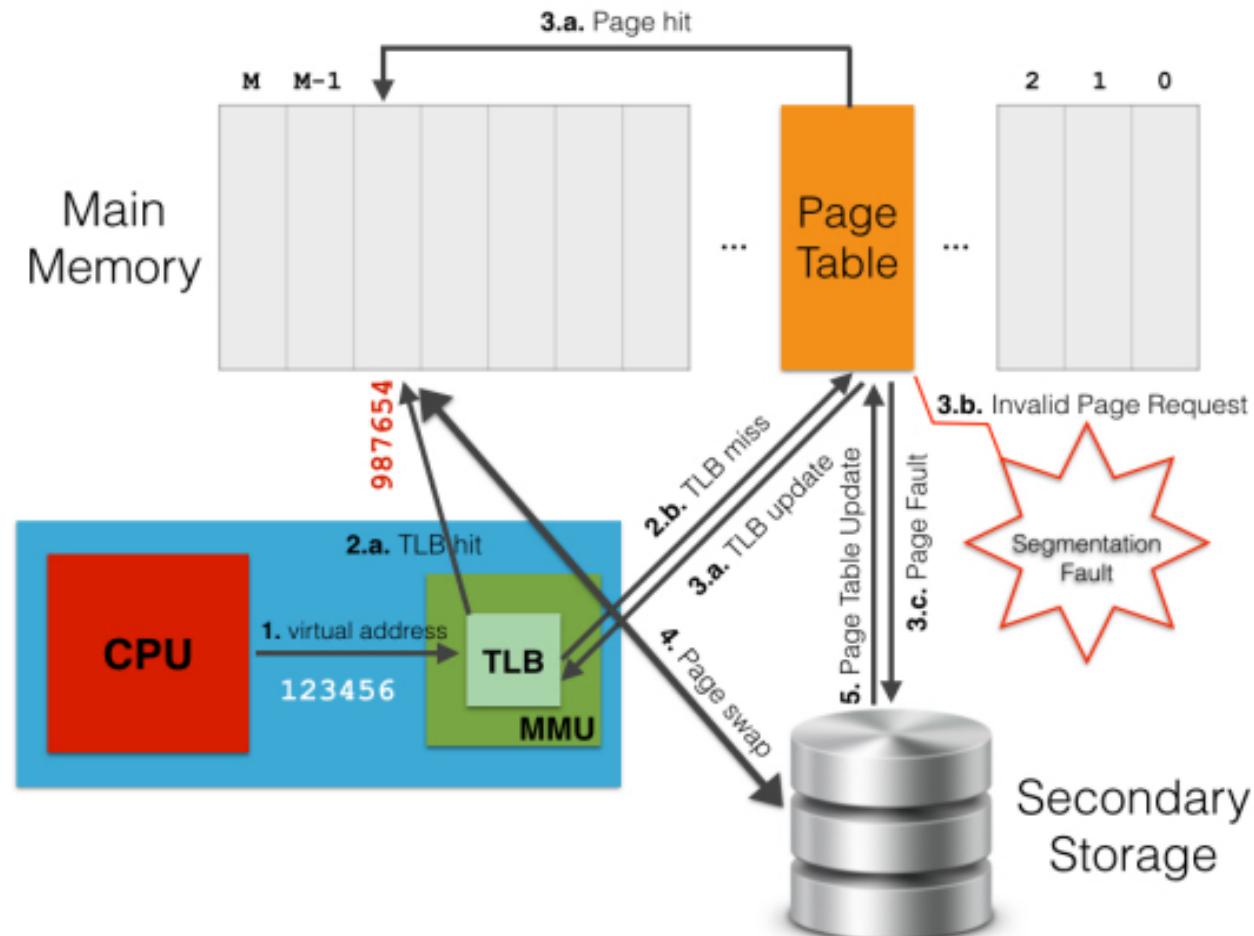
# Lookup Fault

- Page lookup may fail due to **2 reasons**:
  - no valid mapping exists for a virtual address → **segmentation fault**
  - requested page is not in main memory → **page fault**
- Both cases the Page Supervisor (OS kernel) takes over

# Lookup Fault Handling

- In case of a page fault the Page Supervisor:
  - checks if there is enough room in main memory
  - if not, it has to free some space by removing a page from main memory and storing to disk (**swapping**)
    - several page replacement algorithms (e.g., LRU)
    - retrieves the missing page from disk and stores it in main memory
    - updates the Page Table and the TLB cache

# Lookup Fault Handling



# Summary

- Architecture support is key to OS design

# Summary

- Architecture support is key to OS design
- Most of the services provided by the OS to the applications rely on specific HW features

# Summary

- Architecture support is key to OS design
- Most of the services provided by the OS to the applications rely on specific HW features
- The OS is tightly coupled to the HW of the host machine

# Summary

- Architecture support is key to OS design
- Most of the services provided by the OS to the applications rely on specific HW features
- The OS is tightly coupled to the HW of the host machine
- Advice: Keep your Computer Architecture book at hand!